

Meta-Learning Frontiers: Universal, Uncertain, and Unsupervised

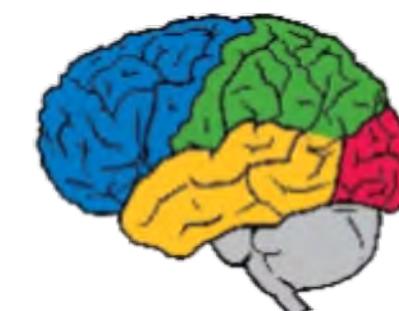
Sergey Levine

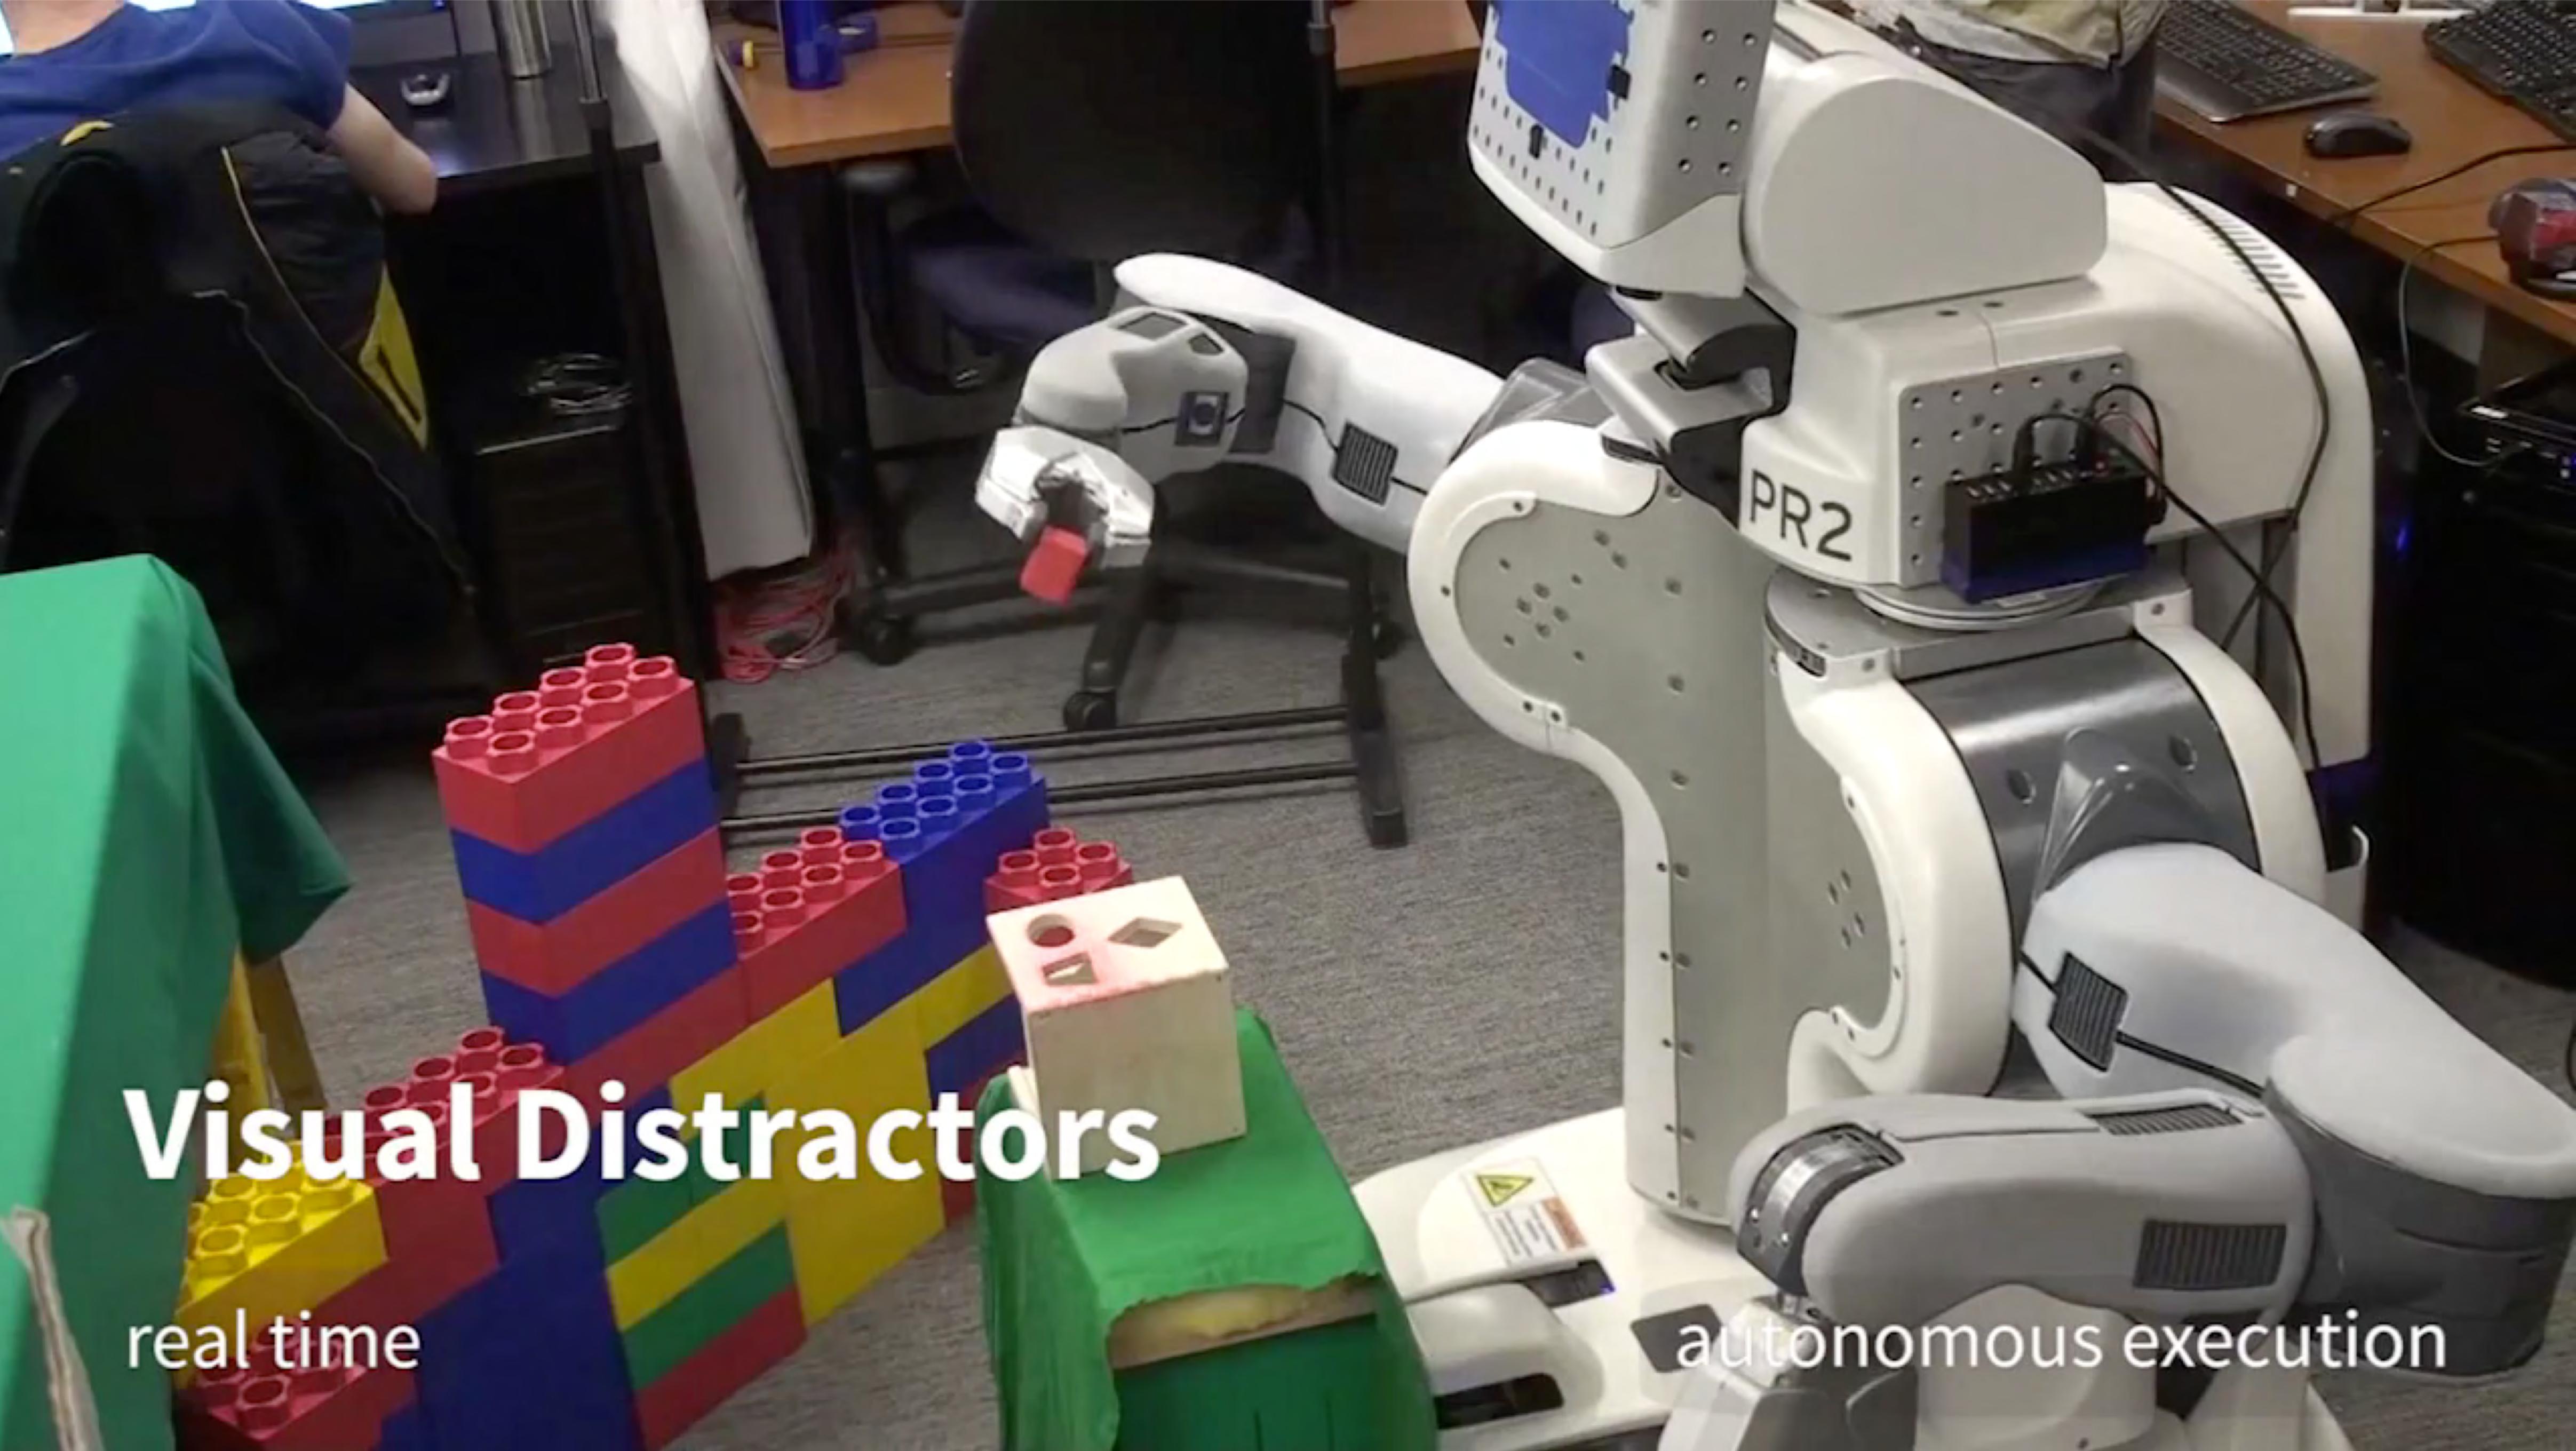
Chelsea Finn

UC Berkeley



Google Brain

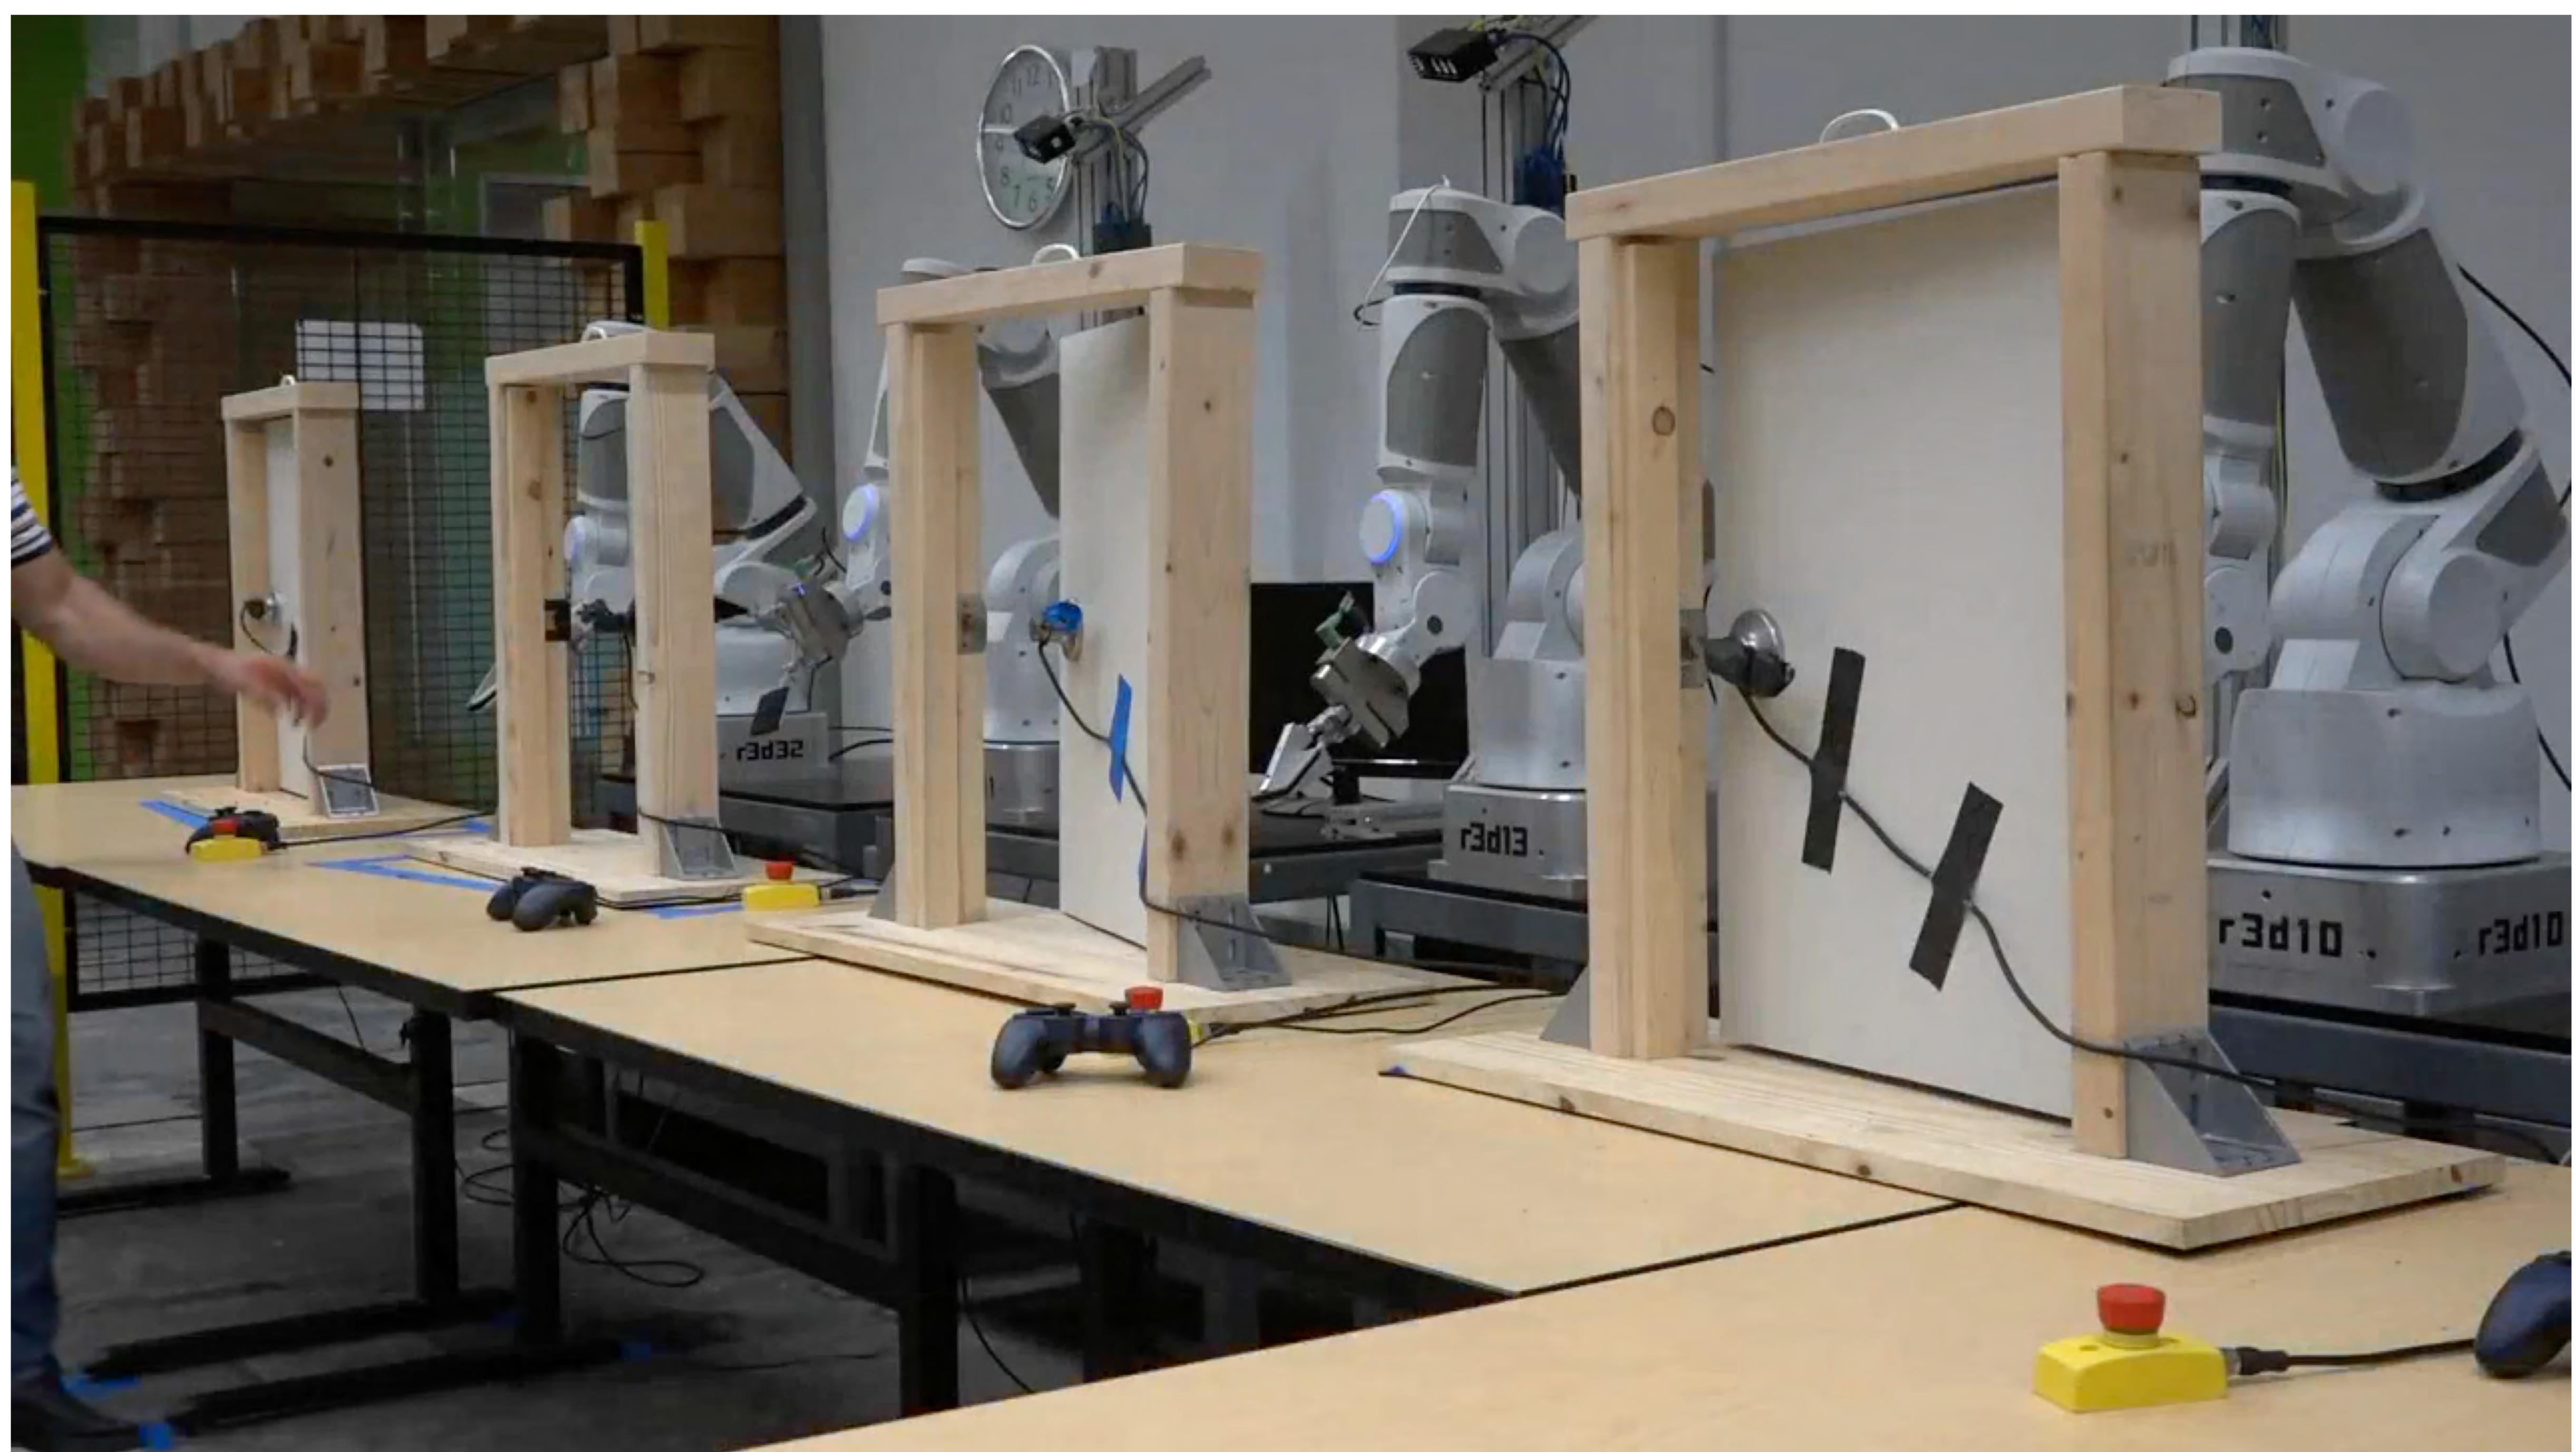




Visual Distractors

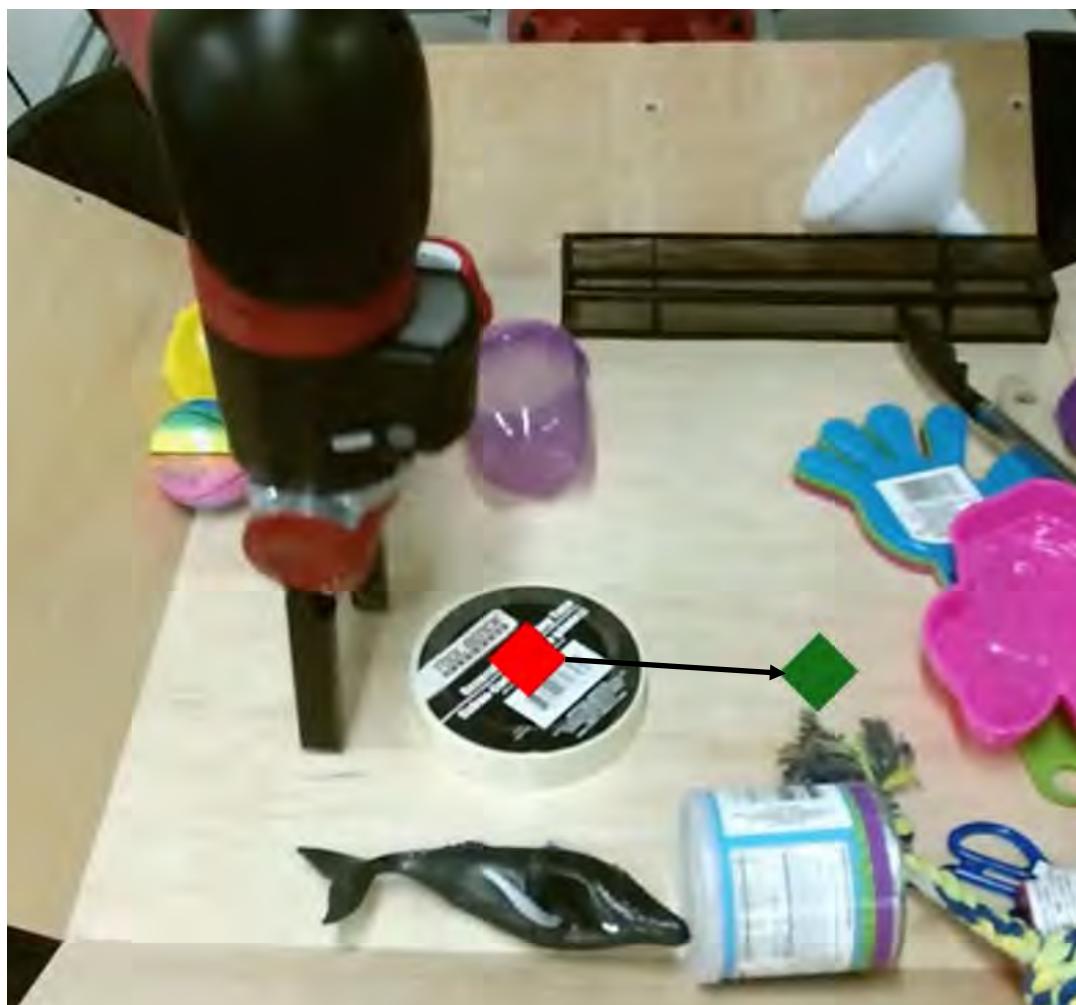
real time

autonomous execution

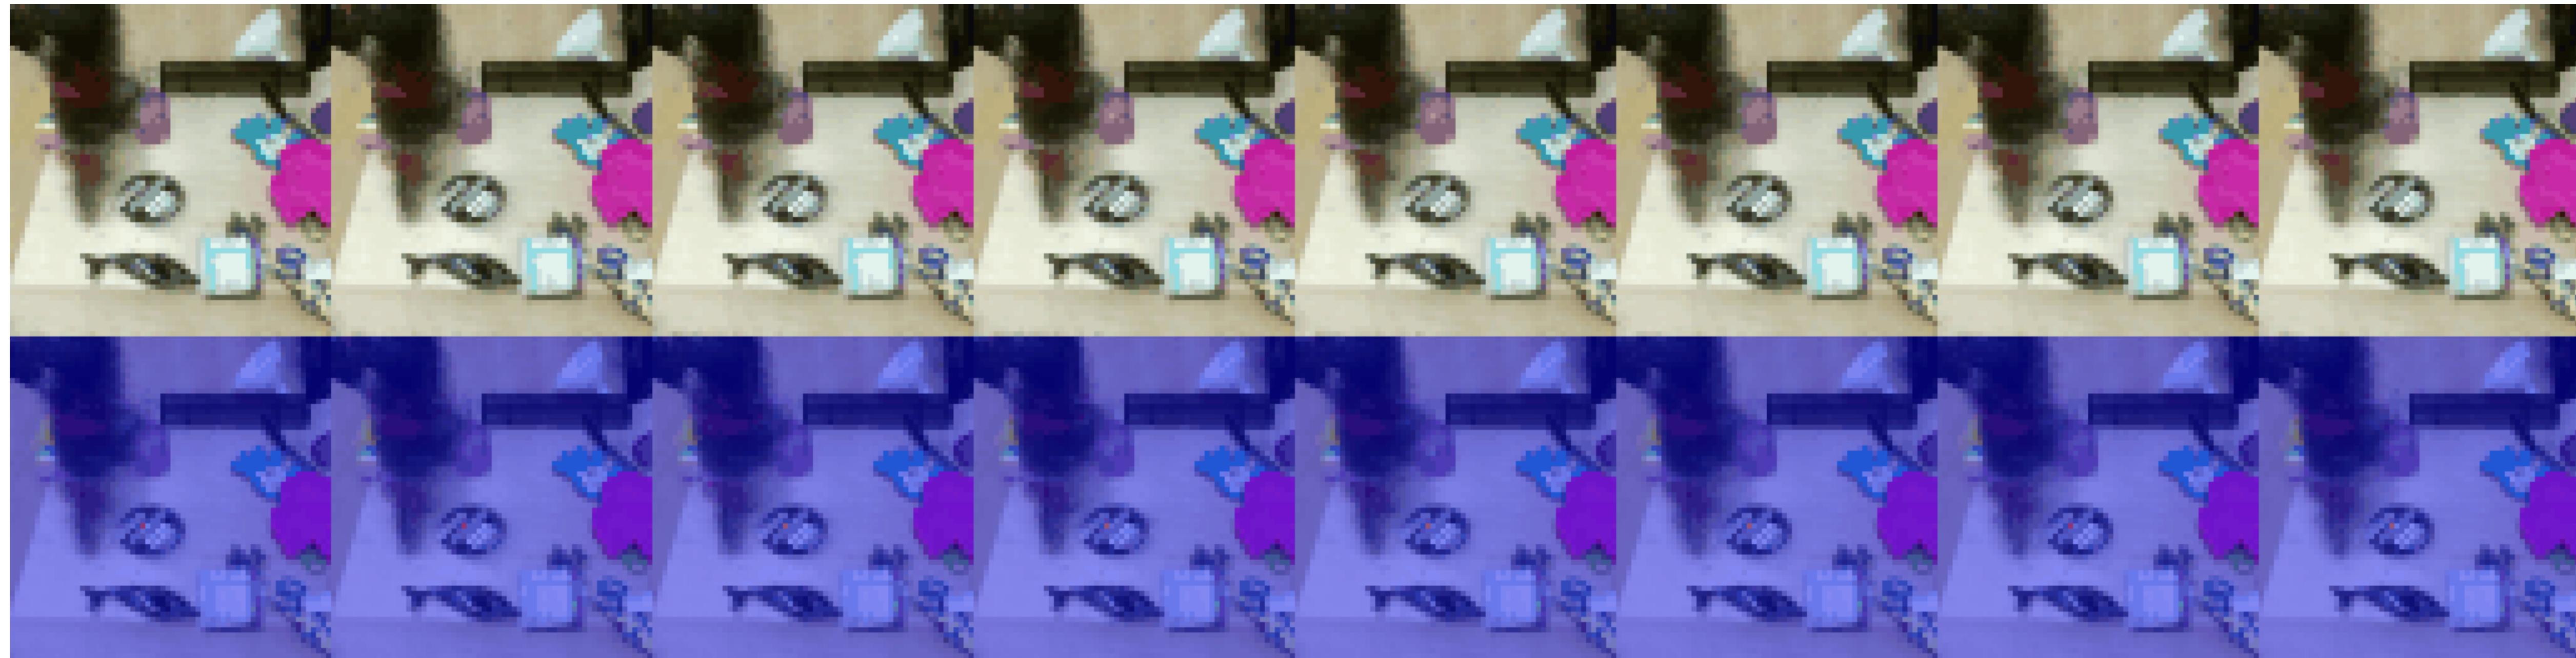


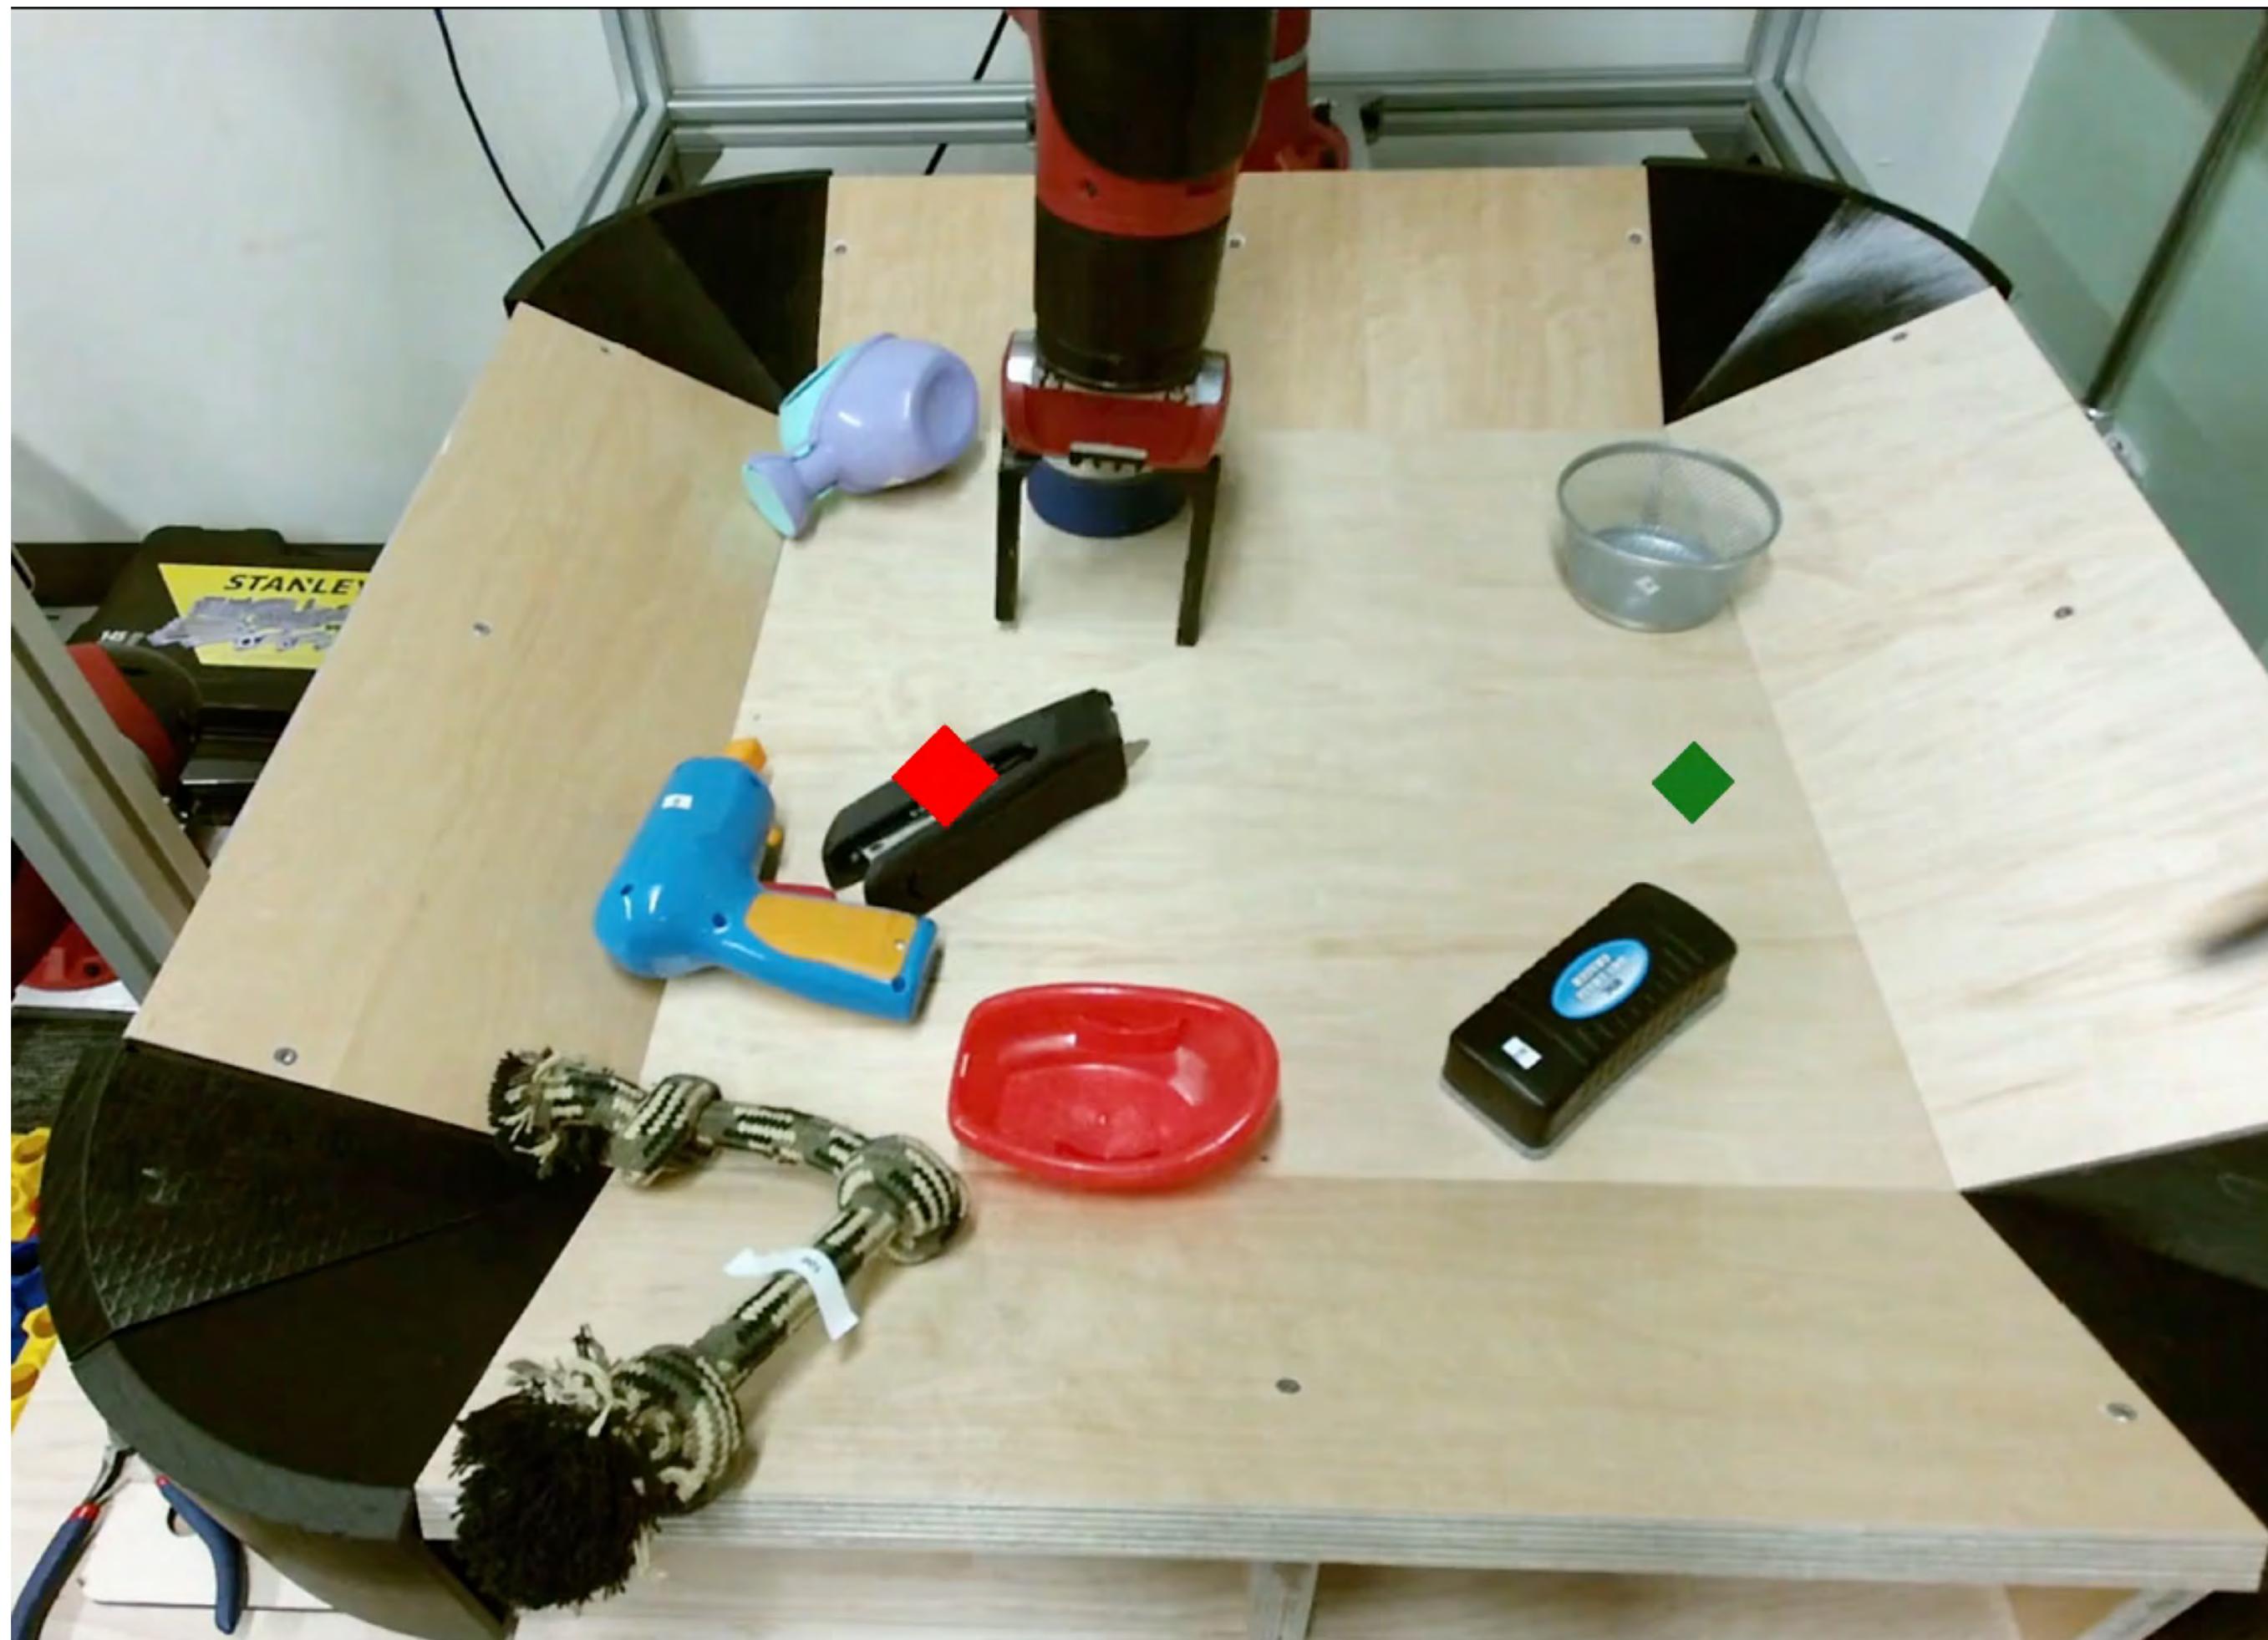
Yahya, Li, Kalakrishnan, Chebotar, Levine, '16

Generalizable model-based RL via video prediction



Designated Pixel ◆
Goal Pixel ◆

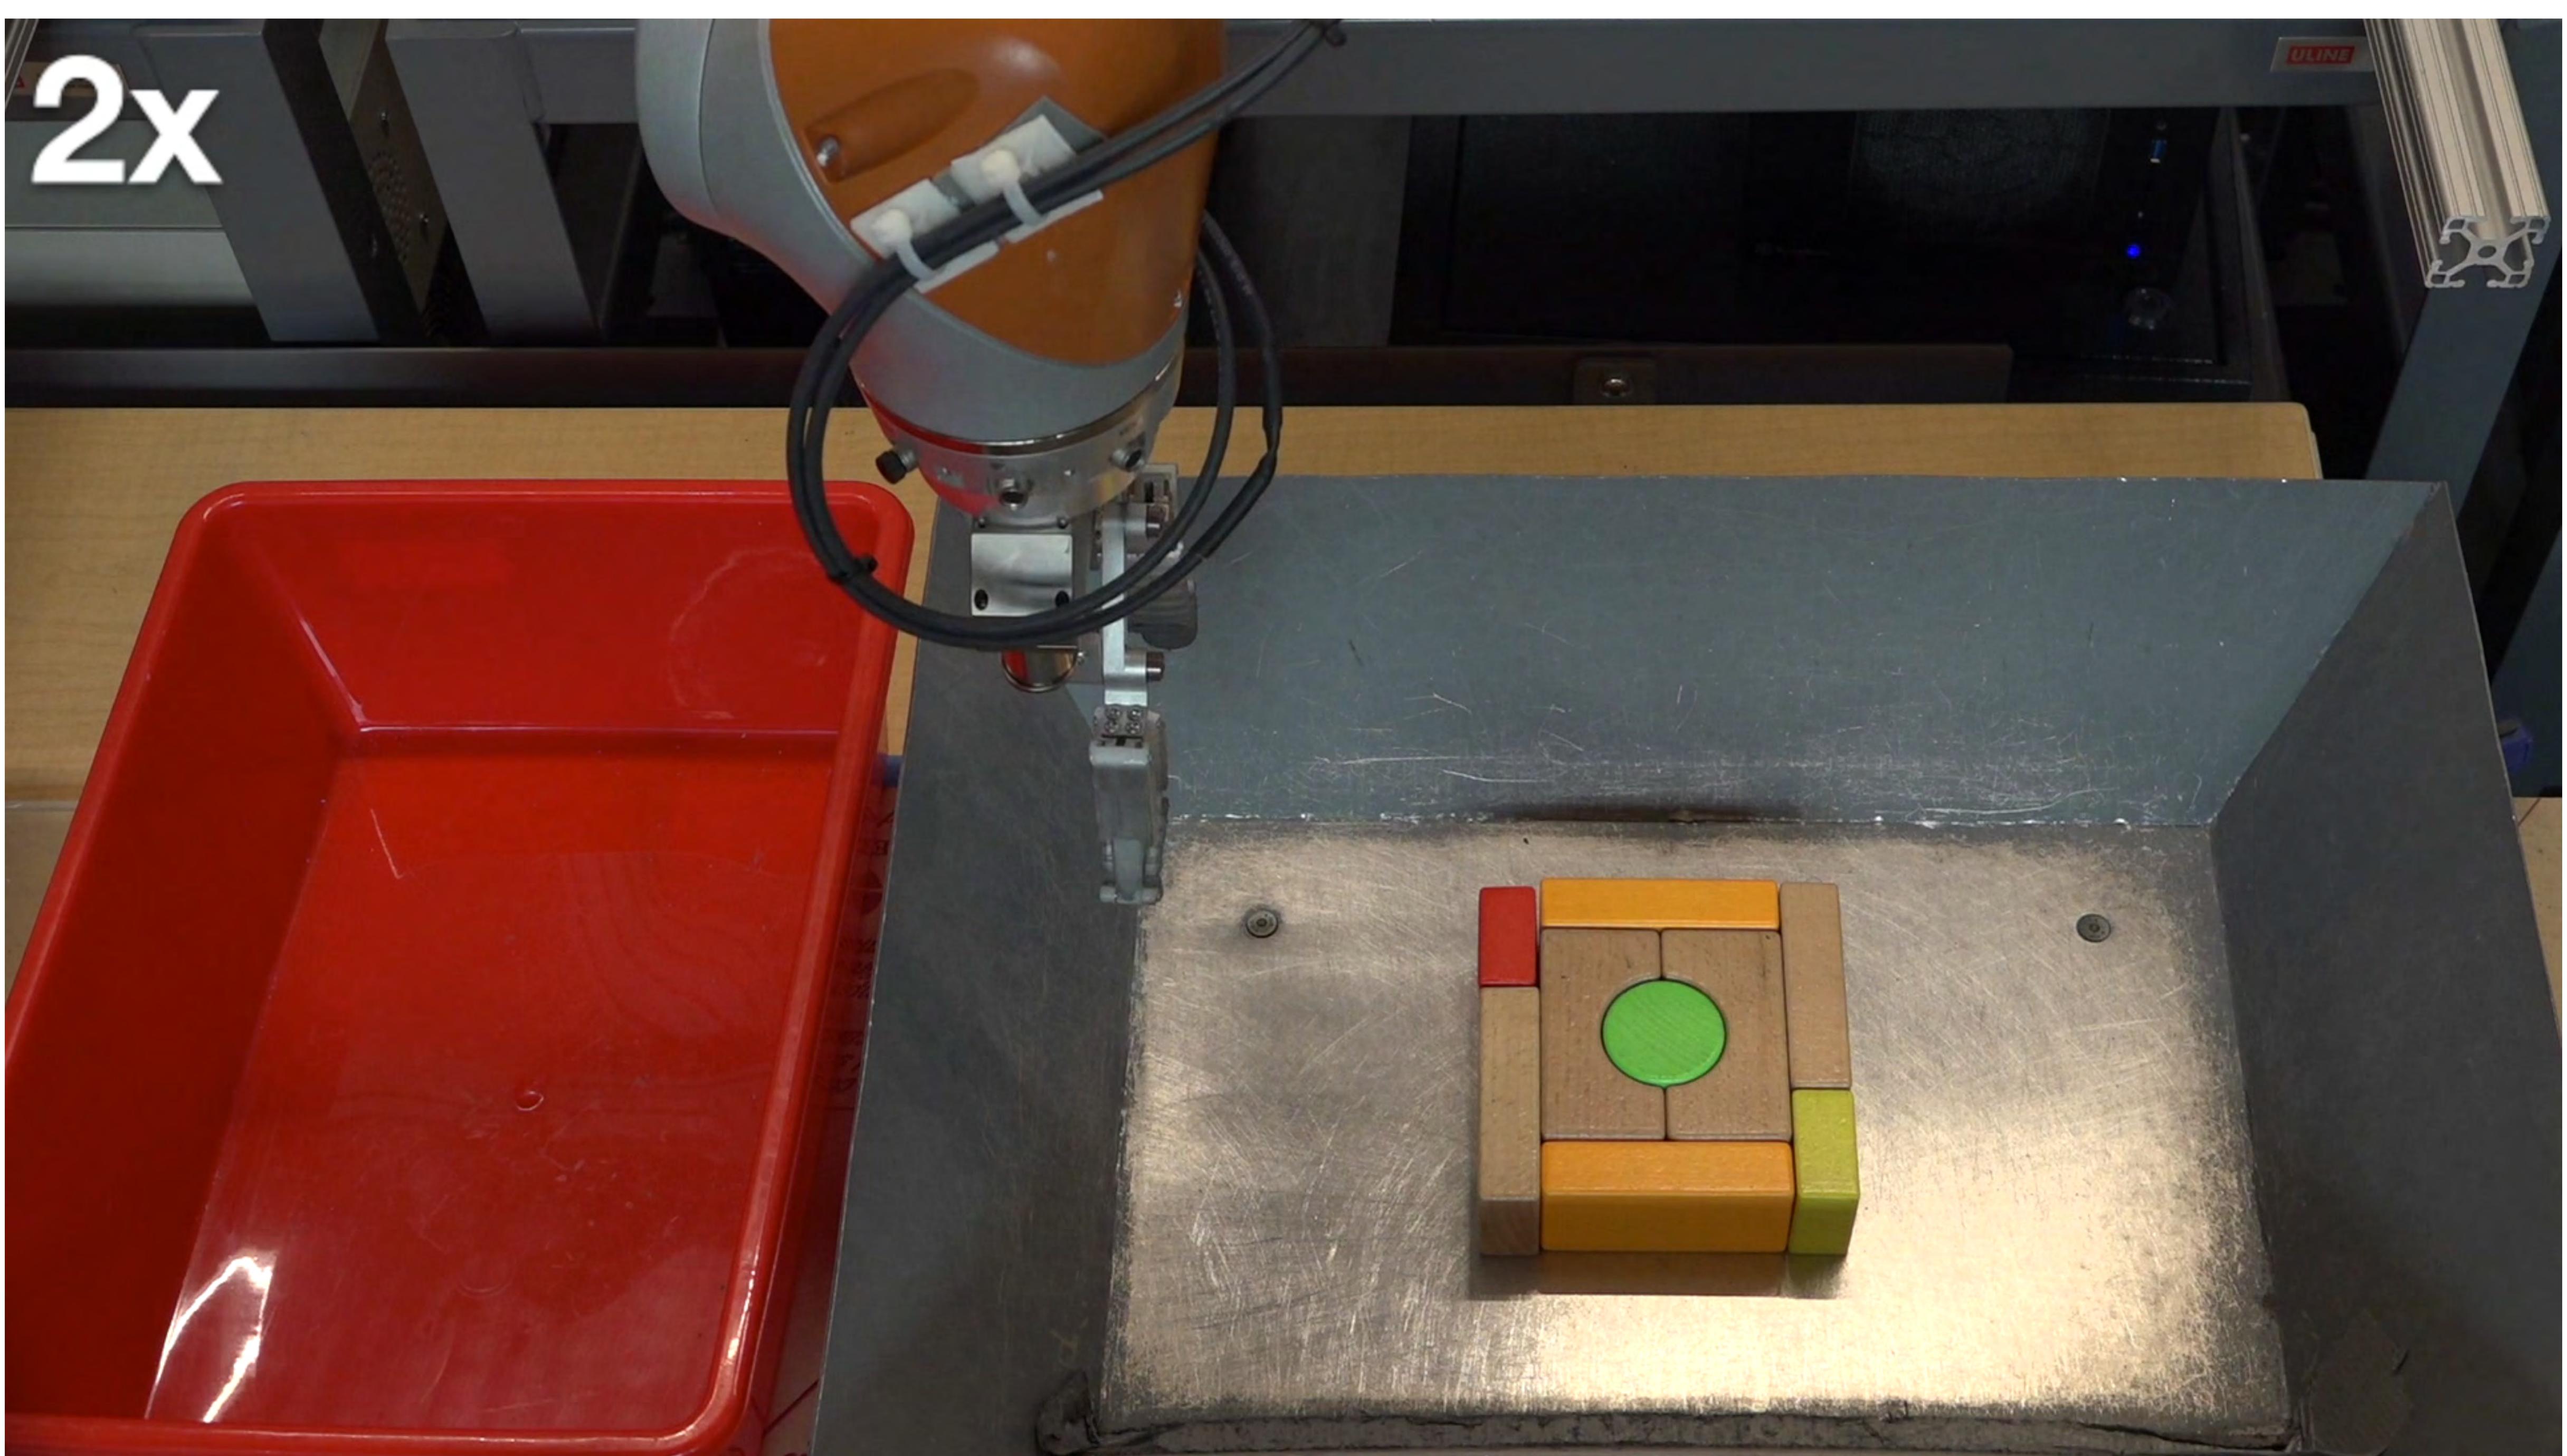


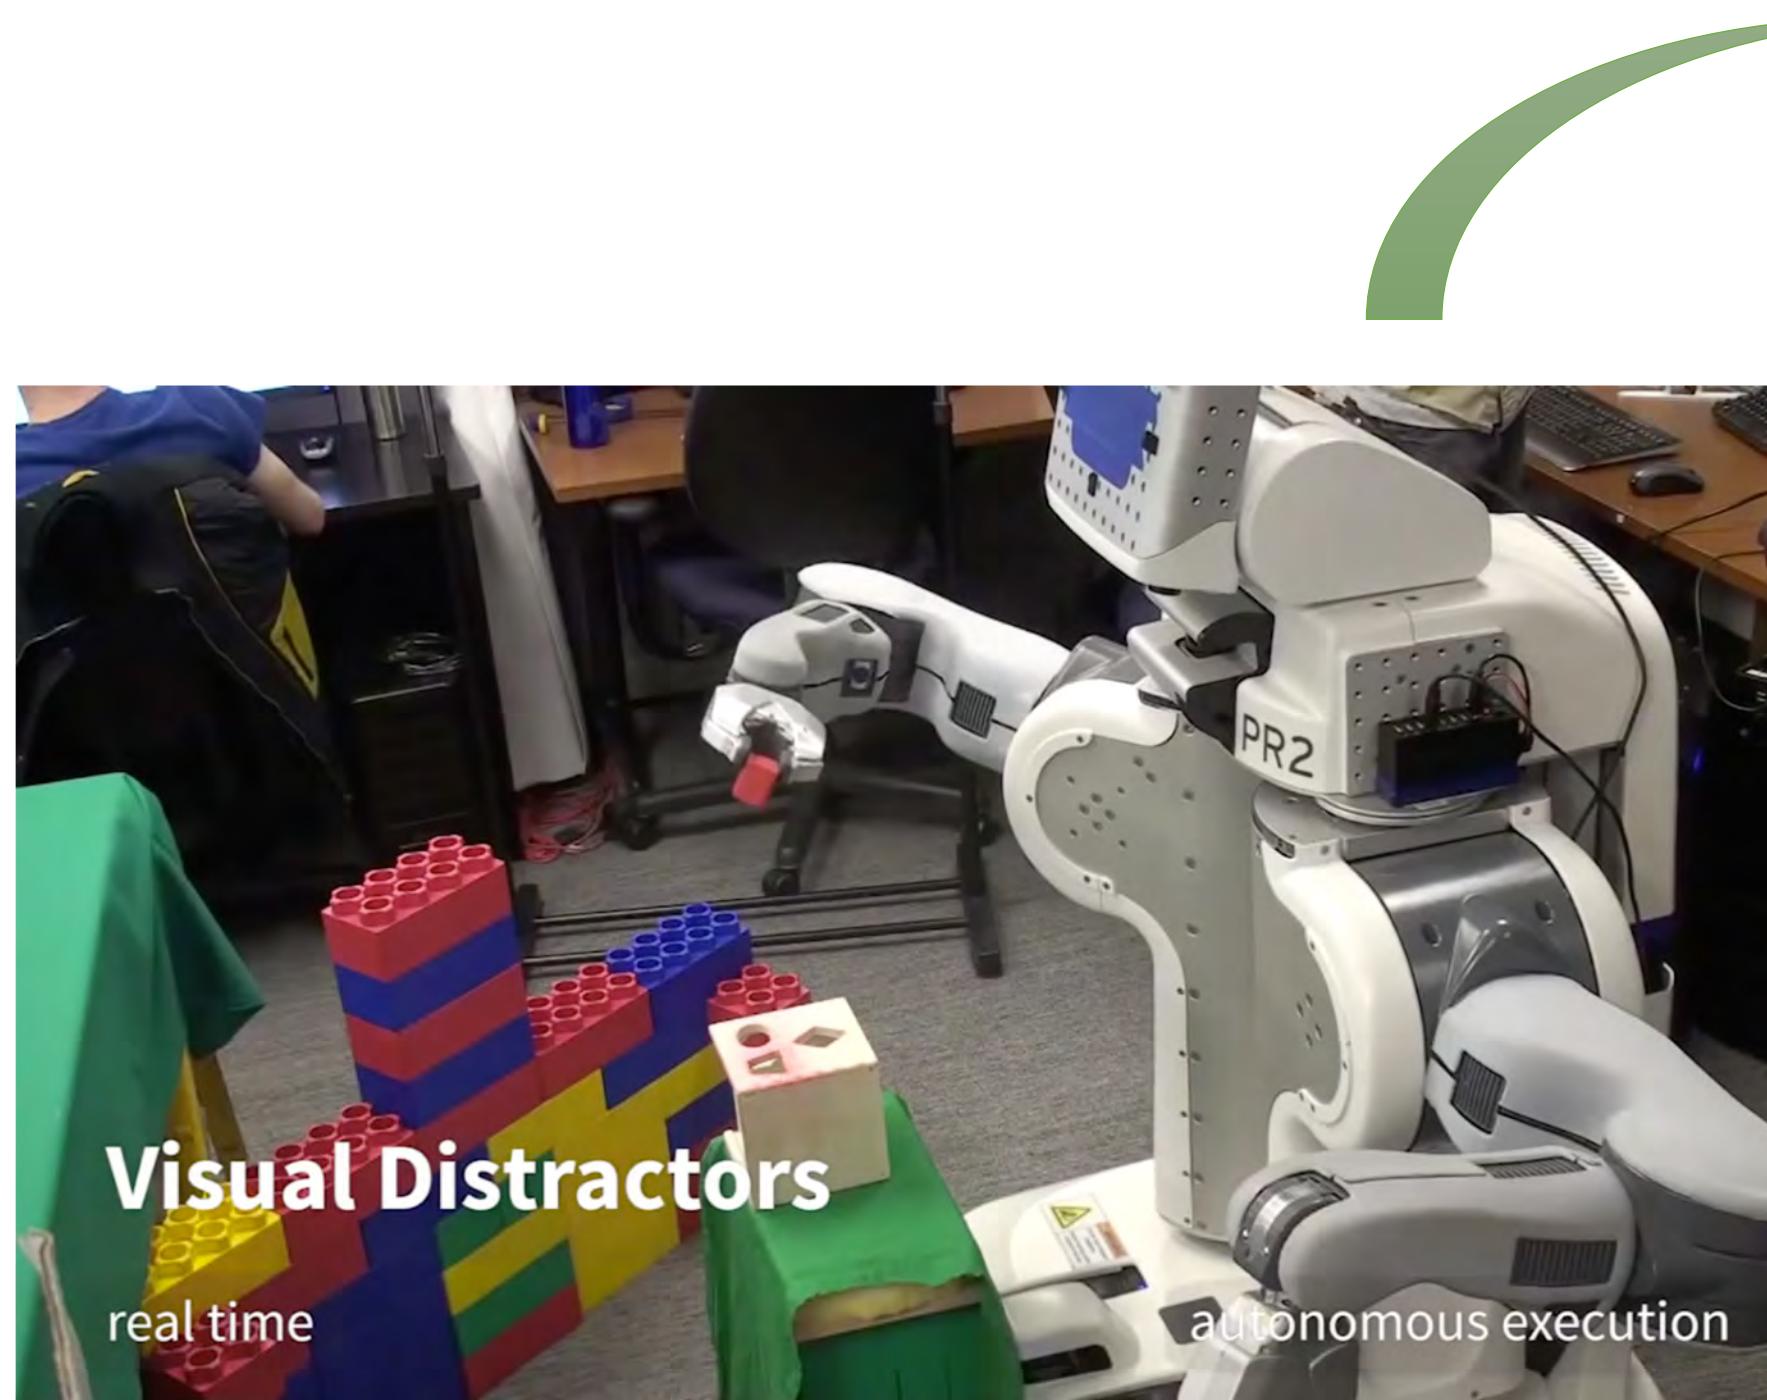




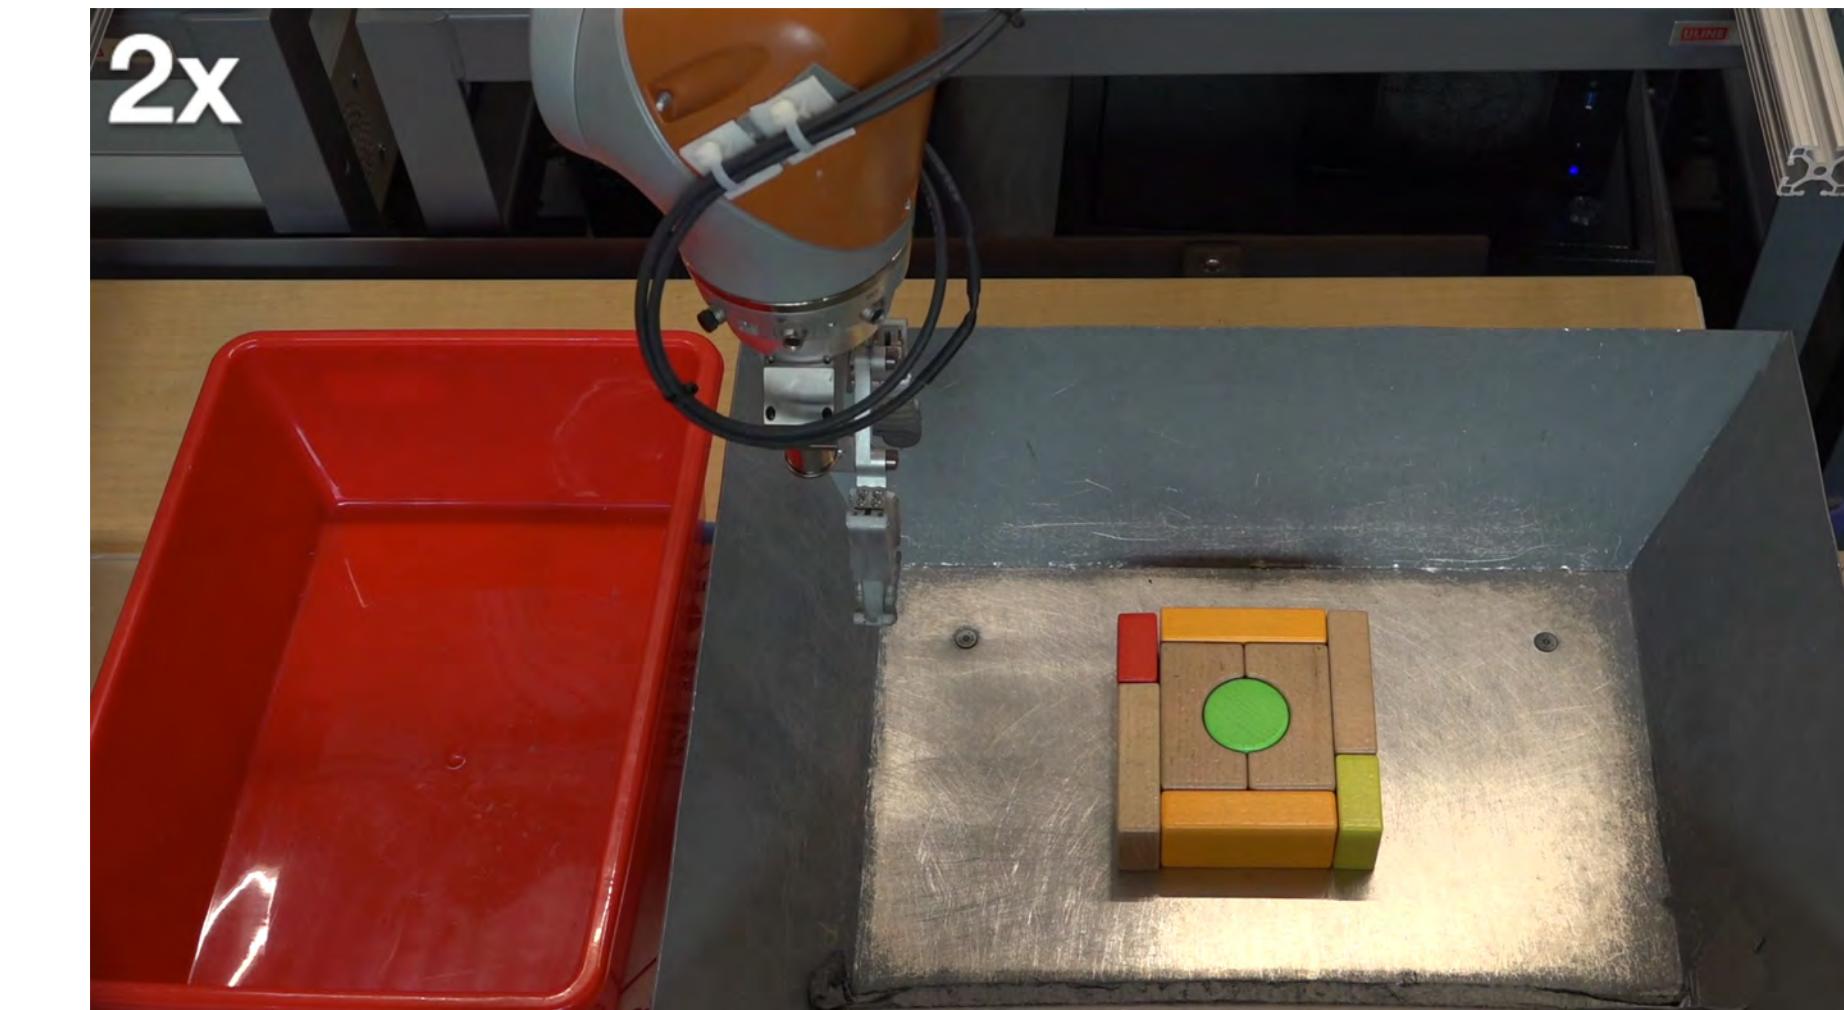
Kalashnikov, Irpan, Pastor, Ibarz, Herzong, Jang, Quillen, Holly, Kalakrishnan, Vanhoucke, Levine. QT-Opt: Scalable Deep Reinforcement Learning of Vision-Based Robotic Manipulation Skills

2x





about four hours



about four weeks, nonstop



people can learn new skills
extremely quickly
how?
we never learn from scratch!

what to transfer?
representations?
models?
can we just optimize for what we
really want?
can we *learn to learn*?

Outline

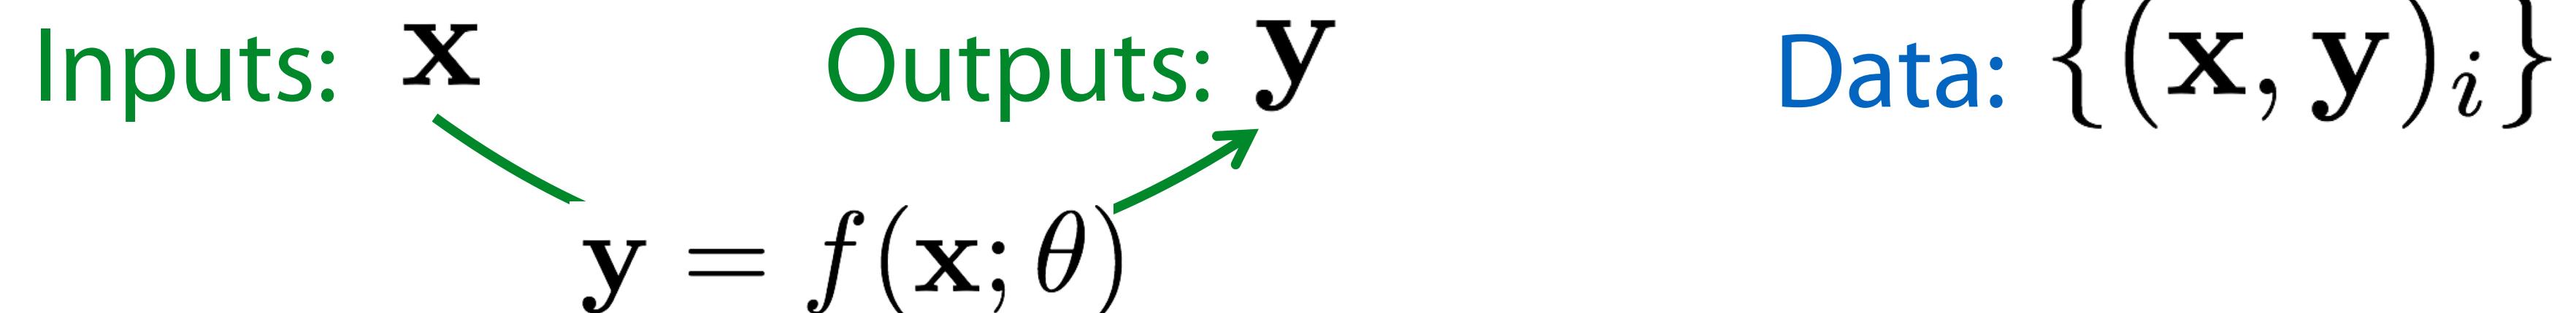
- Meta-Learning Problem Statement
- Model-Agnostic Meta-Learning (MAML)
- Probabilistic Interpretation of MAML
- Meta-Learning with Automated Task Proposals
- Extensions to Robot Imitation & Intent Inference

Outline

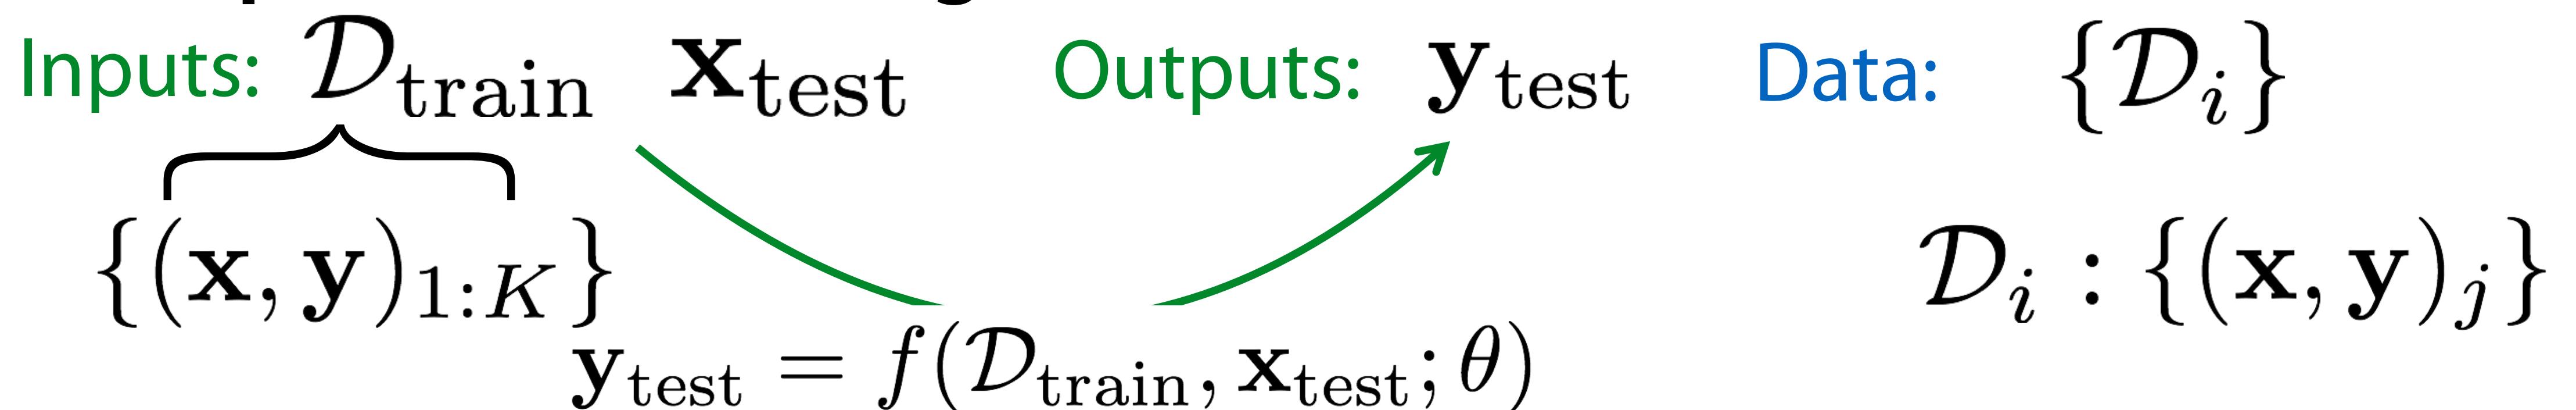
- **Meta-Learning Problem Statement**
- Model-Agnostic Meta-Learning (MAML)
- Probabilistic Interpretation of MAML
- Meta-Learning with Automated Task Proposals
- Extensions to Robot Imitation & Intent Inference

The Meta-Learning Problem

Supervised Learning:



Meta-Supervised Learning:



Why is this view useful?

Reduces the problem to the design & optimization of f .

Example: Few-Shot Classification

Given 1 example of 5 classes:



training data $\mathcal{D}_{\text{train}}$

test set \mathbf{X}_{test}



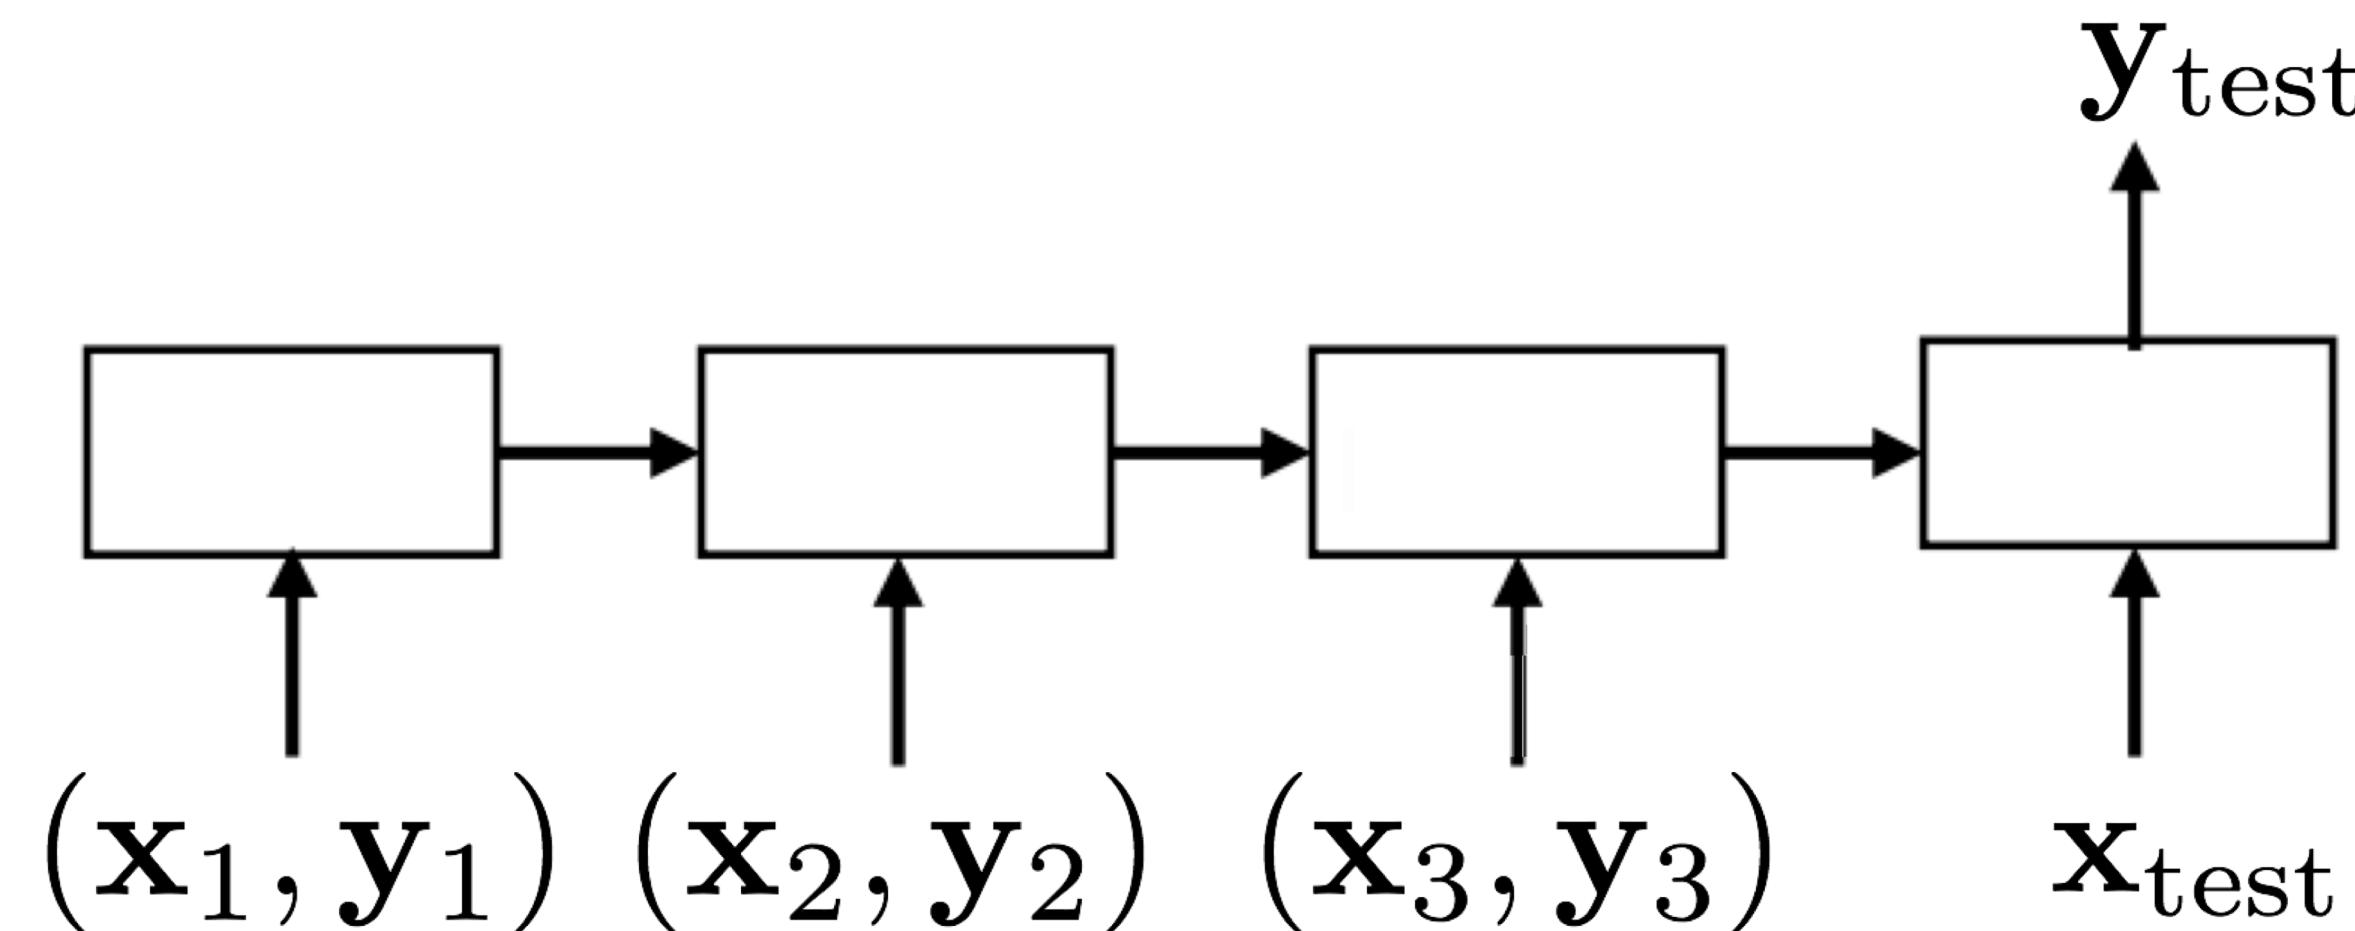
Outline

- Meta-Learning Problem Statement
- **Model-Agnostic Meta-Learning (MAML)**
- Probabilistic Interpretation of MAML
- Meta-Learning with Automated Task Proposals
- Extensions to Robot Imitation & Intent Inference

Design of f ?

$$\mathcal{D}_{\text{train}} \quad \mathbf{x}_{\text{test}} \longrightarrow \mathbf{y}_{\text{test}}$$

Recurrent network $\mathbf{y}_{\text{test}} = f(\mathcal{D}_{\text{train}}, \mathbf{x}_{\text{test}}; \theta)$ Santoro et al. '16, Duan et al. '17, Wang et al. '17,
(LSTM, NTM, Conv) Munkhdalai & Yu '17, Mishra et al. '17, ...



- complex model for complex task of learning
 - impractical data requirements

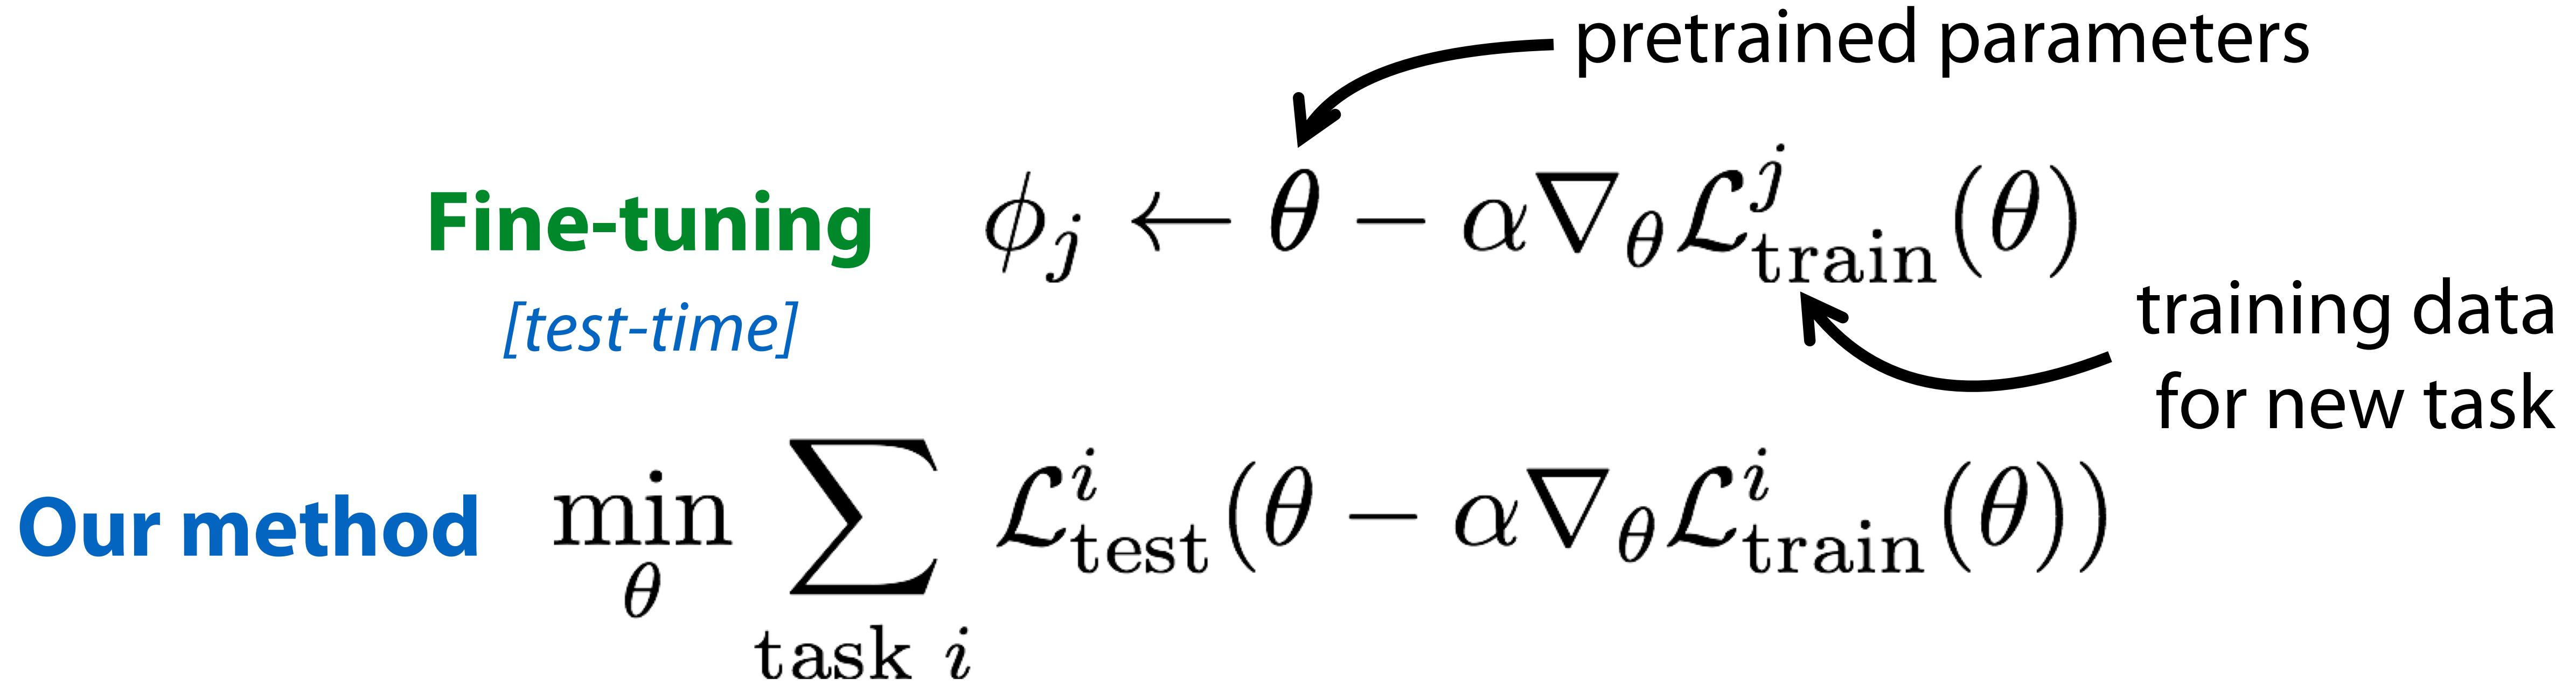
Learning Few-Shot Adaptation

Fine-tuning $\phi_j \leftarrow \theta - \alpha \nabla_{\theta} \mathcal{L}_{\text{train}}^j(\theta)$

[*test-time*] pretrained parameters

Our method $\min_{\theta} \sum_{\text{task } i} \mathcal{L}_{\text{test}}^i(\theta - \alpha \nabla_{\theta} \mathcal{L}_{\text{train}}^i(\theta))$

training data
for new task



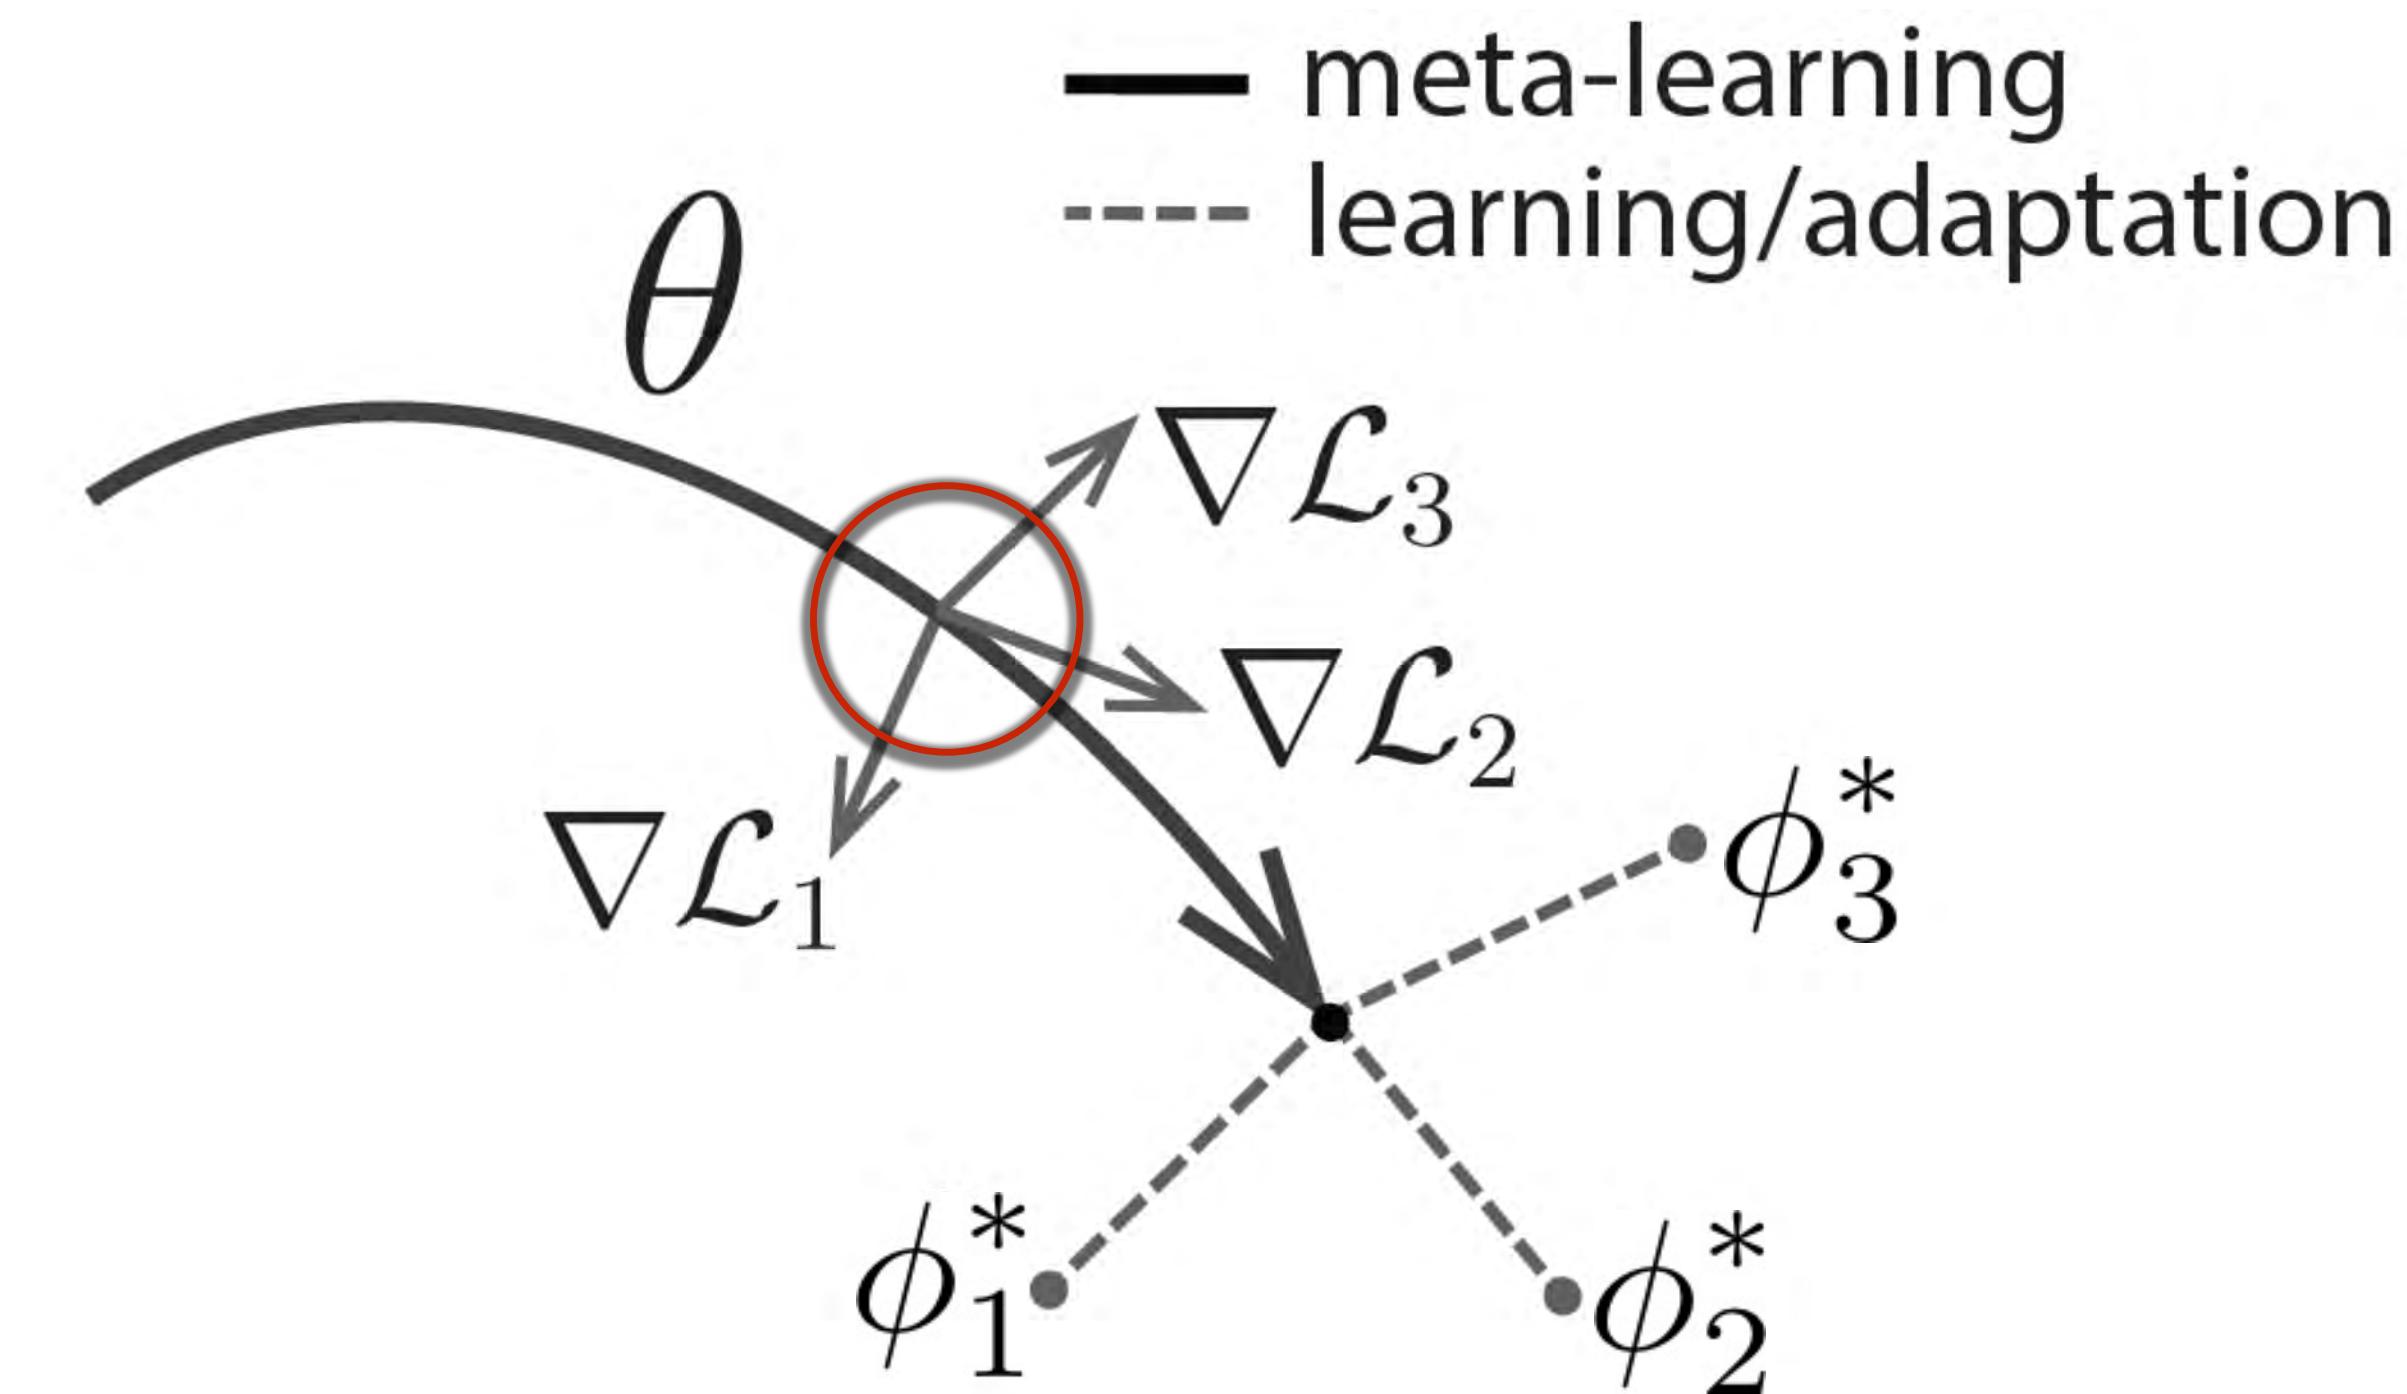
Key idea: Train over many tasks, to learn parameter vector θ that transfers

Learning Few-Shot Adaptation

$$\min_{\theta} \sum_{\text{task } i} \mathcal{L}_{\text{test}}^i(\theta - \alpha \nabla_{\theta} \mathcal{L}_{\text{train}}^i(\theta))$$

θ parameter vector
being meta-learned

ϕ_i^* optimal parameter
vector for task i



Minilmagenet Few-shot Classification



...and the results keep getting better

Minilmagenet few-shot benchmark: 5-shot 5-way

Finn et al. '17: 63.11%

Li et al. '17: 64.03%

Kim et al. '18 (AutoMeta): 76.29%

Program Synthesis

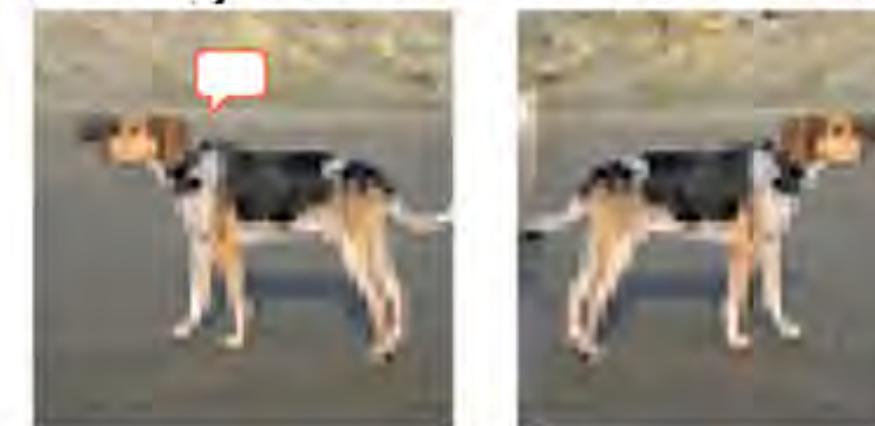
Question:
How many CFL teams are from York College?

SQL:
`SELECT COUNT CFL TEAM FROM CFLDraft WHERE College = "York"`

Result:
2

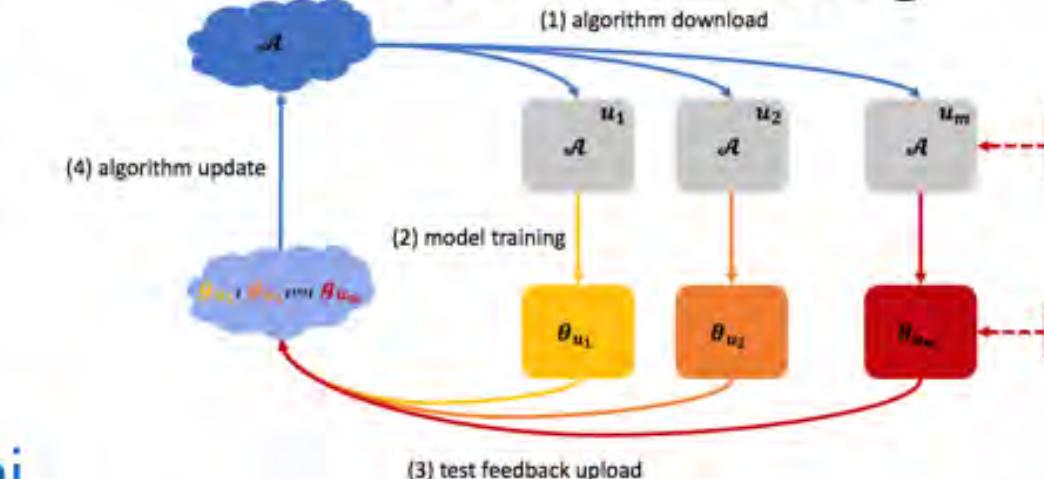
Huang, Wang, Singh,
Yih, He NAACL '18

Learning to Learn Distributions



Reed, Chen, Paine, van den Oord, Eslami,
Rezende, Vinyals, de Freitas ICLR '18

Federated Learning



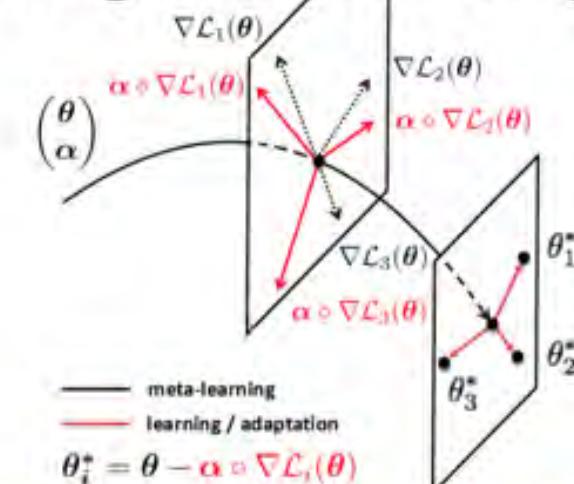
Chen, Dong, Li, He arXiv '18

Multi-Agent Competitions



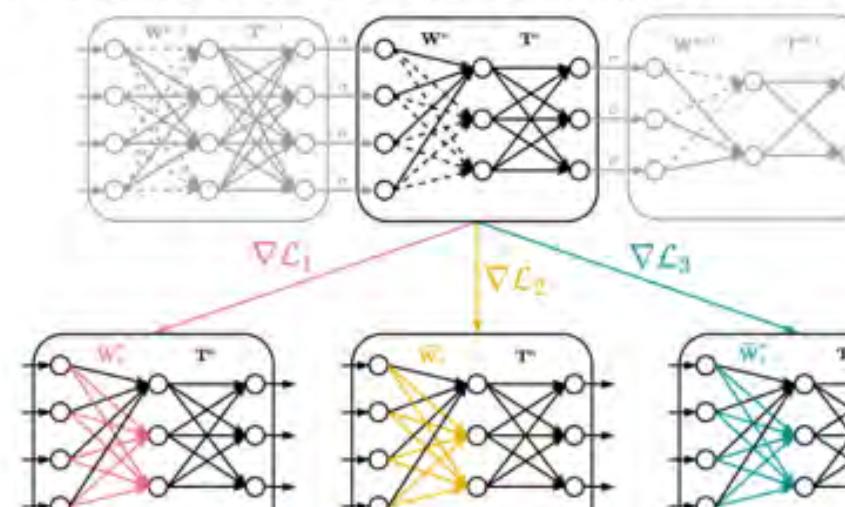
AI-Shedivat, Bansal, Burda, Sutskever
Mordatch, Abbeel ICLR '18

Learning the learning rate



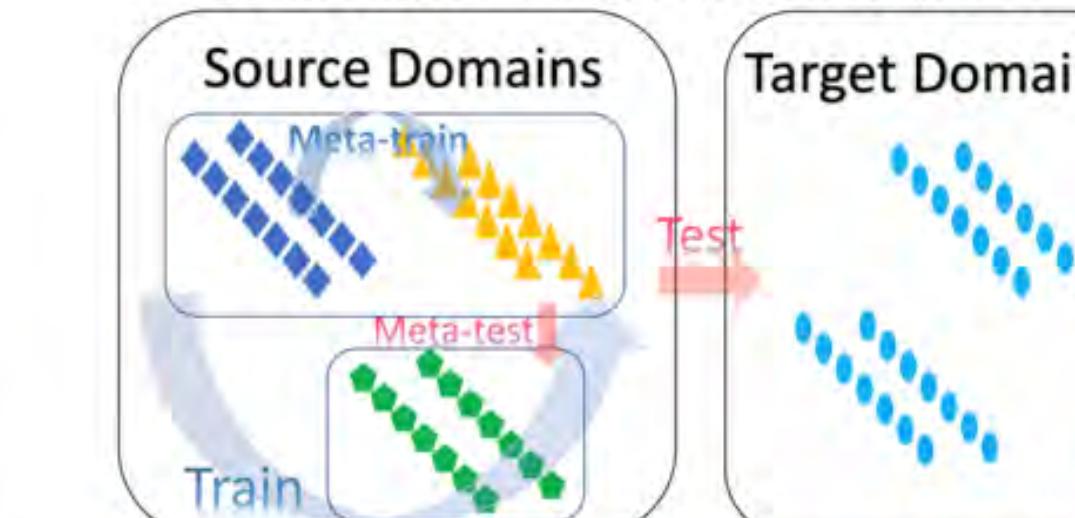
Li, Zhou, Chen, Li arXiv '17

Masked Transformations



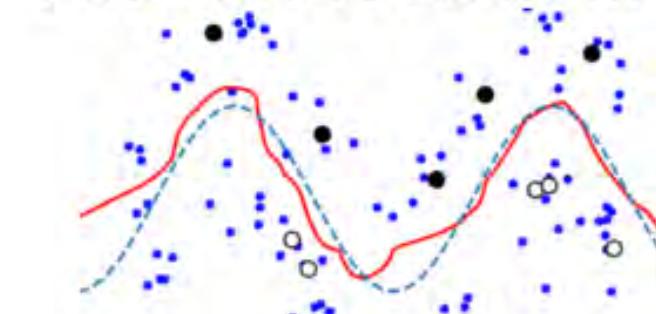
Lee & Choi arXiv '18

Domain Generalization



Li, Yang, Song, Hospedales AAAI '18

Semi-Supervised Few-Shot Learning



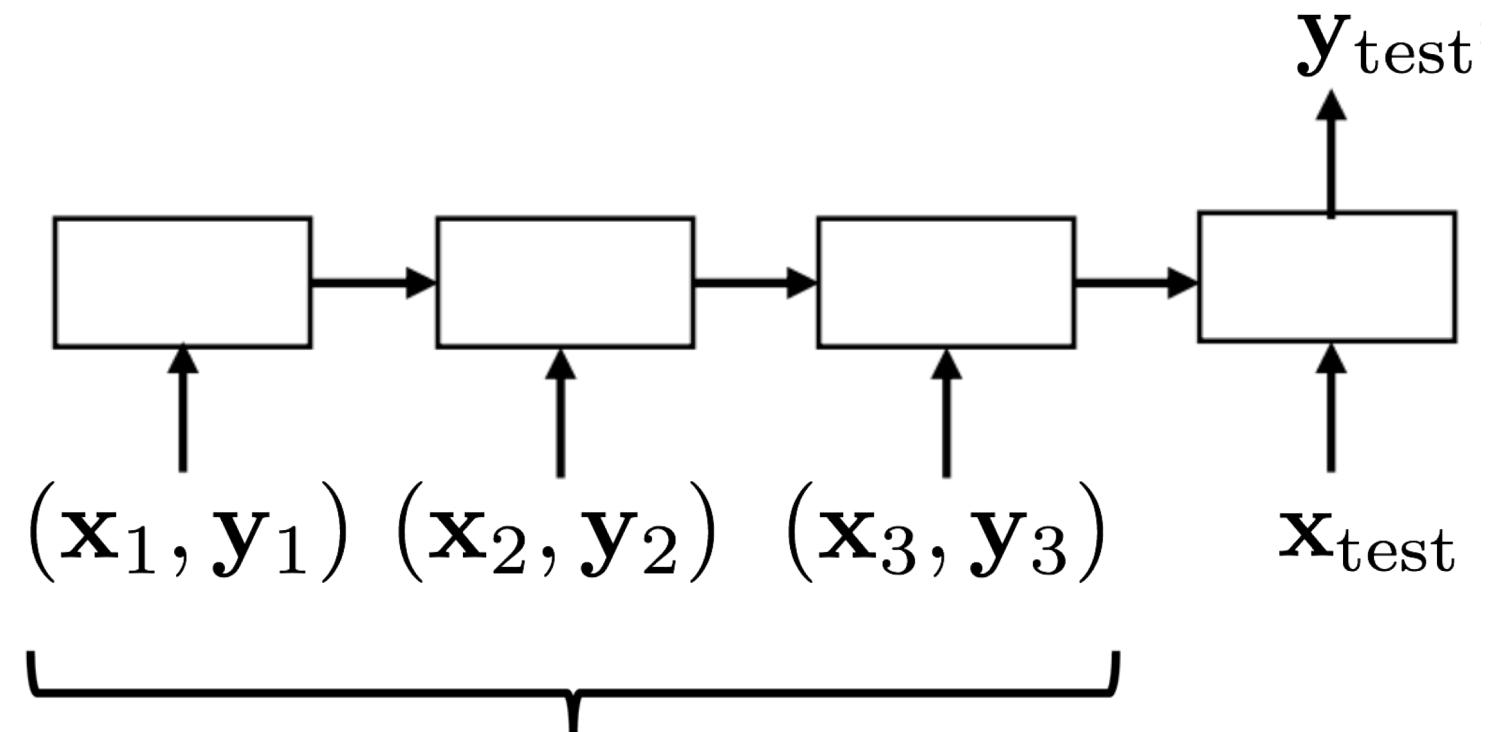
Boney & Ilin ICLR workshop track '18

Design of f ?

$$\mathcal{D}_{\text{train}} \ x_{\text{test}} \xrightarrow{\quad} y_{\text{test}}$$

Recurrent network

$$y_{\text{test}} = f(\mathcal{D}_{\text{train}}, x_{\text{test}}; \theta)$$



network implements the
“learned learning procedure”

Does it converge?

- Sort of?

What does it converge to?

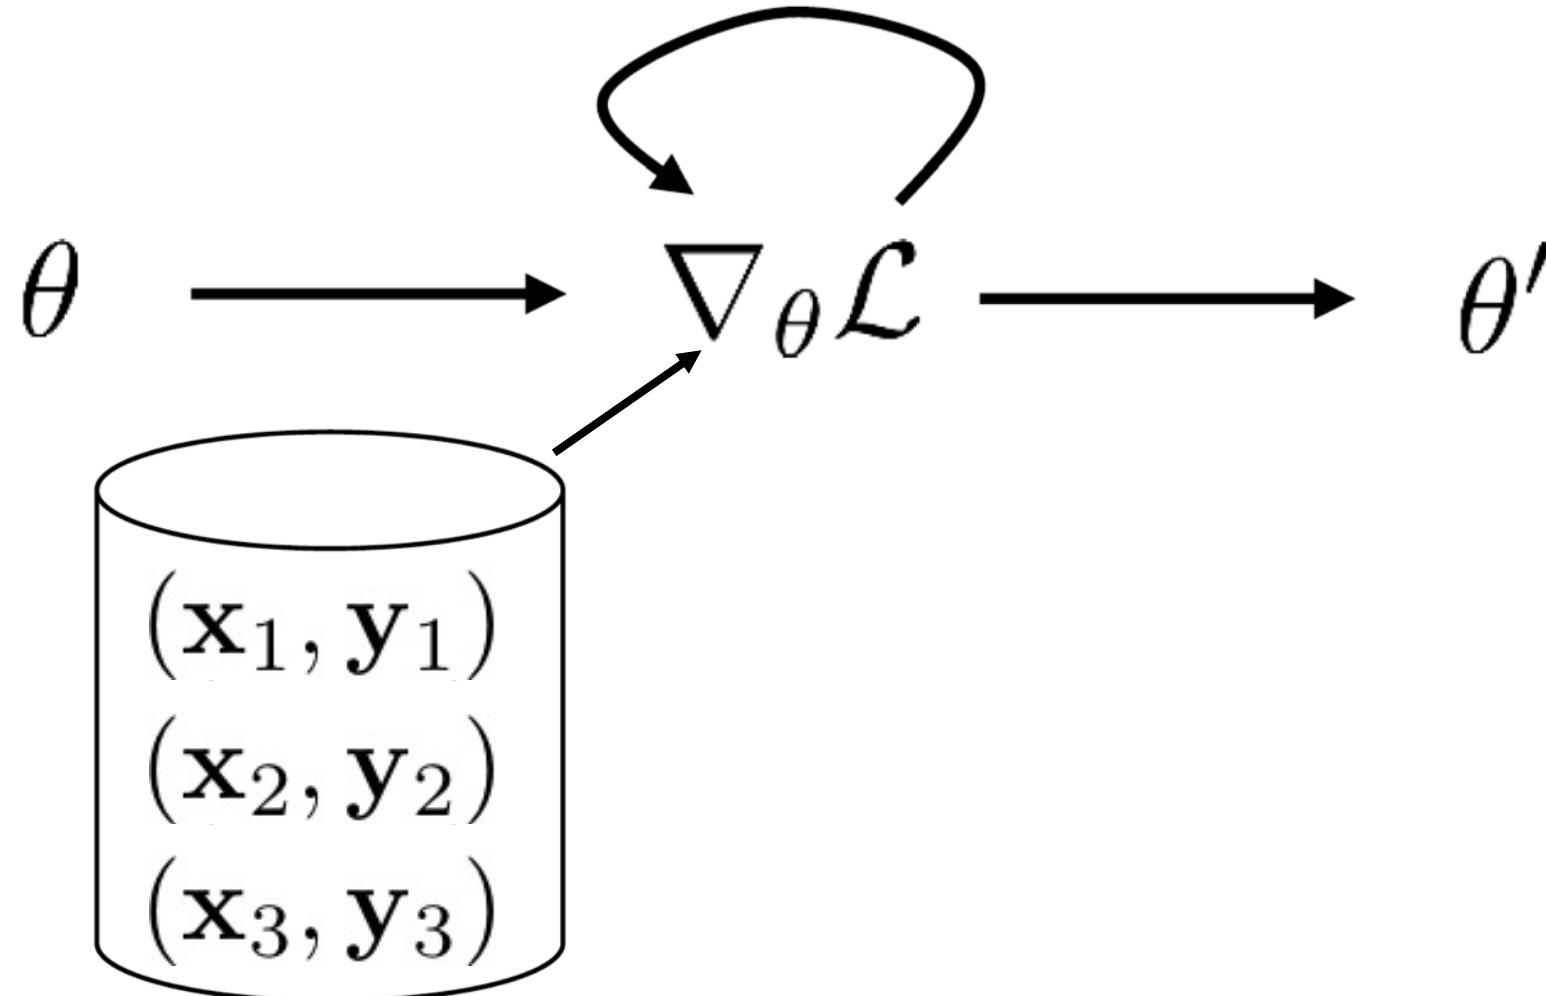
- Who knows...

What to do if not good enough?

- Nothing

MAML

$$y_{\text{test}} = f(x_{\text{test}}; \theta - \alpha \nabla_{\theta} \mathcal{L}(\mathcal{D}_{\text{train}}))$$



Does it converge?

- Yes (it's gradient descent...)

What does it converge to?

- A local optimum (it's gradient descent...)

What to do if not good enough?

- Keep taking gradient steps (it's gradient descent..)

Does this structure come at a cost?

Recurrent network

$$\mathbf{y}_{\text{test}} = f(\mathcal{D}_{\text{train}}, \mathbf{x}_{\text{test}}; \theta)$$

MAML

$$\mathbf{y}_{\text{test}} = f(\mathbf{x}_{\text{test}}; \theta - \alpha \nabla_{\theta} \mathcal{L}(\mathcal{D}_{\text{train}}))$$

Does this structure come at a cost?

For a sufficiently deep f ,

MAML function can approximate any function of $\mathcal{D}_{\text{train}}, \mathbf{x}_{\text{test}}$

Finn & Levine, ICLR 2018

Assumptions:

- nonzero α
- loss function gradient does not lose information about the label
- datapoints in $\mathcal{D}_{\text{train}}$ are unique

Why is this interesting?

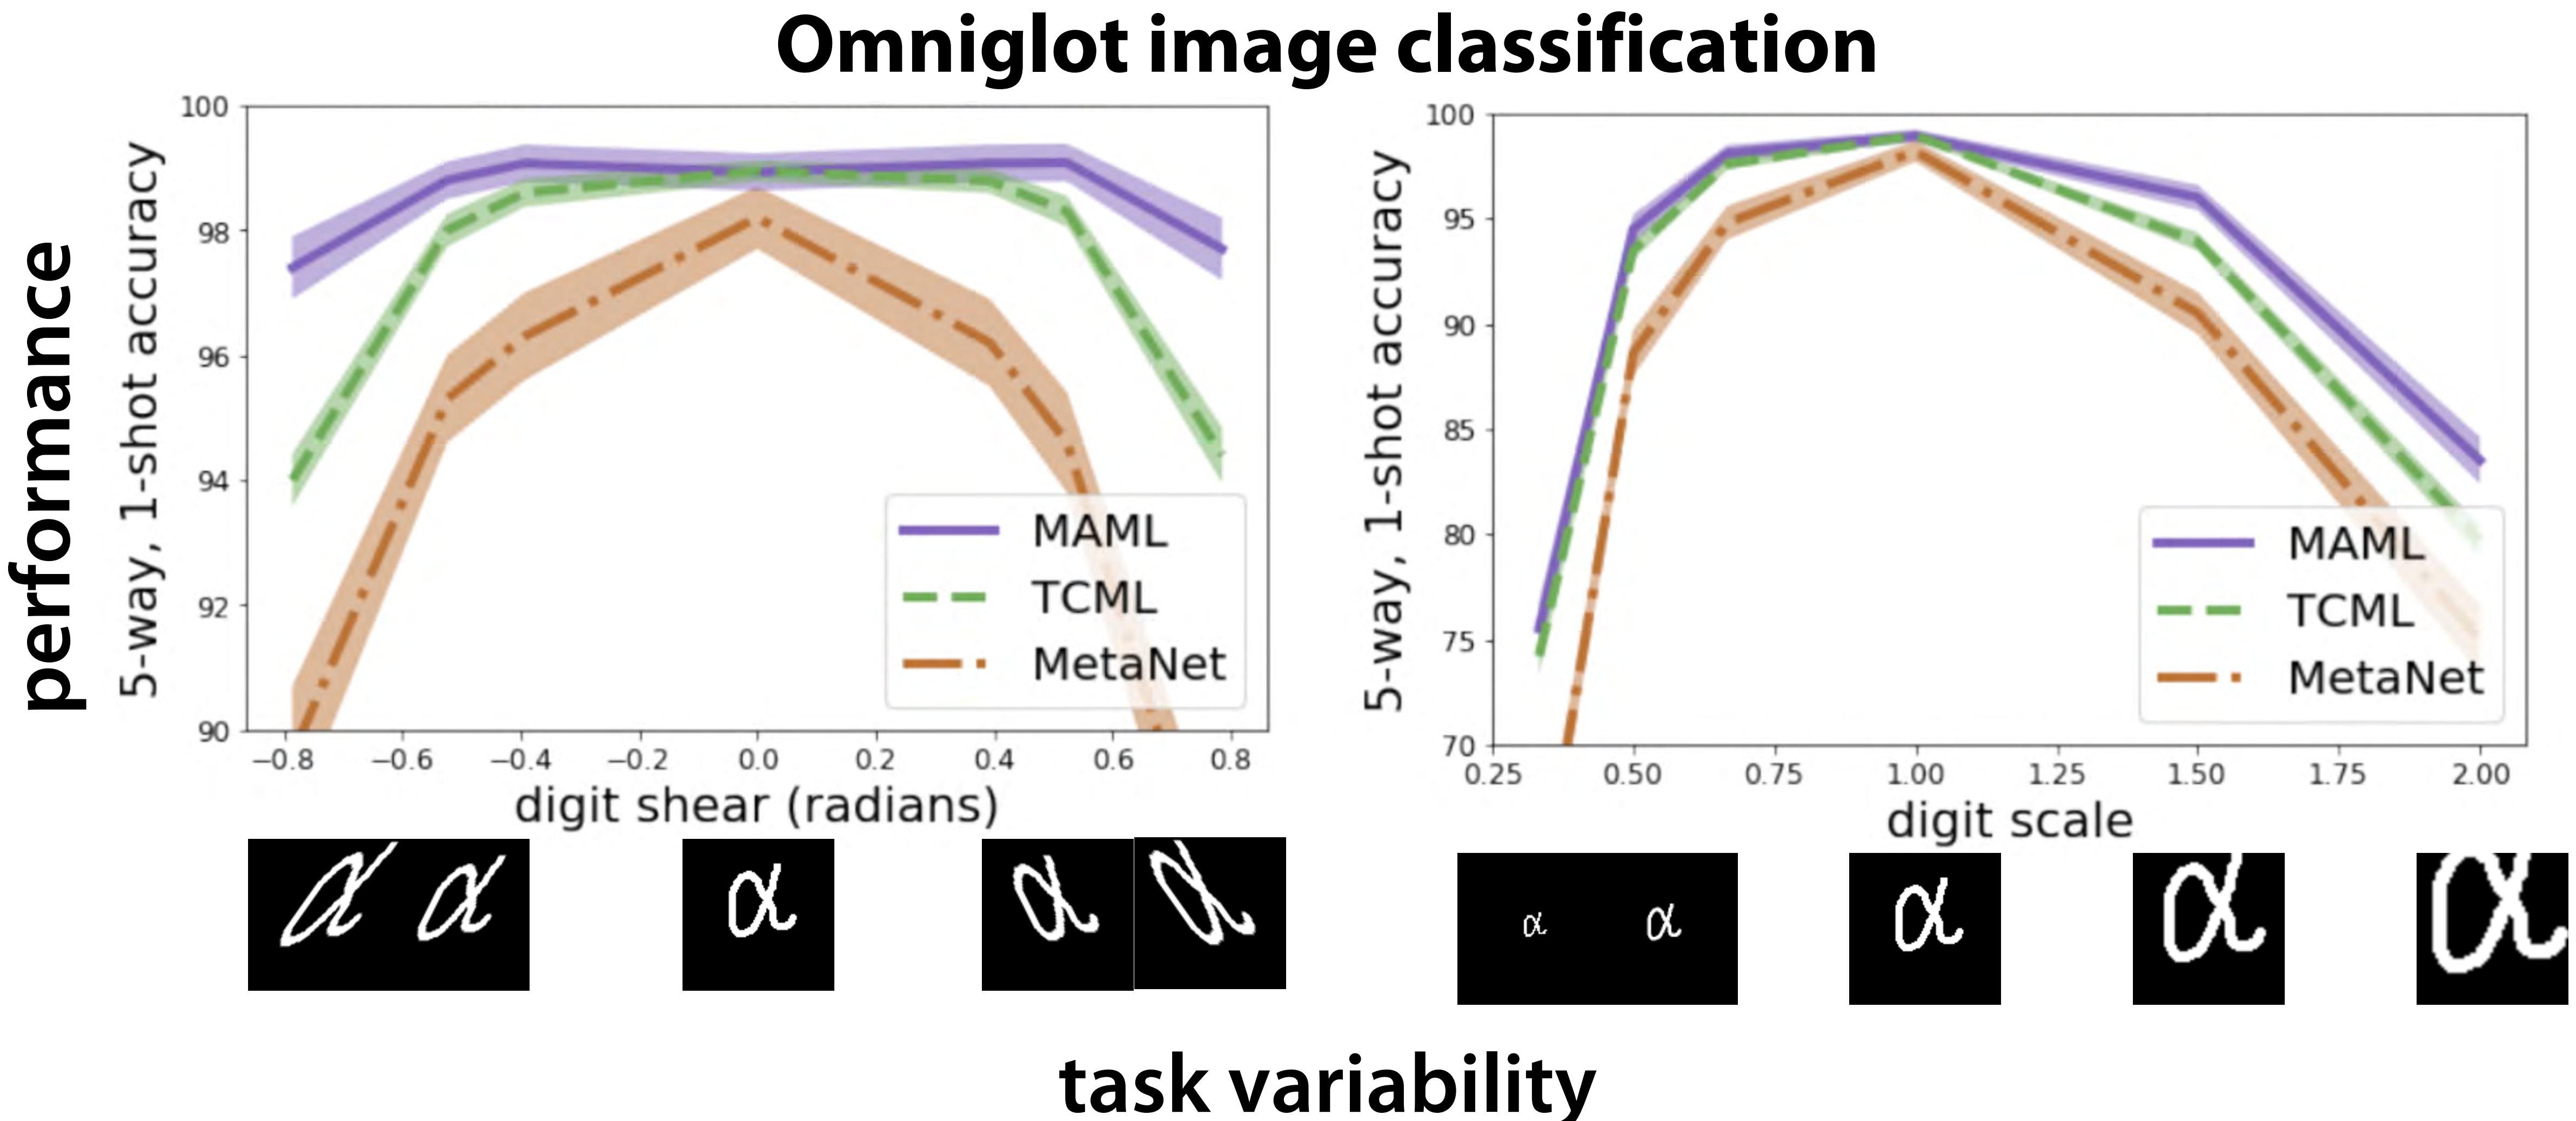
MAML has benefit of inductive bias without losing expressive power.

Is this structure useful?

How well can methods generalize to similar, but extrapolated tasks?

The world is non-stationary.

MAML **TCML**, **MetaNetworks**

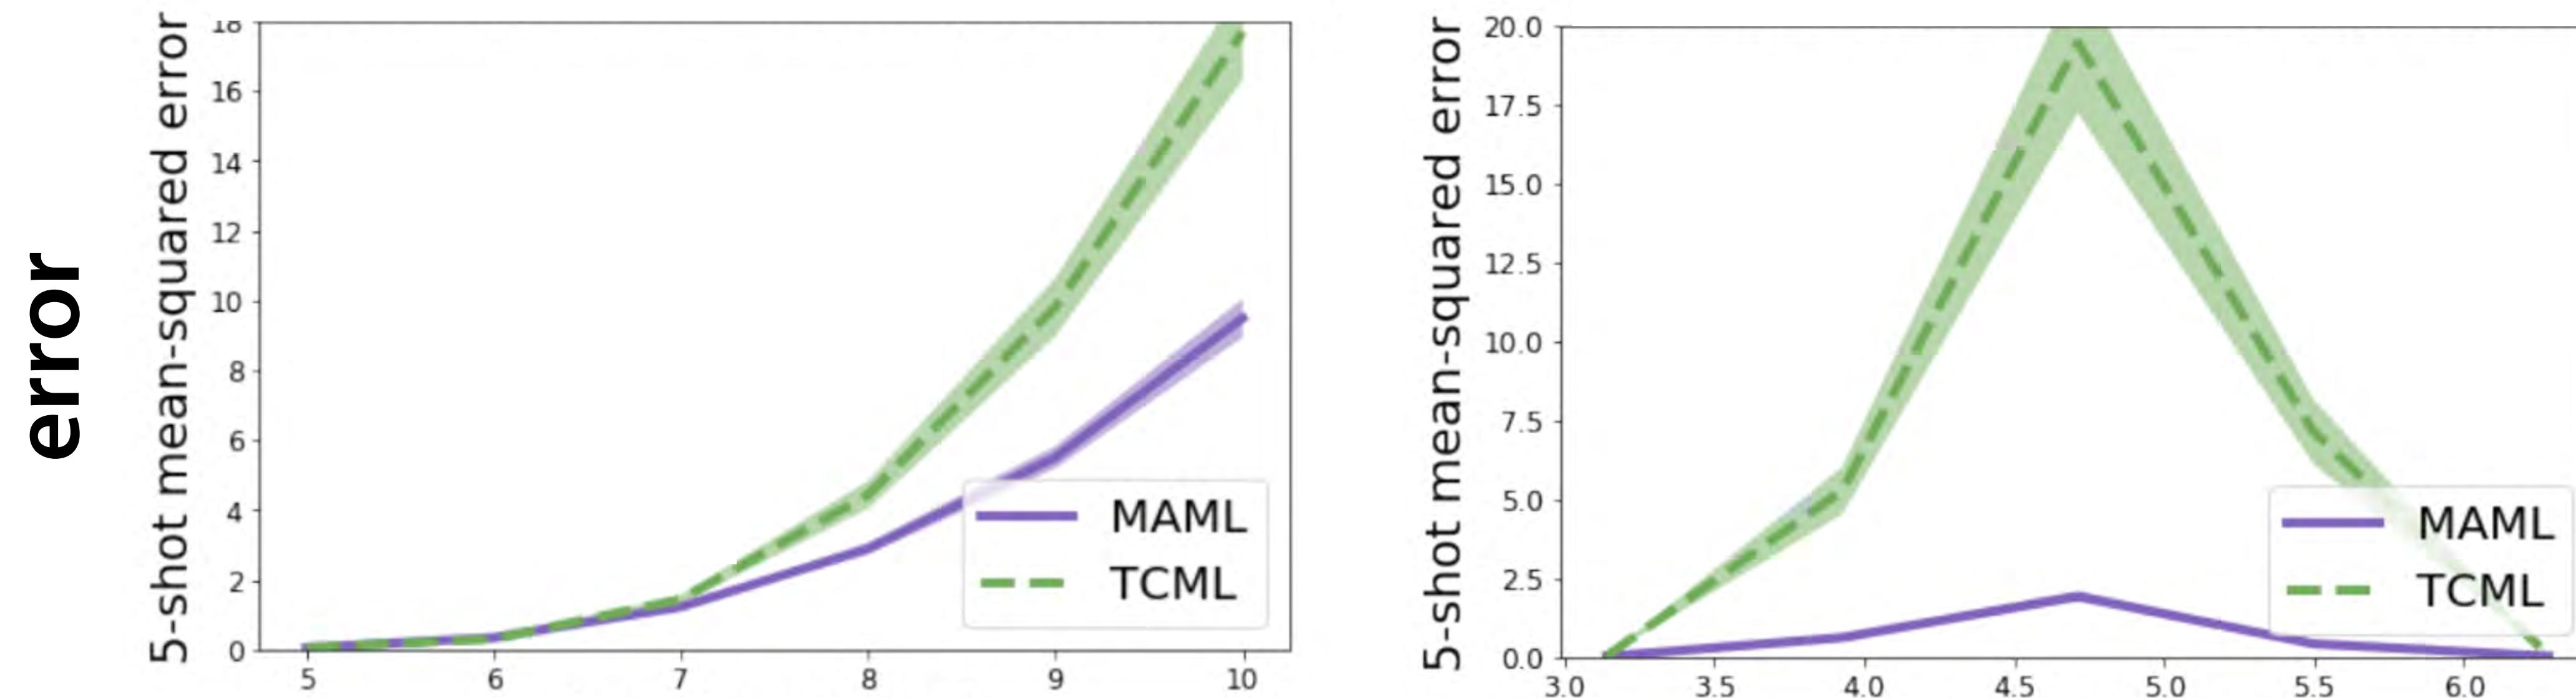


How well can methods generalize to similar, but extrapolated tasks?

The world is non-stationary.

MAML **TCML**

Sinusoid curve regression



Takeaway: Strategies learned with MAML consistently generalize better to out-of-distribution tasks

Outline

- Meta-Learning Problem Statement
- Model-Agnostic Meta-Learning (MAML)
- **Probabilistic Interpretation of MAML**
- Meta-Learning with Automated Task Proposals
- Extensions to Robot Imitation & Intent Inference

A probabilistic interpretation

that's nice but is it... *Bayesian?*

kind of... but it can be more Bayesian

A useful property:

start from $\phi = \theta$ and follow gradient of $\log p(\mathbf{Y}|\mathbf{X}, \phi)$ for K steps

this is equivalent to MAP on $p(\phi|\mathbf{X}, \mathbf{Y})$

for a prior $p(\phi|\theta) = \mathcal{N}(\theta, \Sigma)$

and for a *linear* model $E[\mathbf{Y}] = \mathbf{X}^T \phi$

(Santos, 1996)

A probabilistic interpretation

start from $\phi = \theta$ and follow gradient of $\log p(\mathbf{Y}|\mathbf{X}, \phi)$ for K steps

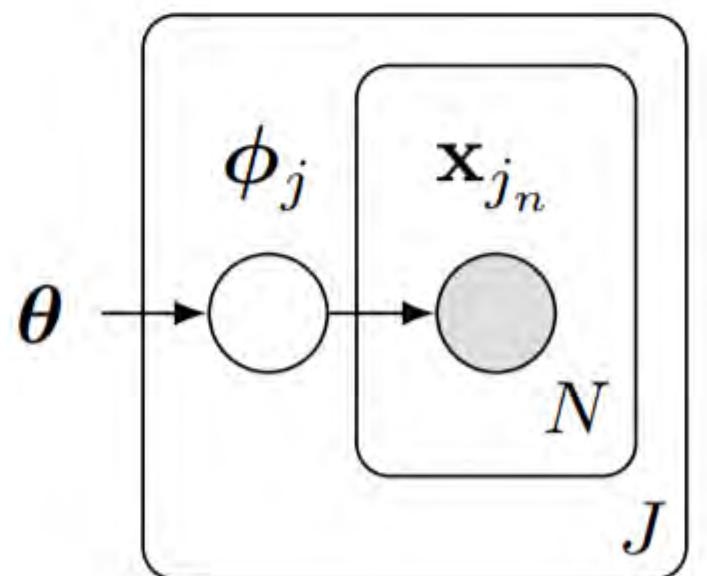
this is equivalent to MAP on $p(\phi|\mathbf{X}, \mathbf{Y})$

for a prior $p(\phi|\theta) = \mathcal{N}(\theta, \Sigma)$

and for a *linear* model $E[\mathbf{Y}] = \mathbf{X}^T \phi$

MAML adaptation:

$$\phi_i = \theta - \alpha \nabla_{\theta} \mathcal{L}(\theta, \mathcal{D}_{\text{train}}^i)$$



MAP inference in this model

$$p(\phi_i|\theta) = \mathcal{N}(\theta, \Sigma)$$

estimate Hessian for
neural nets with KFAC

can we do better than MAP?

can use Laplace estimate:

$$-\log p(\mathbf{X} | \boldsymbol{\theta}) \approx \sum_i \left[-\log p(\mathbf{x}_j | \hat{\phi}_j) - \log p(\hat{\phi}_j | \boldsymbol{\theta}) + \frac{1}{2} \log \det(\mathbf{H}_j) \right]$$

Modeling ambiguity



- ✓ Smiling,
- ✓ Wearing Hat,
- ✗ Young

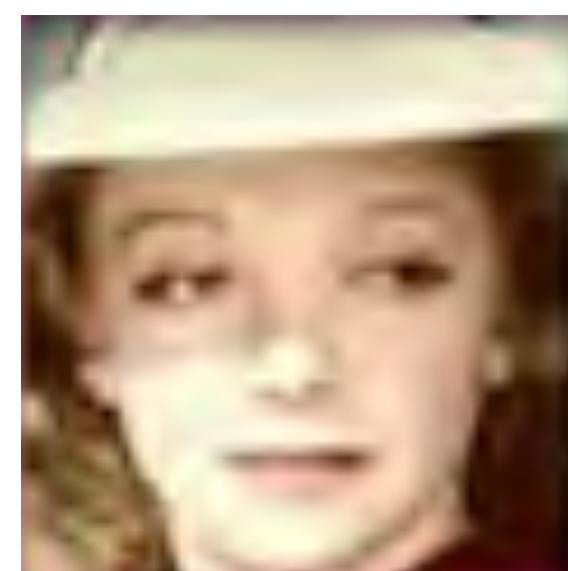


?

Can we learn to generate hypotheses about the underlying function?



- ✓ Smiling,
 - ✓ Wearing Hat,
 - ✓ Young
- ;
- ✗ Smiling,
 - ✓ Wearing Hat,
 - ✓ Young



?

Important for:

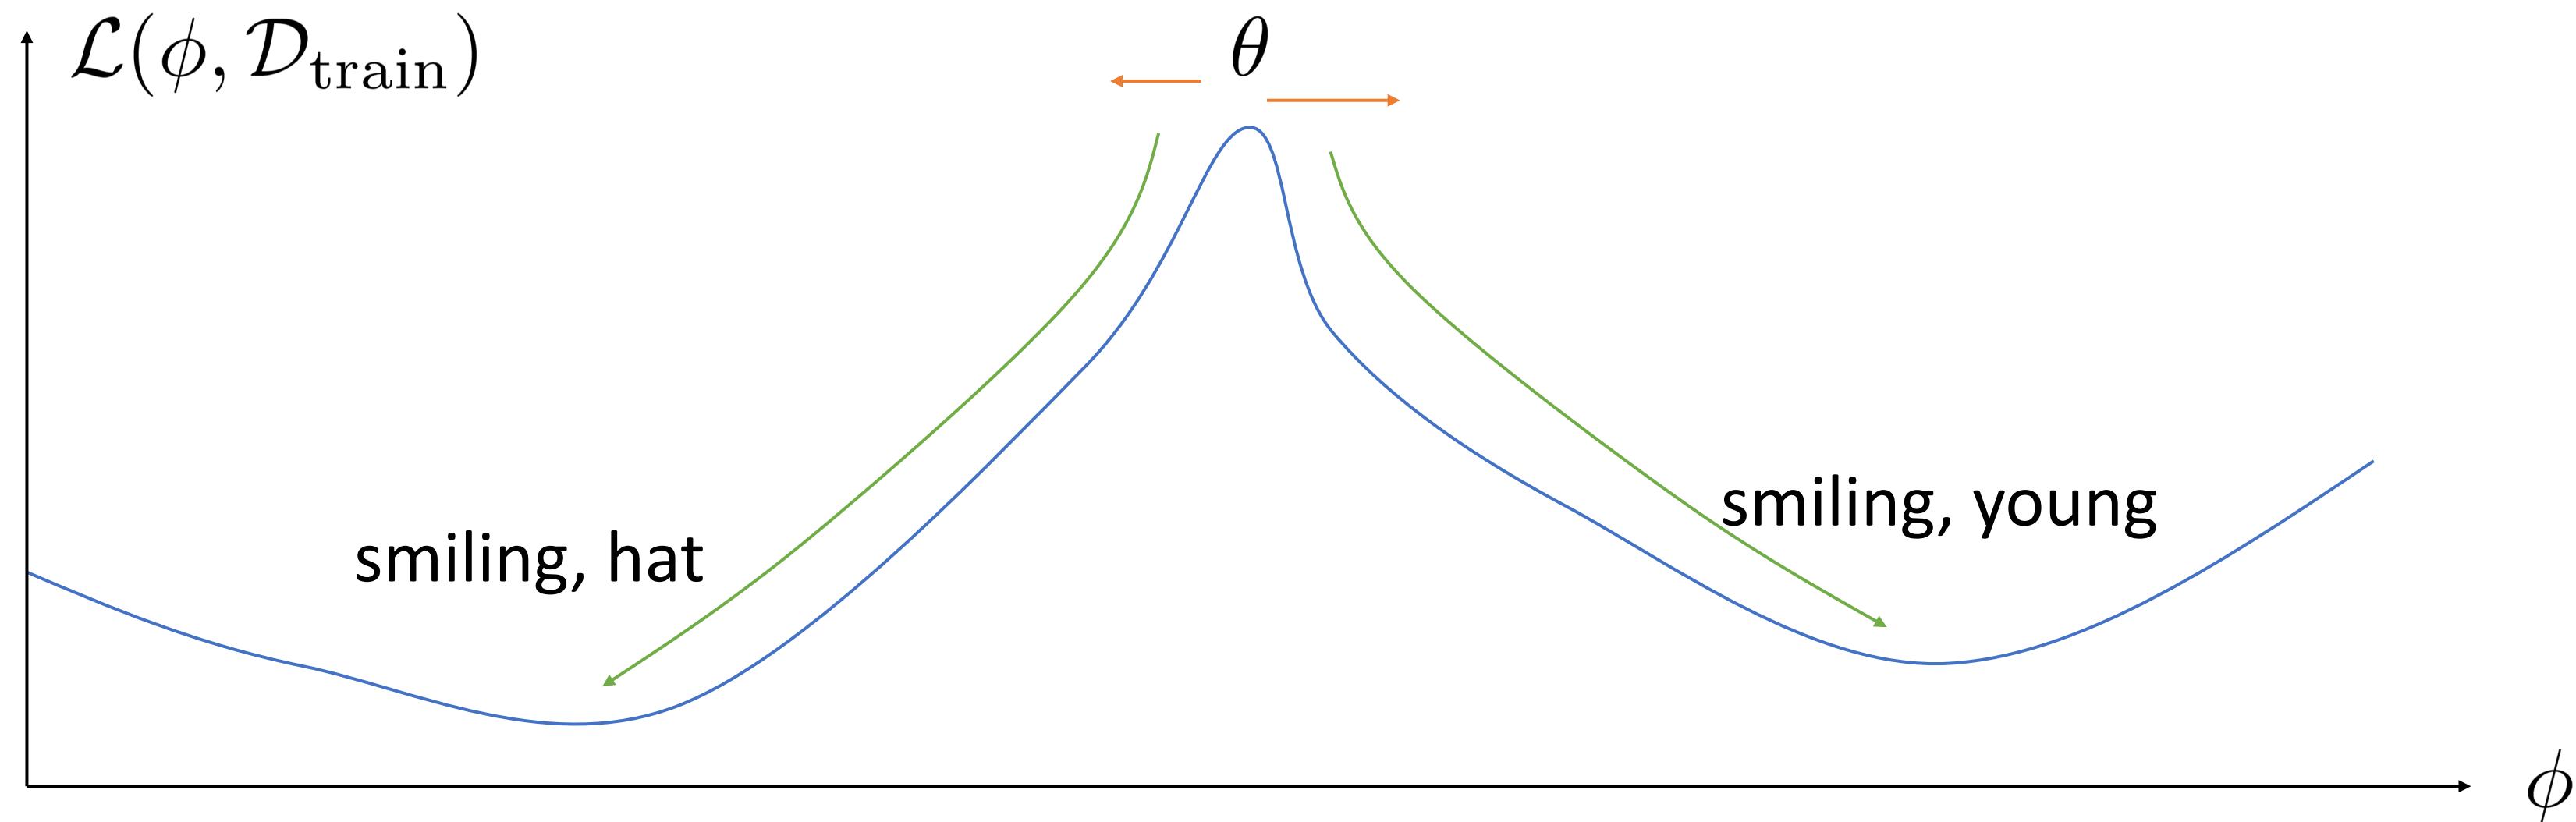
- learning to **actively learn**
- **safety-critical few-shot learning** (e.g. medical imaging)
- learning to **explore** in meta-RL

Modeling ambiguity



Can we *sample* classifiers?

Intuition: we want to learn a prior where a random *kick* can put us in different modes



- ✓ Smiling,
- ✓ Wearing Hat,
- ✓ Young

$$\phi \leftarrow \theta + \epsilon$$

$$\phi \leftarrow \phi + \alpha \nabla_{\phi} \mathcal{L}(\phi, \mathcal{D}_{\text{train}})$$

Meta-learning with ambiguity



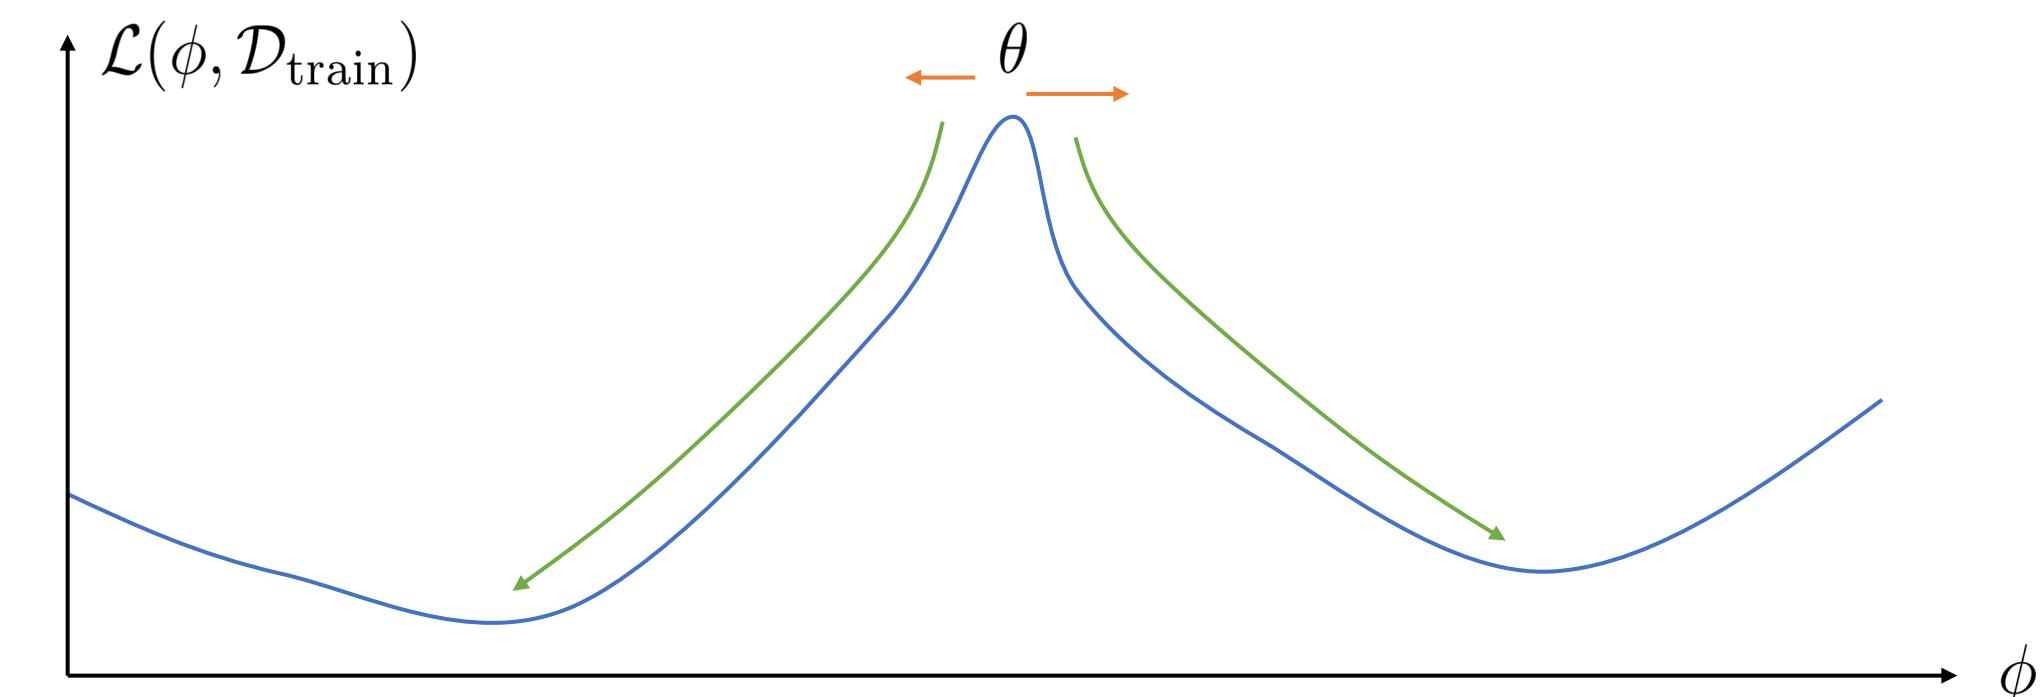
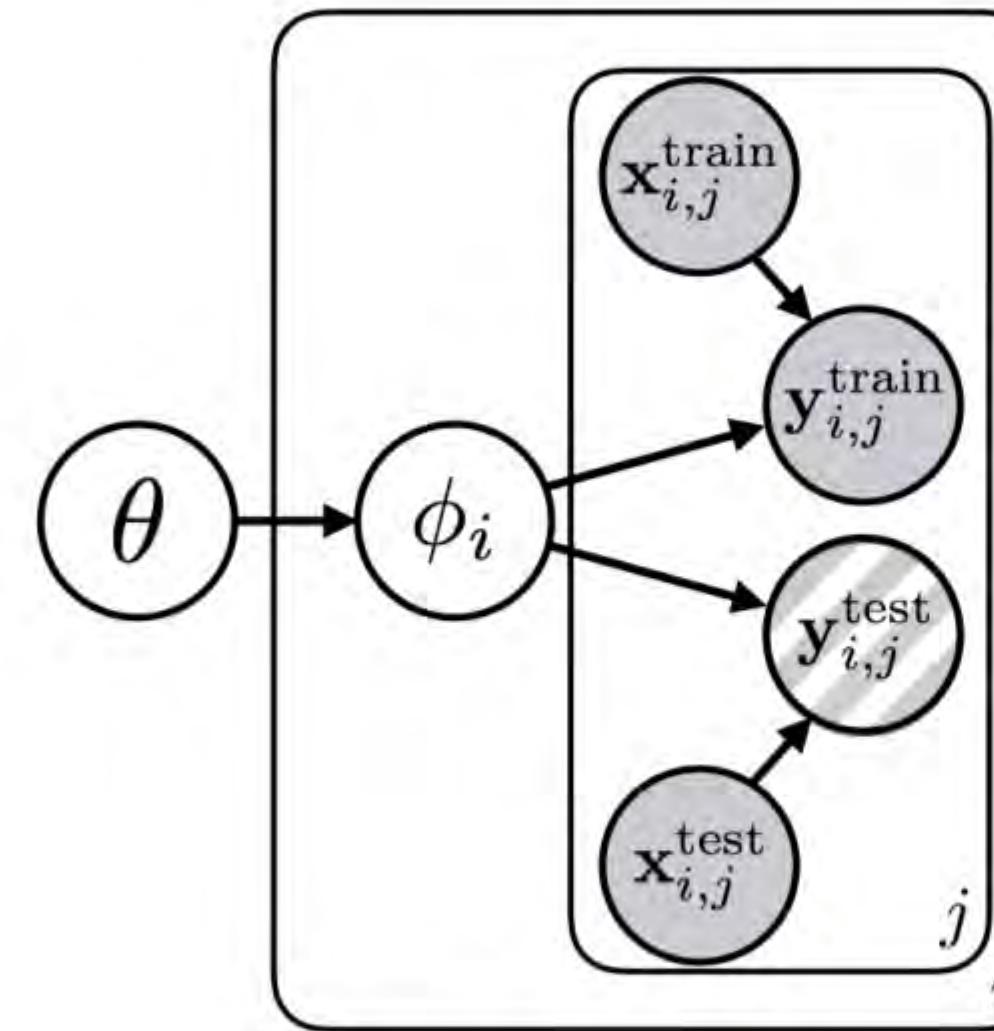
$$\theta \sim p(\theta) = \mathcal{N}(\mu_\theta, \Sigma_\theta)$$

$$\phi_i \sim p(\phi_i | \theta)$$

Goal: sample $\phi_i \sim p(\phi_i | x_i^{\text{train}}, y_i^{\text{train}}, \cancel{x_i^{\text{test}}})$

$$\log p(y_i^{\text{train}} | x_i^{\text{train}}, \phi_i)$$

$$\log p(y_i^{\text{test}} | x_i^{\text{test}}, \phi_i)$$



Sampling parameter vectors

$$\theta \sim p(\theta) = \mathcal{N}(\mu_\theta, \Sigma_\theta)$$

$$\log p(y_i^{\text{train}} | x_i^{\text{train}}, \phi_i)$$

$$\phi_i \sim p(\phi_i | \theta)$$

$$\log p(y_i^{\text{test}} | x_i^{\text{test}}, \phi_i)$$

Goal: sample $\phi_i \sim p(\phi_i | x_i^{\text{train}}, y_i^{\text{train}})$

$$p(\phi_i | x_i^{\text{train}}, y_i^{\text{train}}) \propto \int p(\theta)p(\phi_i | \theta)p(y_i^{\text{train}} | x_i^{\text{train}}, \phi_i)d\theta$$

⇒ this is completely intractable!

what if we knew $p(\phi_i | \theta, x_i^{\text{train}}, y_i^{\text{train}})$?

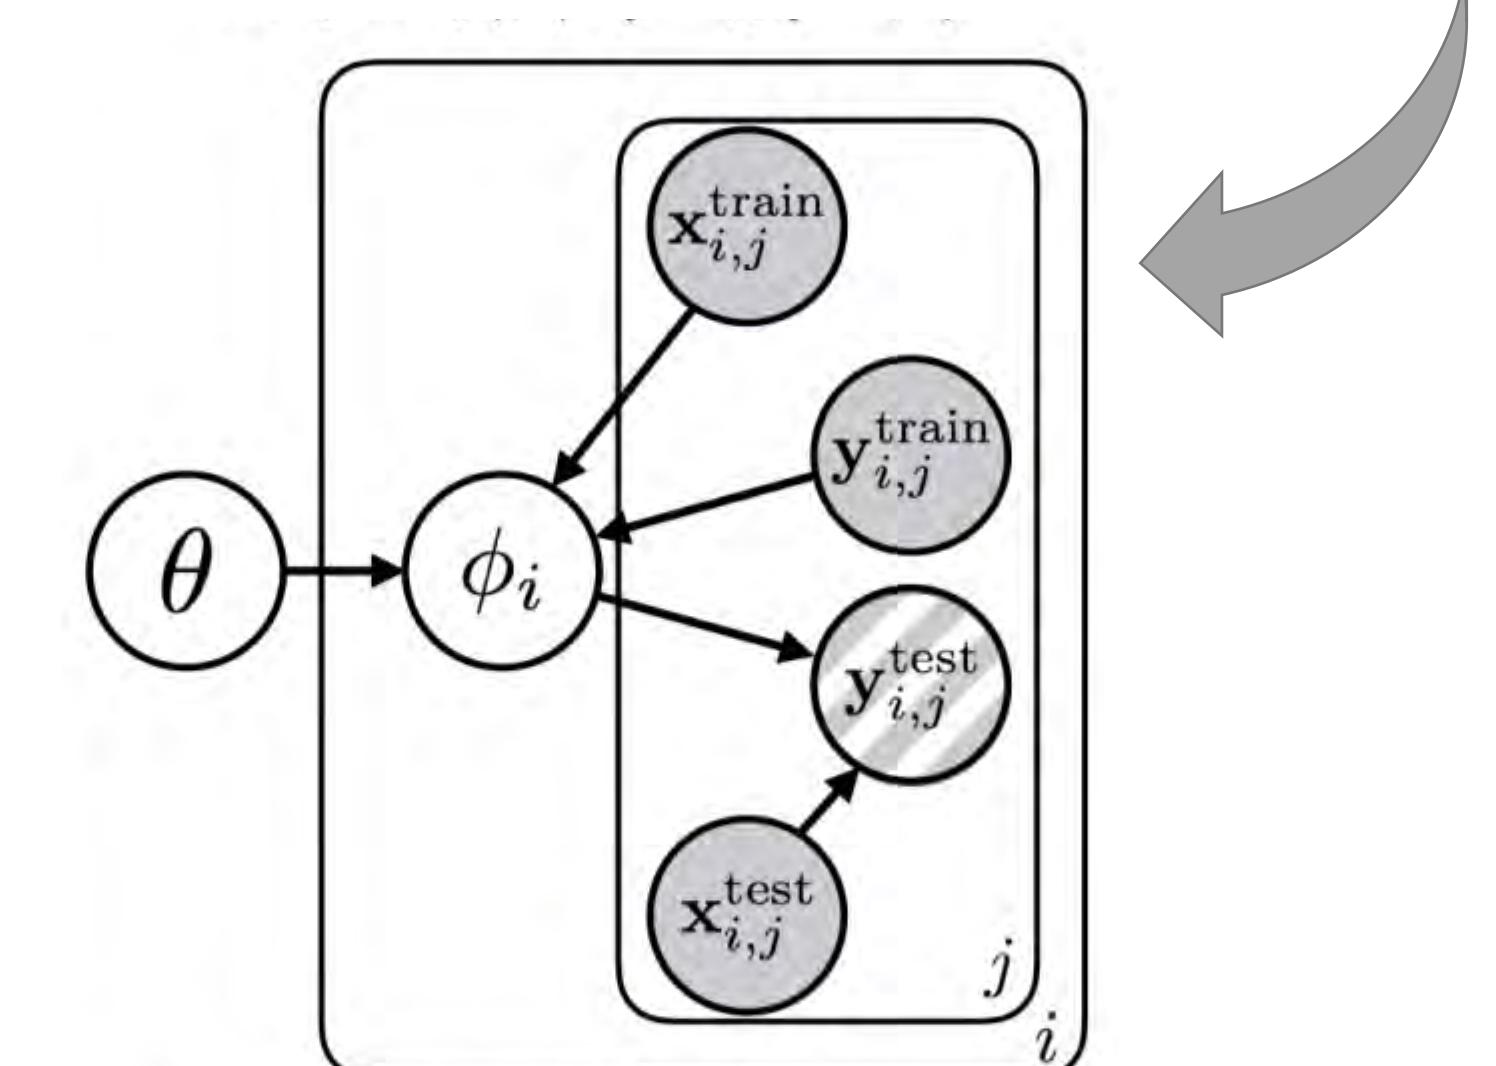
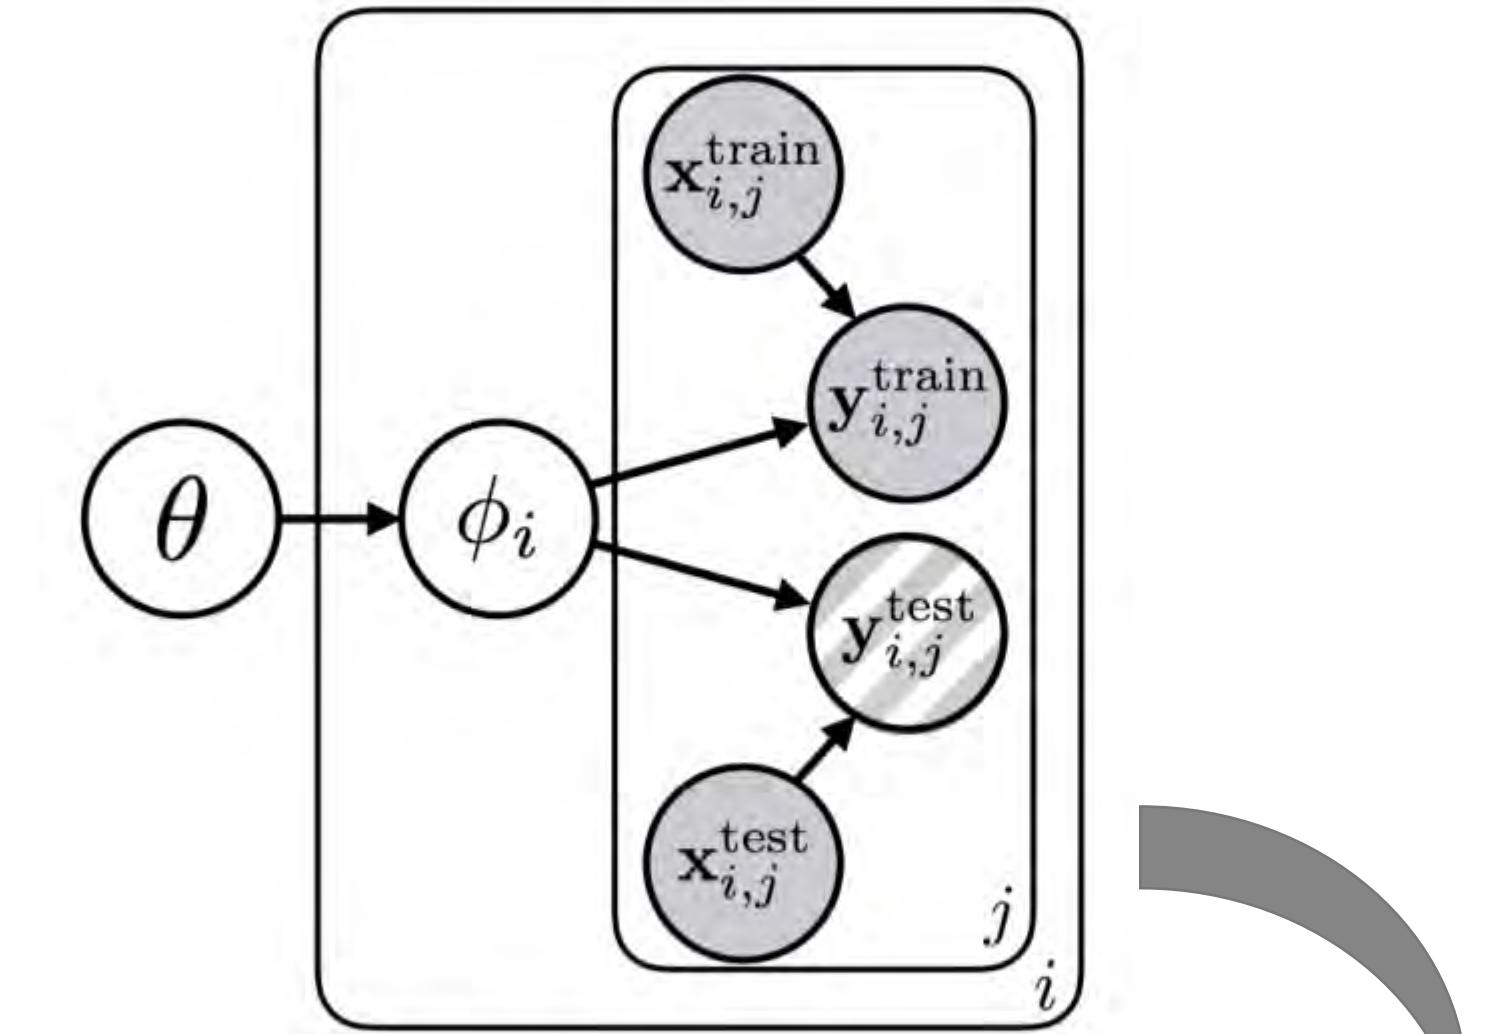
⇒ now sampling is easy! just use ancestral sampling!

key idea: $p(\phi_i | \theta, x_i^{\text{train}}, y_i^{\text{train}}) \approx \delta(\hat{\phi}_i)$

this is **extremely** crude

but **extremely** convenient!

$$\hat{\phi}_i \approx \theta + \alpha \nabla_\theta \log p(y_i^{\text{train}} | x_i^{\text{train}}, \theta)$$



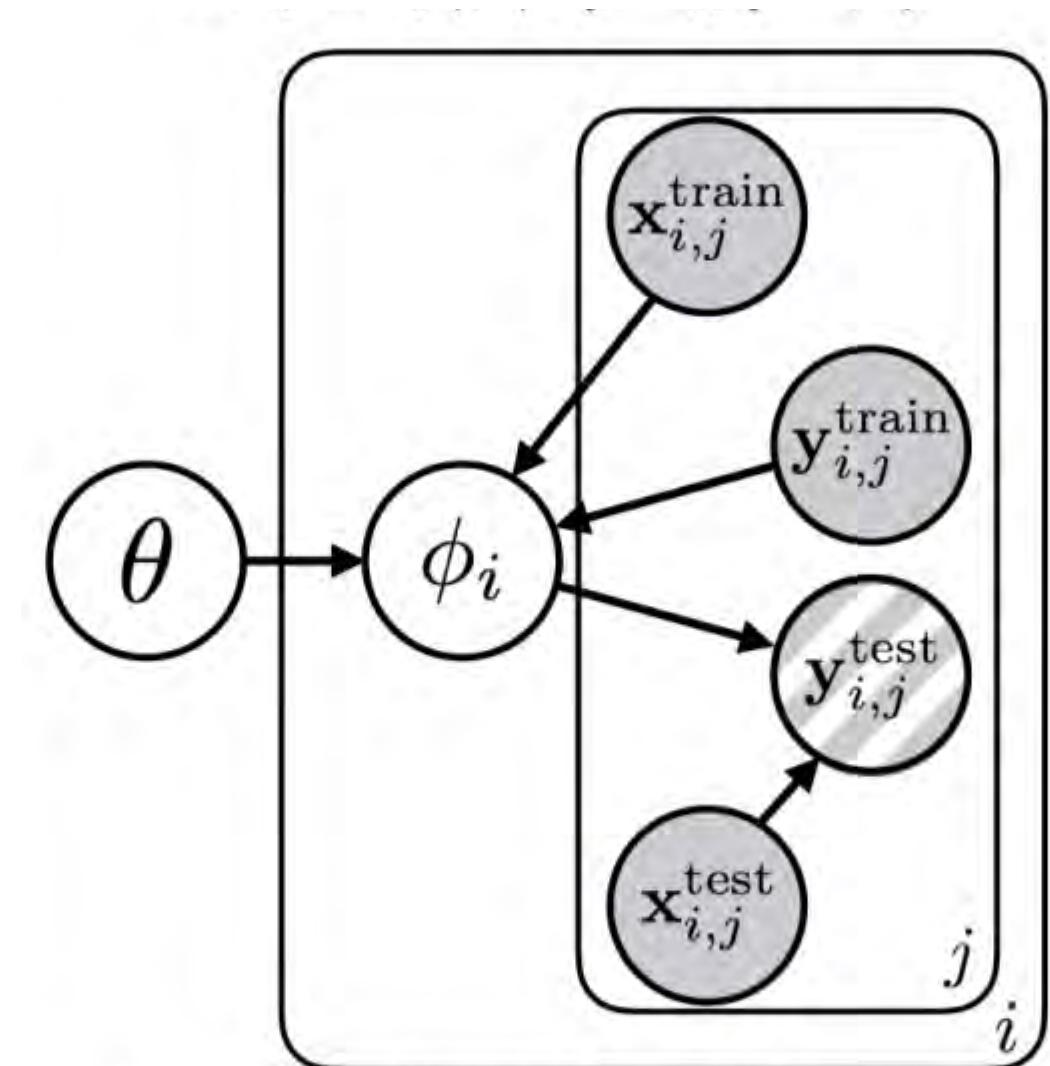
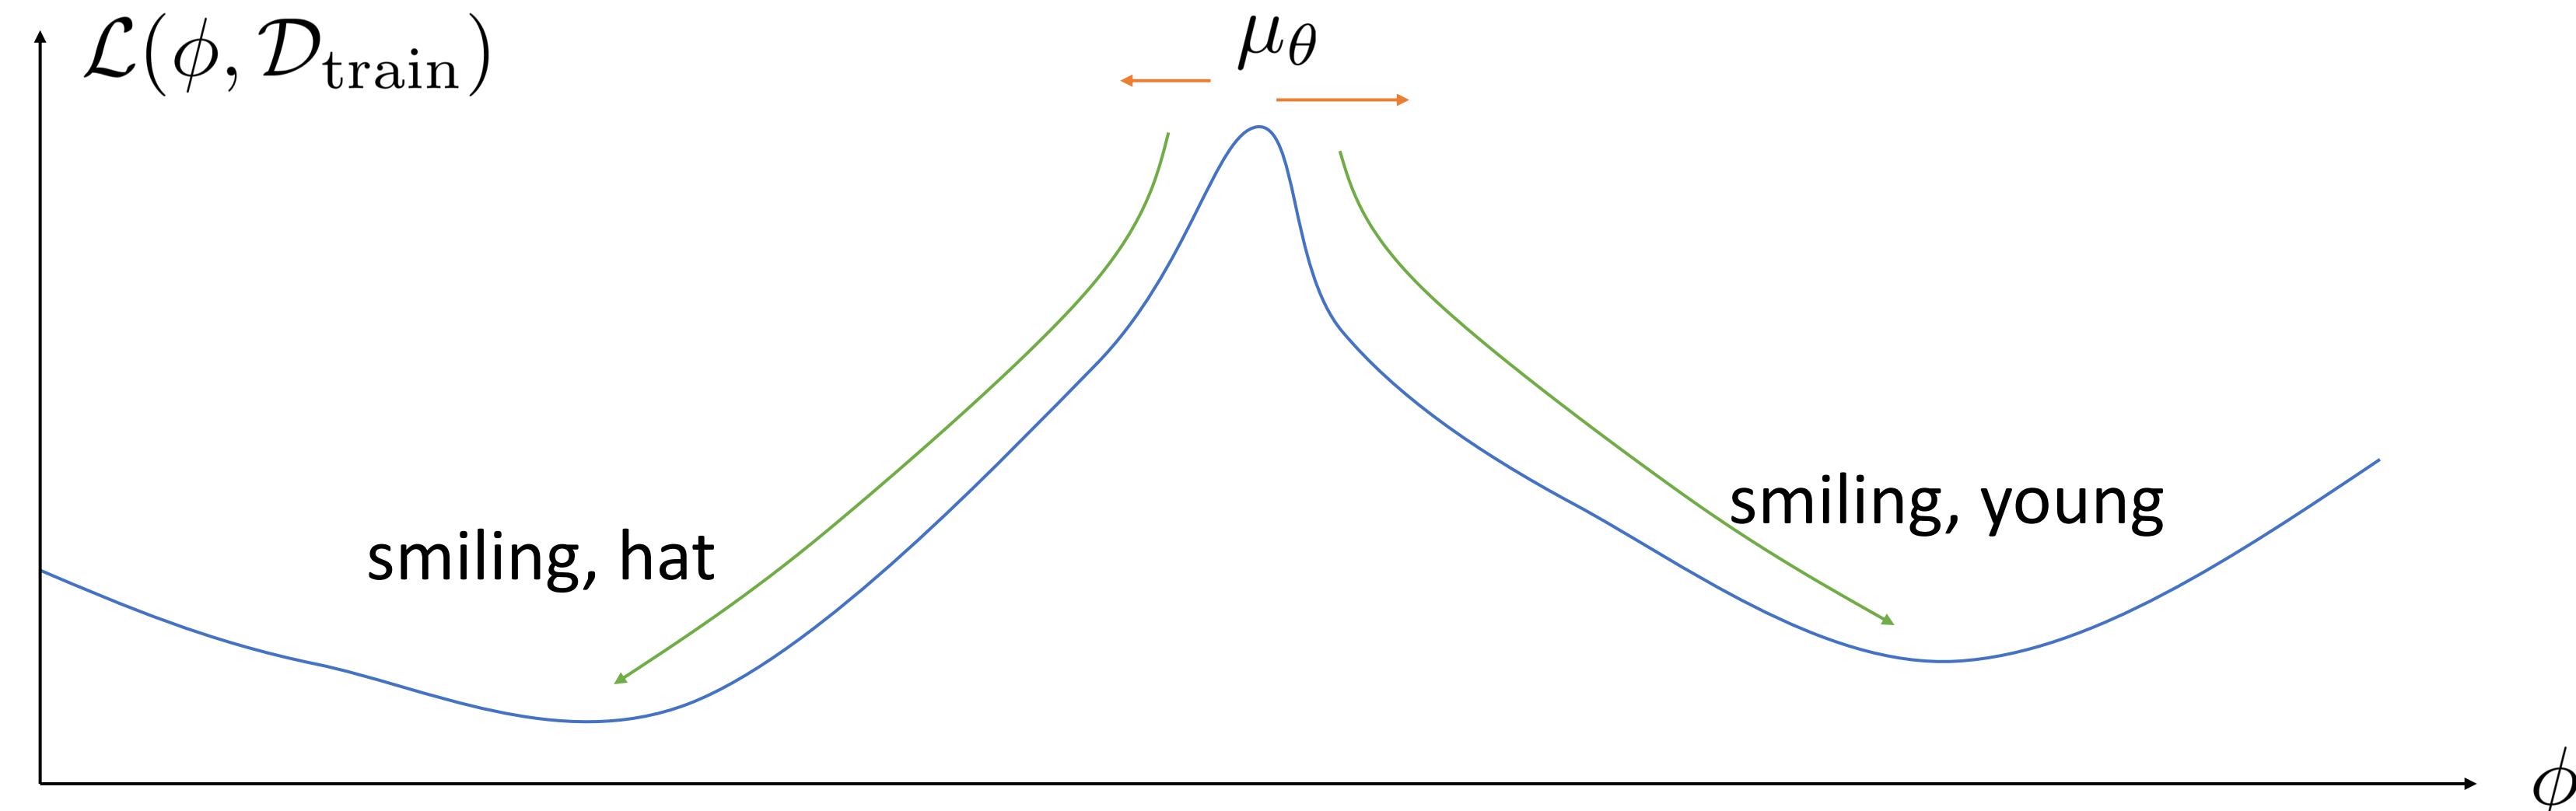
Sampling parameter vectors

$$\theta \sim p(\theta) = \mathcal{N}(\mu_\theta, \Sigma_\theta)$$

key idea: $p(\phi_i|\theta, x_i^{\text{train}}, y_i^{\text{train}}) \approx \delta(\hat{\phi}_i) \quad \hat{\phi}_i \approx \theta + \alpha \nabla_\theta \log p(y_i^{\text{train}}|x_i^{\text{train}}, \theta)$

What does ancestral sampling look like?

1. $\theta \sim \mathcal{N}(\mu_\theta, \Sigma_\theta)$
2. $\phi_i \sim p(\phi_i|\theta, x_i^{\text{train}}, y_i^{\text{train}}) \approx \hat{\phi}_i = \theta + \alpha \nabla_\theta \log p(y_i^{\text{train}}|x_i^{\text{train}}, \theta)$



Training the model – not so simple!

$$\theta \sim p(\theta) = \mathcal{N}(\mu_\theta, \Sigma_\theta)$$

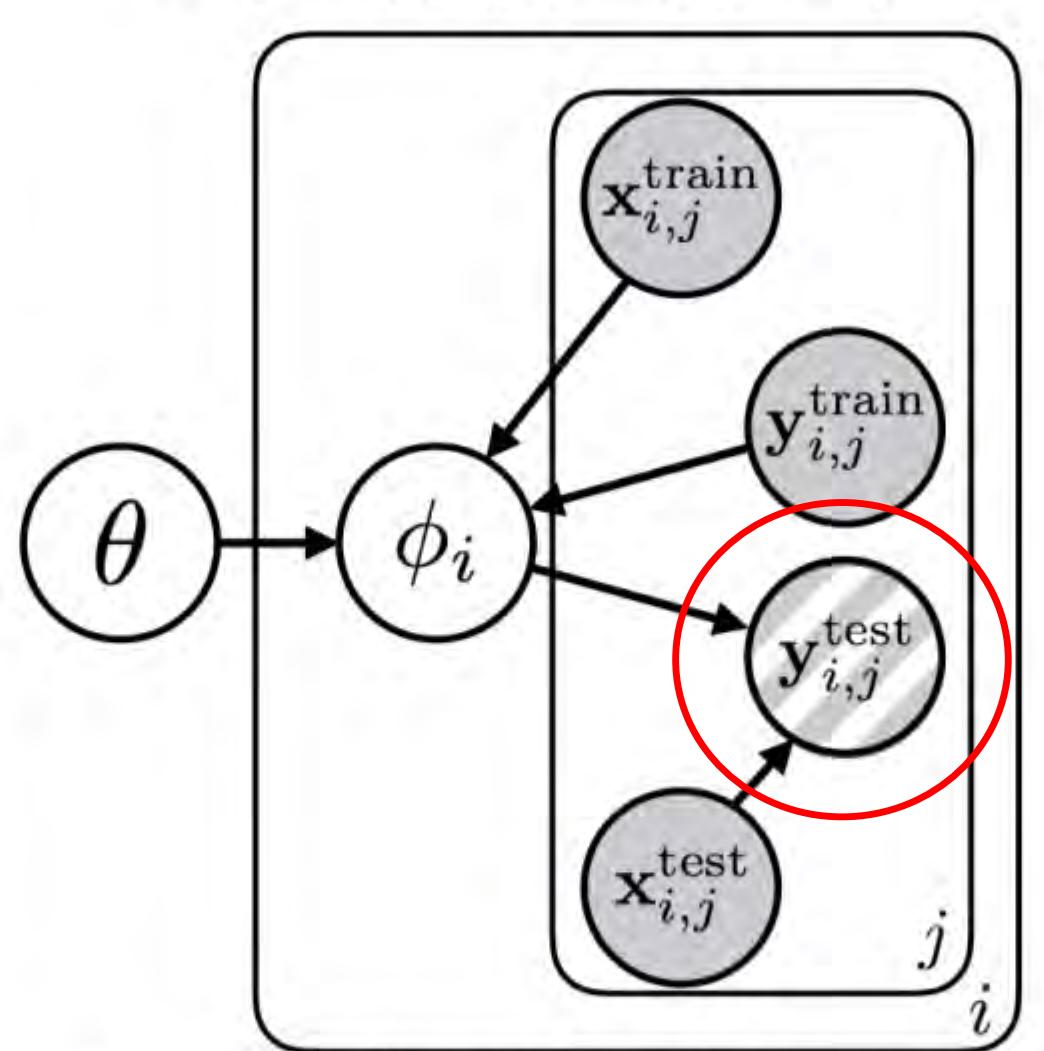
key idea: $p(\phi_i|\theta, x_i^{\text{train}}, y_i^{\text{train}}) \approx \delta(\hat{\phi}_i) \quad \hat{\phi}_i \approx \theta + \alpha \nabla_\theta \log p(y_i^{\text{train}}|x_i^{\text{train}}, \theta)$

Can't do ancestral sampling anymore!

need to figure out $p(\theta|x_i^{\text{test}}, y_i^{\text{test}})$ (intractable inference problem)

approximate using $q(\theta|x_i^{\text{test}}, y_i^{\text{test}})$, train via variational bound

$$\begin{aligned} & \log p(x_i^{\text{train}}, y_i^{\text{train}}, x_i^{\text{test}}, y_i^{\text{test}} | \mu_\theta, \Sigma_\theta) \\ & \leq E_{\theta \sim q} \left[\log \int p(\theta) p(\phi_i | \theta) p(y_i^{\text{train}} | x_i^{\text{train}}, \phi_i) p(y_i^{\text{test}} | x_i^{\text{test}}, \phi_i) d\phi_i - \log q(\theta | x_i^{\text{test}}, y_i^{\text{test}}) \right] \\ & \approx E_{\theta \sim q} \left[\log p(\theta) + \log p(y_i^{\text{test}} | x_i^{\text{test}}, \hat{\phi}_i(\theta)) - \log q(\theta | x_i^{\text{test}}, y_i^{\text{test}}) \right] \\ & = E_{\theta \sim q} \left[\log p(y_i^{\text{test}} | x_i^{\text{test}}, \hat{\phi}_i(\theta)) \right] - D_{\text{KL}}(q(\theta | x_i^{\text{test}}, y_i^{\text{test}}) \| p(\theta)) \end{aligned}$$



Training the model with amortized inference

need to figure out $p(\theta|x_i^{\text{test}}, y_i^{\text{test}})$ (intractable inference problem)

approximate using $q(\theta|x_i^{\text{test}}, y_i^{\text{test}})$, train via variational bound

$$E_{\theta \sim q} \left[\log p(y_i^{\text{test}}|x_i^{\text{test}}, \hat{\phi}_i(\theta)) \right] - D_{\text{KL}}(q(\theta|x_i^{\text{test}}, y_i^{\text{test}}) \| p(\theta))$$

how do we represent q ?

big neural network that outputs parameters? too intractable...

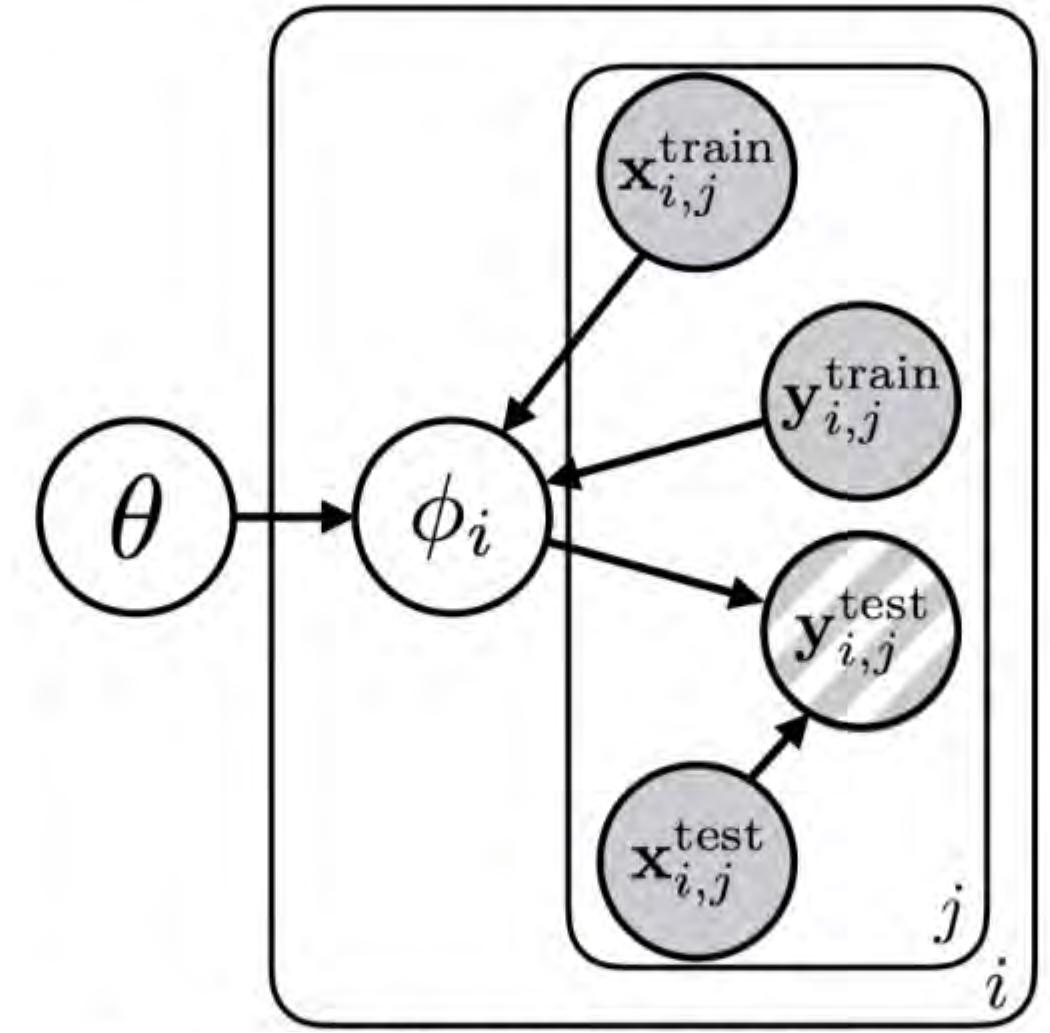
idea: $q(\theta|x_i^{\text{test}}, y_i^{\text{test}}) = \mathcal{N}(\mu_\theta + \alpha \nabla_{\mu_\theta} \log p(y_i^{\text{test}}|x_i^{\text{test}}, \mu_\theta), \Sigma_q)$

intuitive interpretation: *cheat* during meta-training by peeking at test set

...but minimize divergence against the prior

$$1. \theta \leftarrow \mu_\theta + \alpha \nabla_{\mu_\theta} \log p(y_i^{\text{test}}|x_i^{\text{test}}, \mu_\theta) + \epsilon$$

$$2. \phi_i \sim p(\phi_i|\theta, x_i^{\text{train}}, y_i^{\text{train}}) \approx \hat{\phi}_i = \theta + \alpha \nabla_\theta \log p(y_i^{\text{train}}|x_i^{\text{train}}, \theta)$$



Summary

need to figure out $p(\theta|x_i^{\text{test}}, y_i^{\text{test}})$ (intractable inference problem)

approximate using $q(\theta|x_i^{\text{test}}, y_i^{\text{test}})$, train via variational bound

$$E_{\theta \sim q} \left[\log p(y_i^{\text{test}}|x_i^{\text{test}}, \hat{\phi}_i(\theta)) \right] - D_{\text{KL}}(q(\theta|x_i^{\text{test}}, y_i^{\text{test}}) \| p(\theta))$$

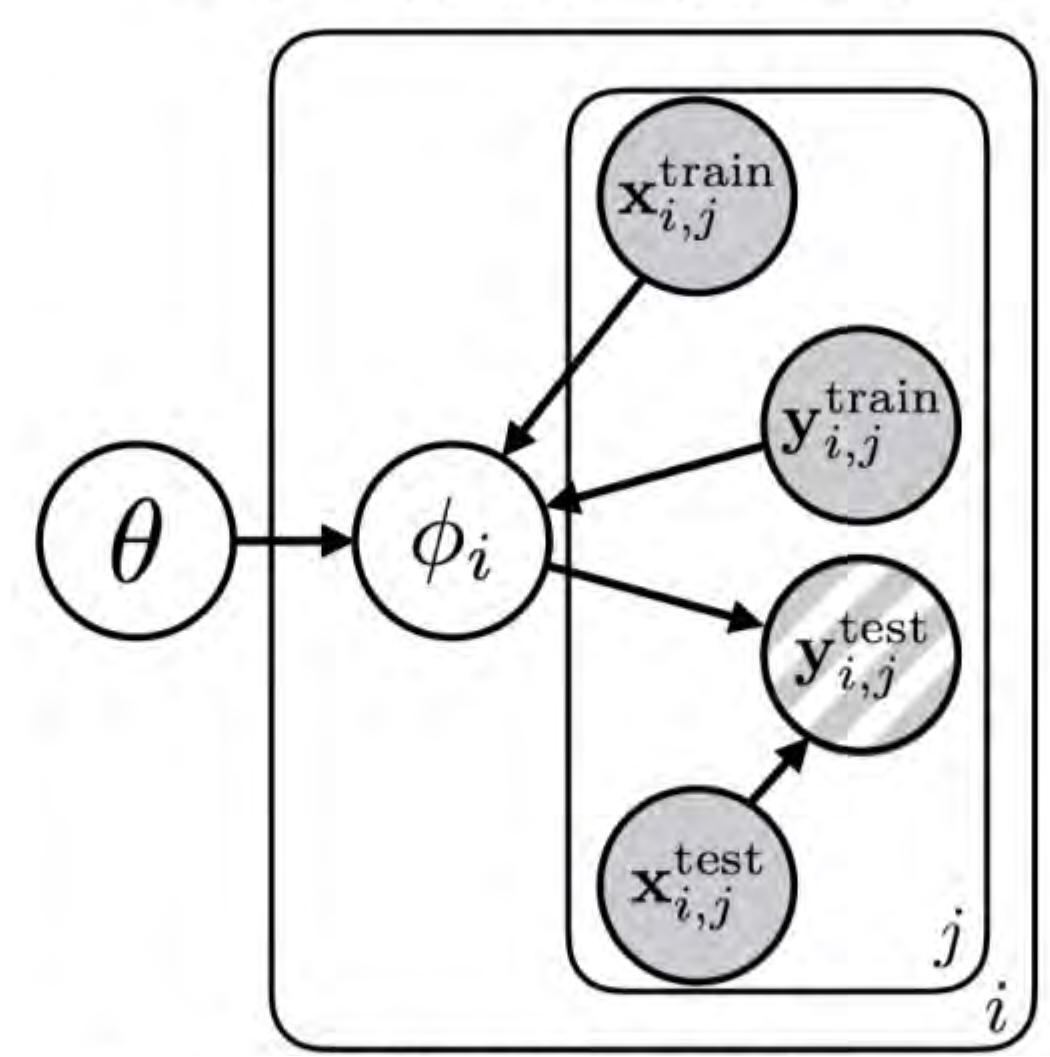
During meta-training:

1. $\theta \leftarrow \mu_\theta + \alpha \nabla_{\mu_\theta} \log p(y_i^{\text{test}}|x_i^{\text{test}}, \mu_\theta) + \epsilon$
2. $\phi_i \sim p(\phi_i|\theta, x_i^{\text{train}}, y_i^{\text{train}}) \approx \hat{\phi}_i = \theta + \alpha \nabla_\theta \log p(y_i^{\text{train}}|x_i^{\text{train}}, \theta)$

take gradient step on variational bound w.r.t. θ , Σ , and Σ_q

During meta-testing:

1. $\theta \sim \mathcal{N}(\mu_\theta, \Sigma_\theta)$
2. $\phi_i \sim p(\phi_i|\theta, x_i^{\text{train}}, y_i^{\text{train}}) \approx \hat{\phi}_i = \theta + \alpha \nabla_\theta \log p(y_i^{\text{train}}|x_i^{\text{train}}, \theta)$

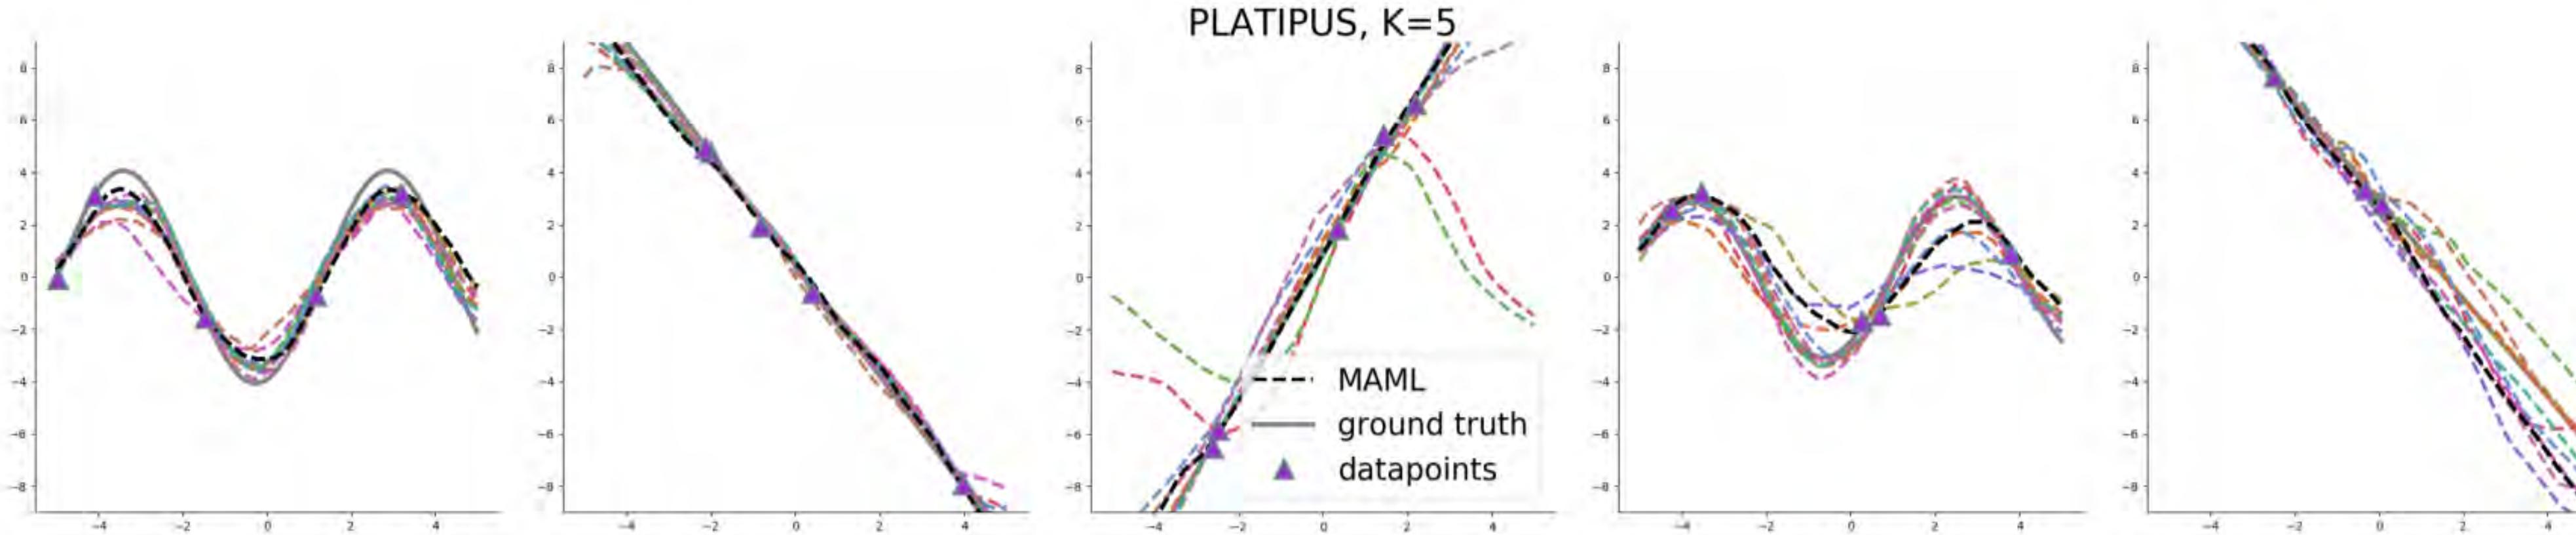


PLATIPUS

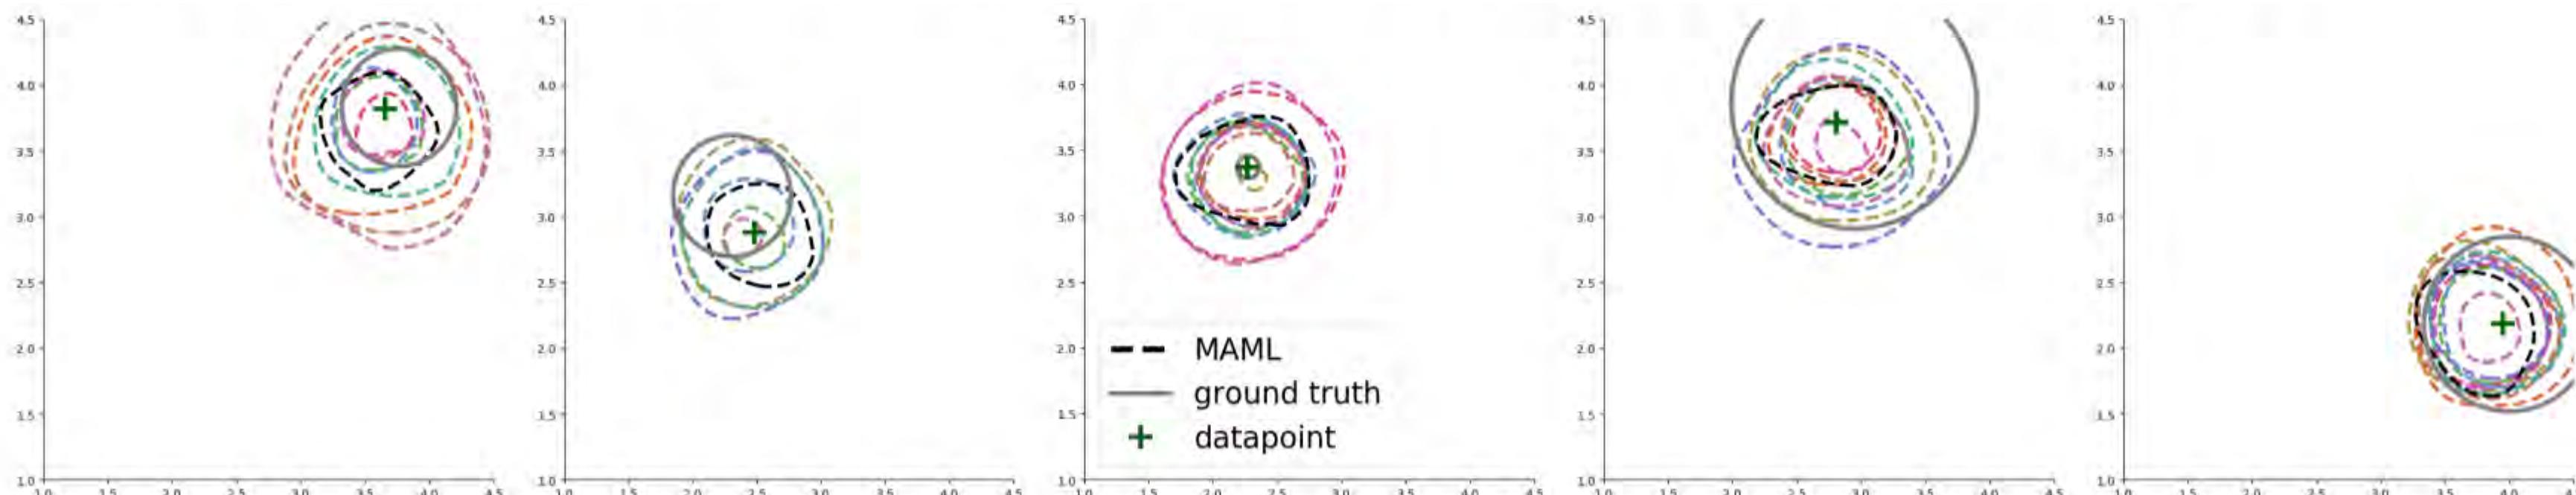
Probabilistic LATent model for Incorporating Priors and Uncertainty in few-Shot learning



Ambiguous regression:

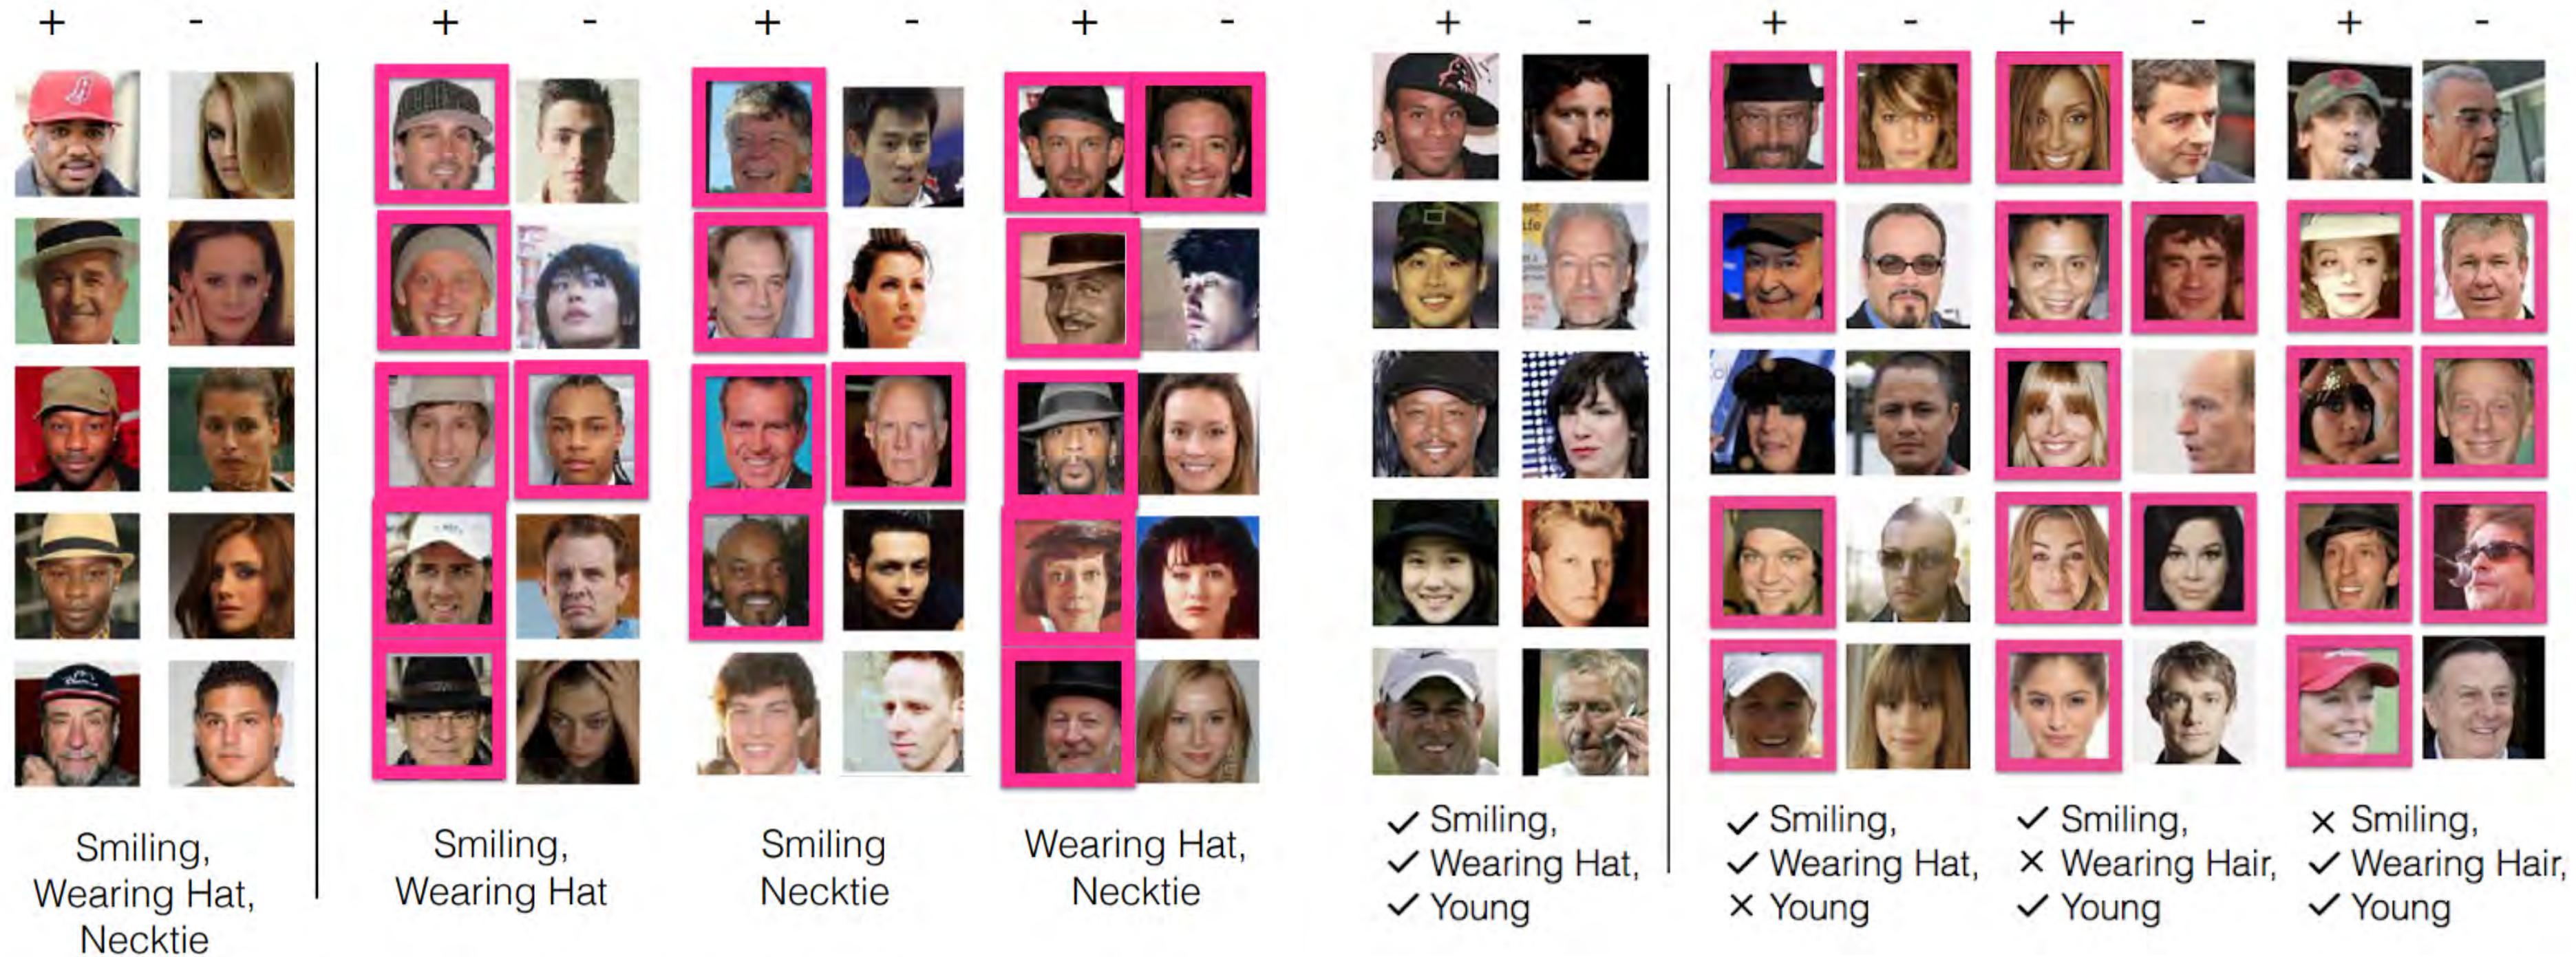


Ambiguous classification:



PLATIPUS

Probabilistic LATent model for Incorporating Priors and Uncertainty in few-Shot learning



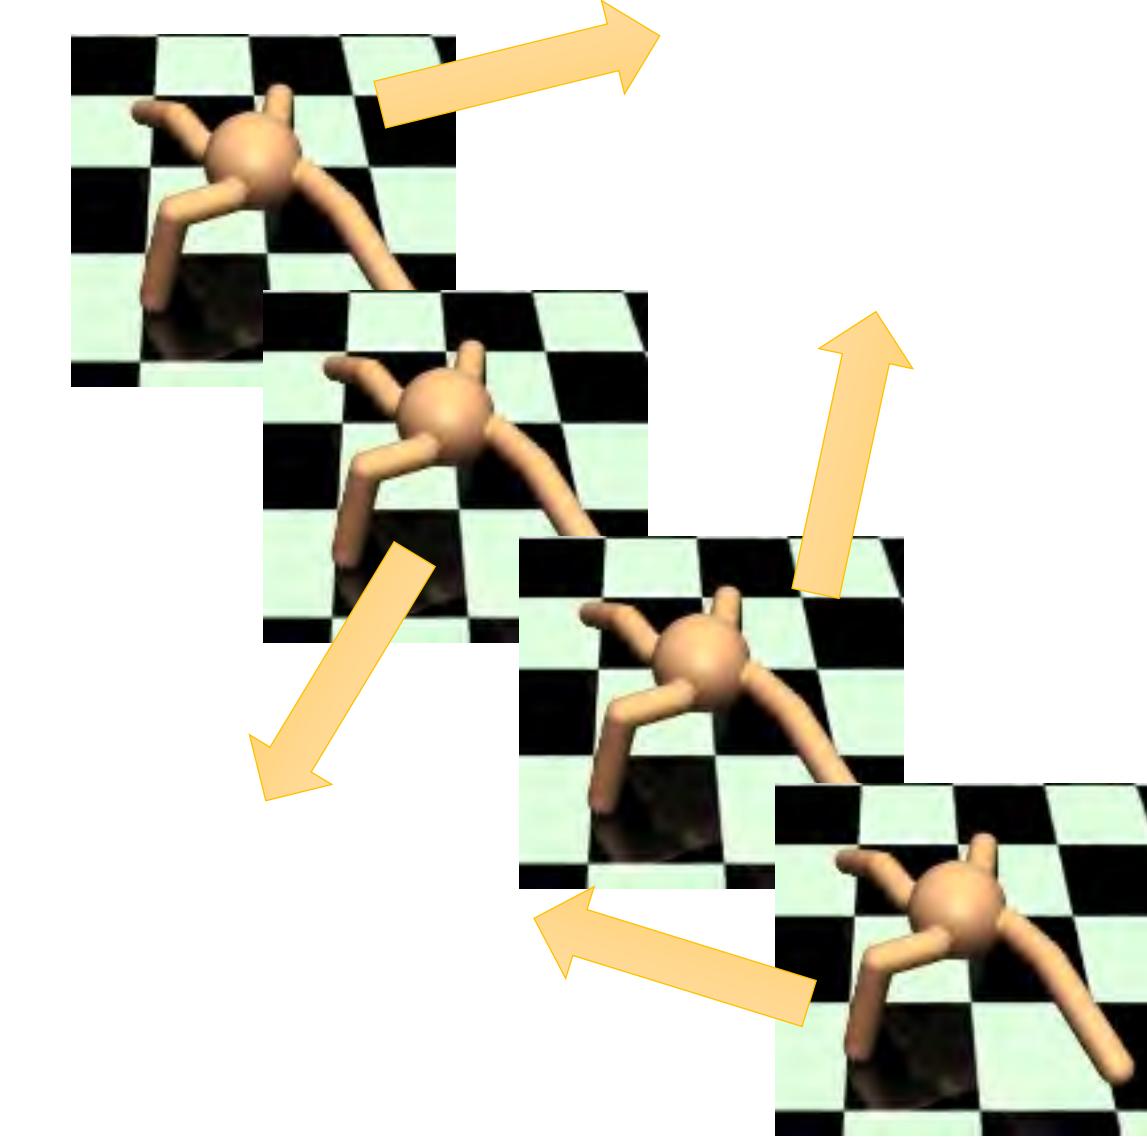
Ambiguous celebA (5-shot)		
	Accuracy	Coverage (max=3)
MAML	$69.26 \pm 2.18\%$	1.00 ± 0.0
MAML + noise	$54.73 \pm 0.8\%$	2.60 ± 0.12
PLATIPUS (ours)	$69.97 \pm 1.32\%$	2.62 ± 0.11

Outline

- Meta-Learning Problem Statement
- Model-Agnostic Meta-Learning (MAML)
- Probabilistic Interpretation of MAML
- Meta-Learning with Automated Task Proposals**
- Extensions to Robot Imitation & Intent Inference

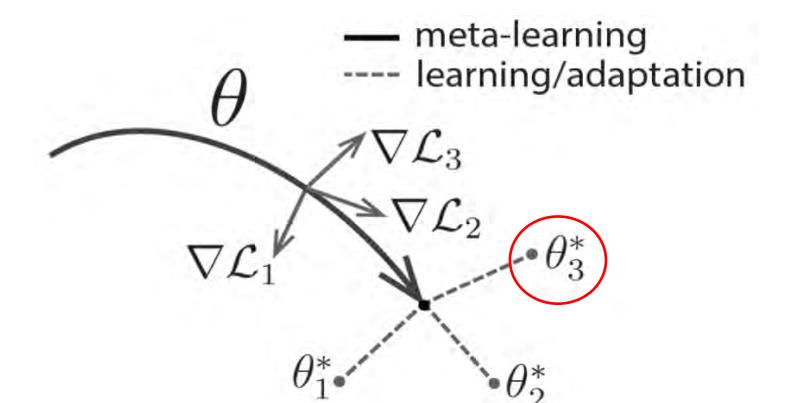
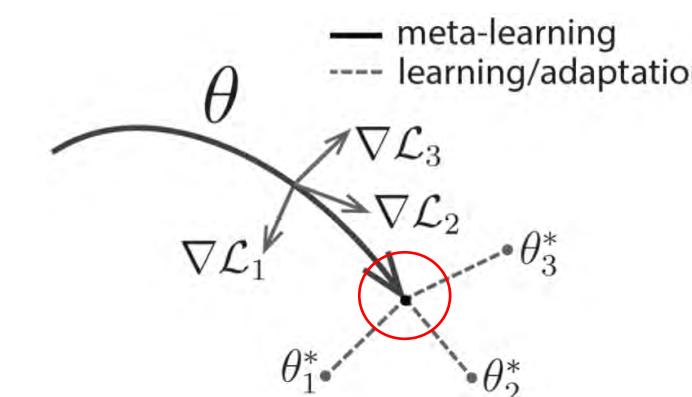
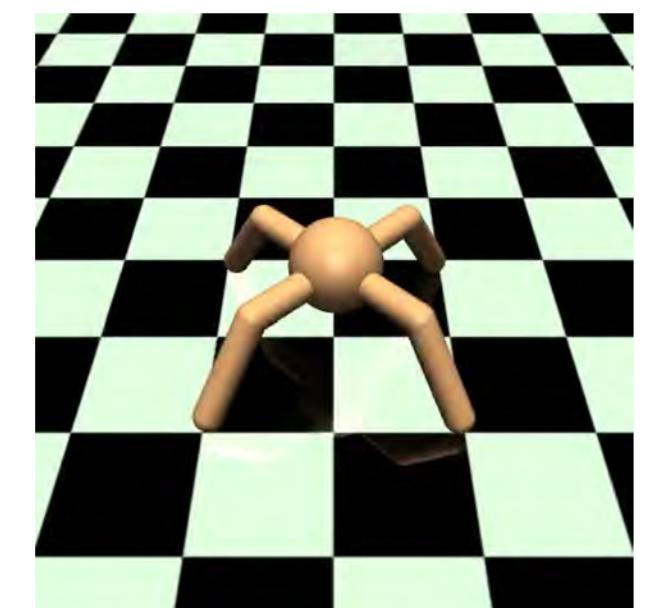
Let's Talk about Meta-Overfitting

- Meta learning requires task distributions
- When there are too few meta-training tasks, we can *meta-overfit*
- Specifying task distributions is hard, especially for meta-RL!
- Can we propose tasks *automatically*?

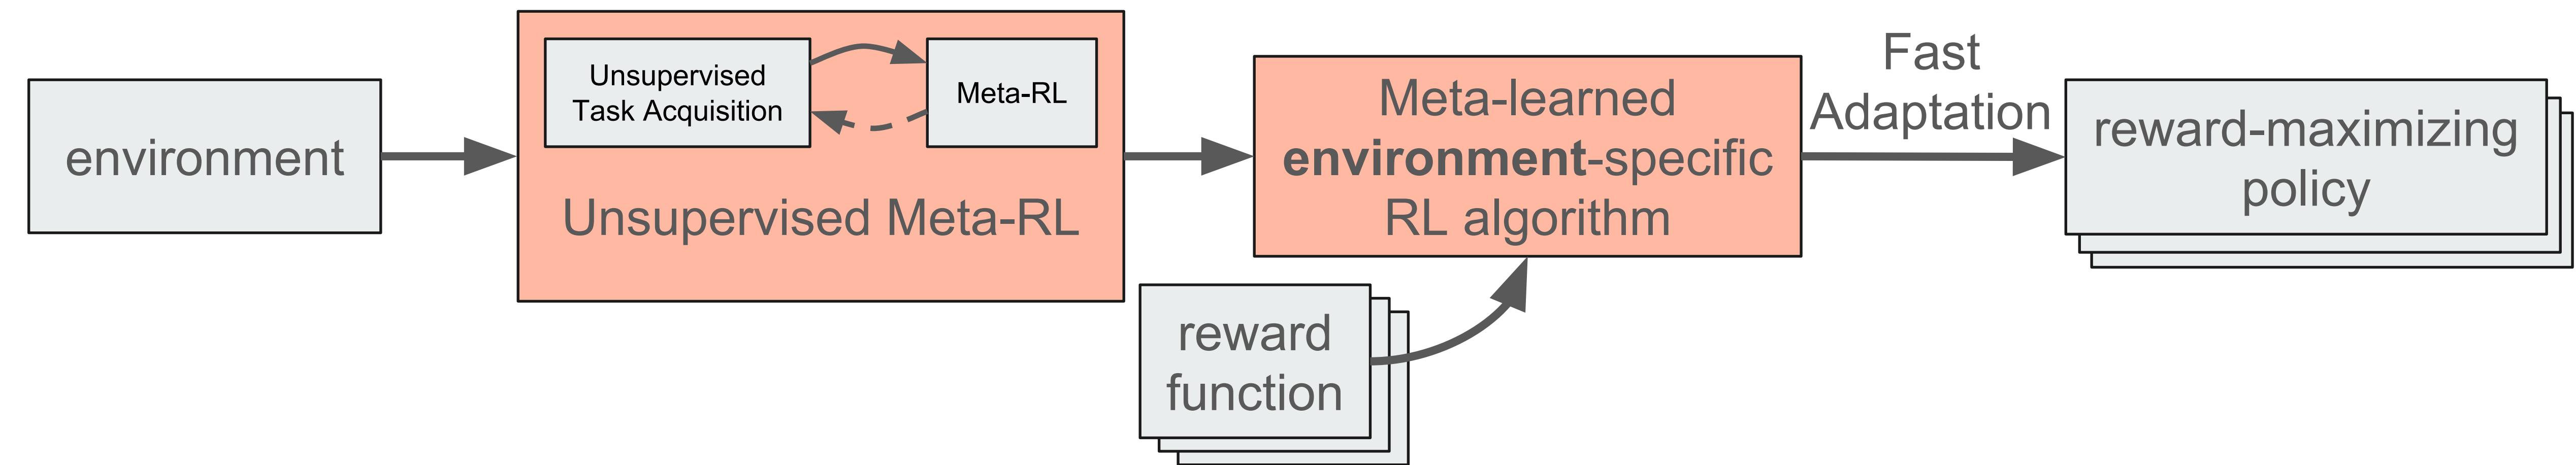


after MAML training

after 1 gradient step



A General Recipe for Unsupervised RL

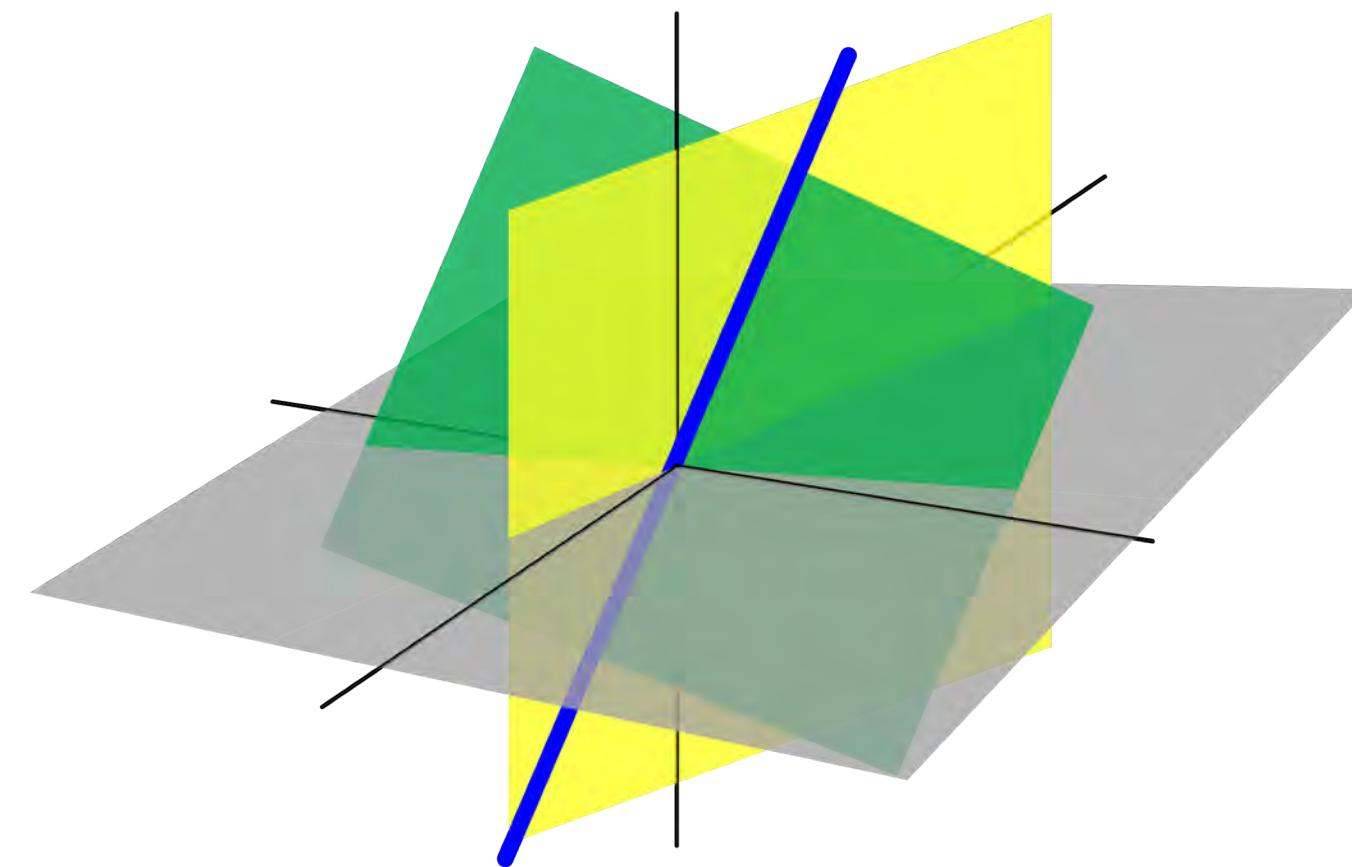
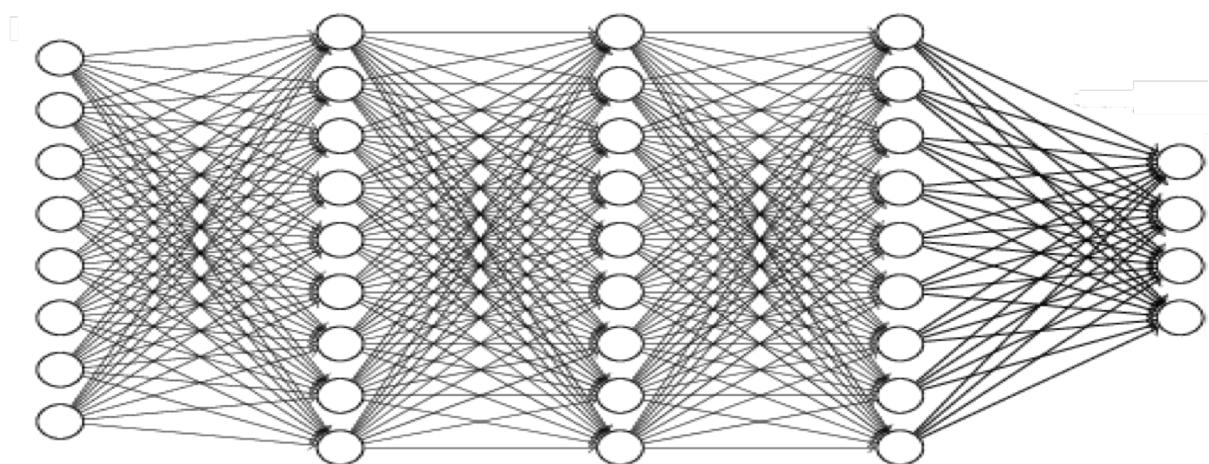


Random Task Proposals

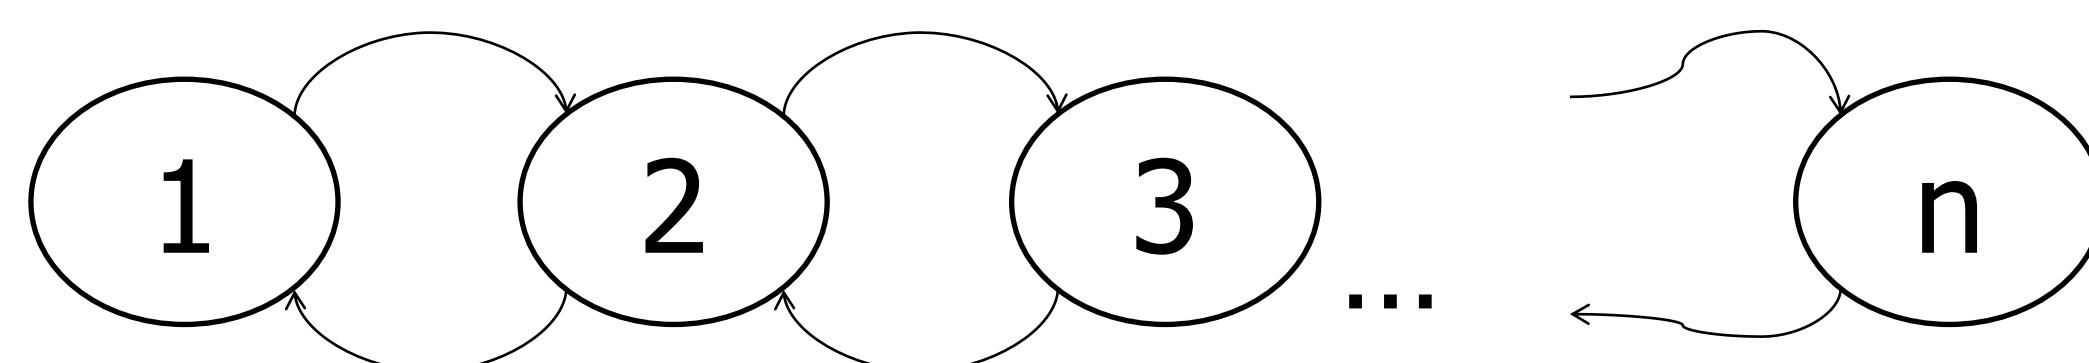
- Use randomly initialize discriminators for reward functions

$$R(s, z) = \log p_D(z|s)$$

D → randomly initialized network

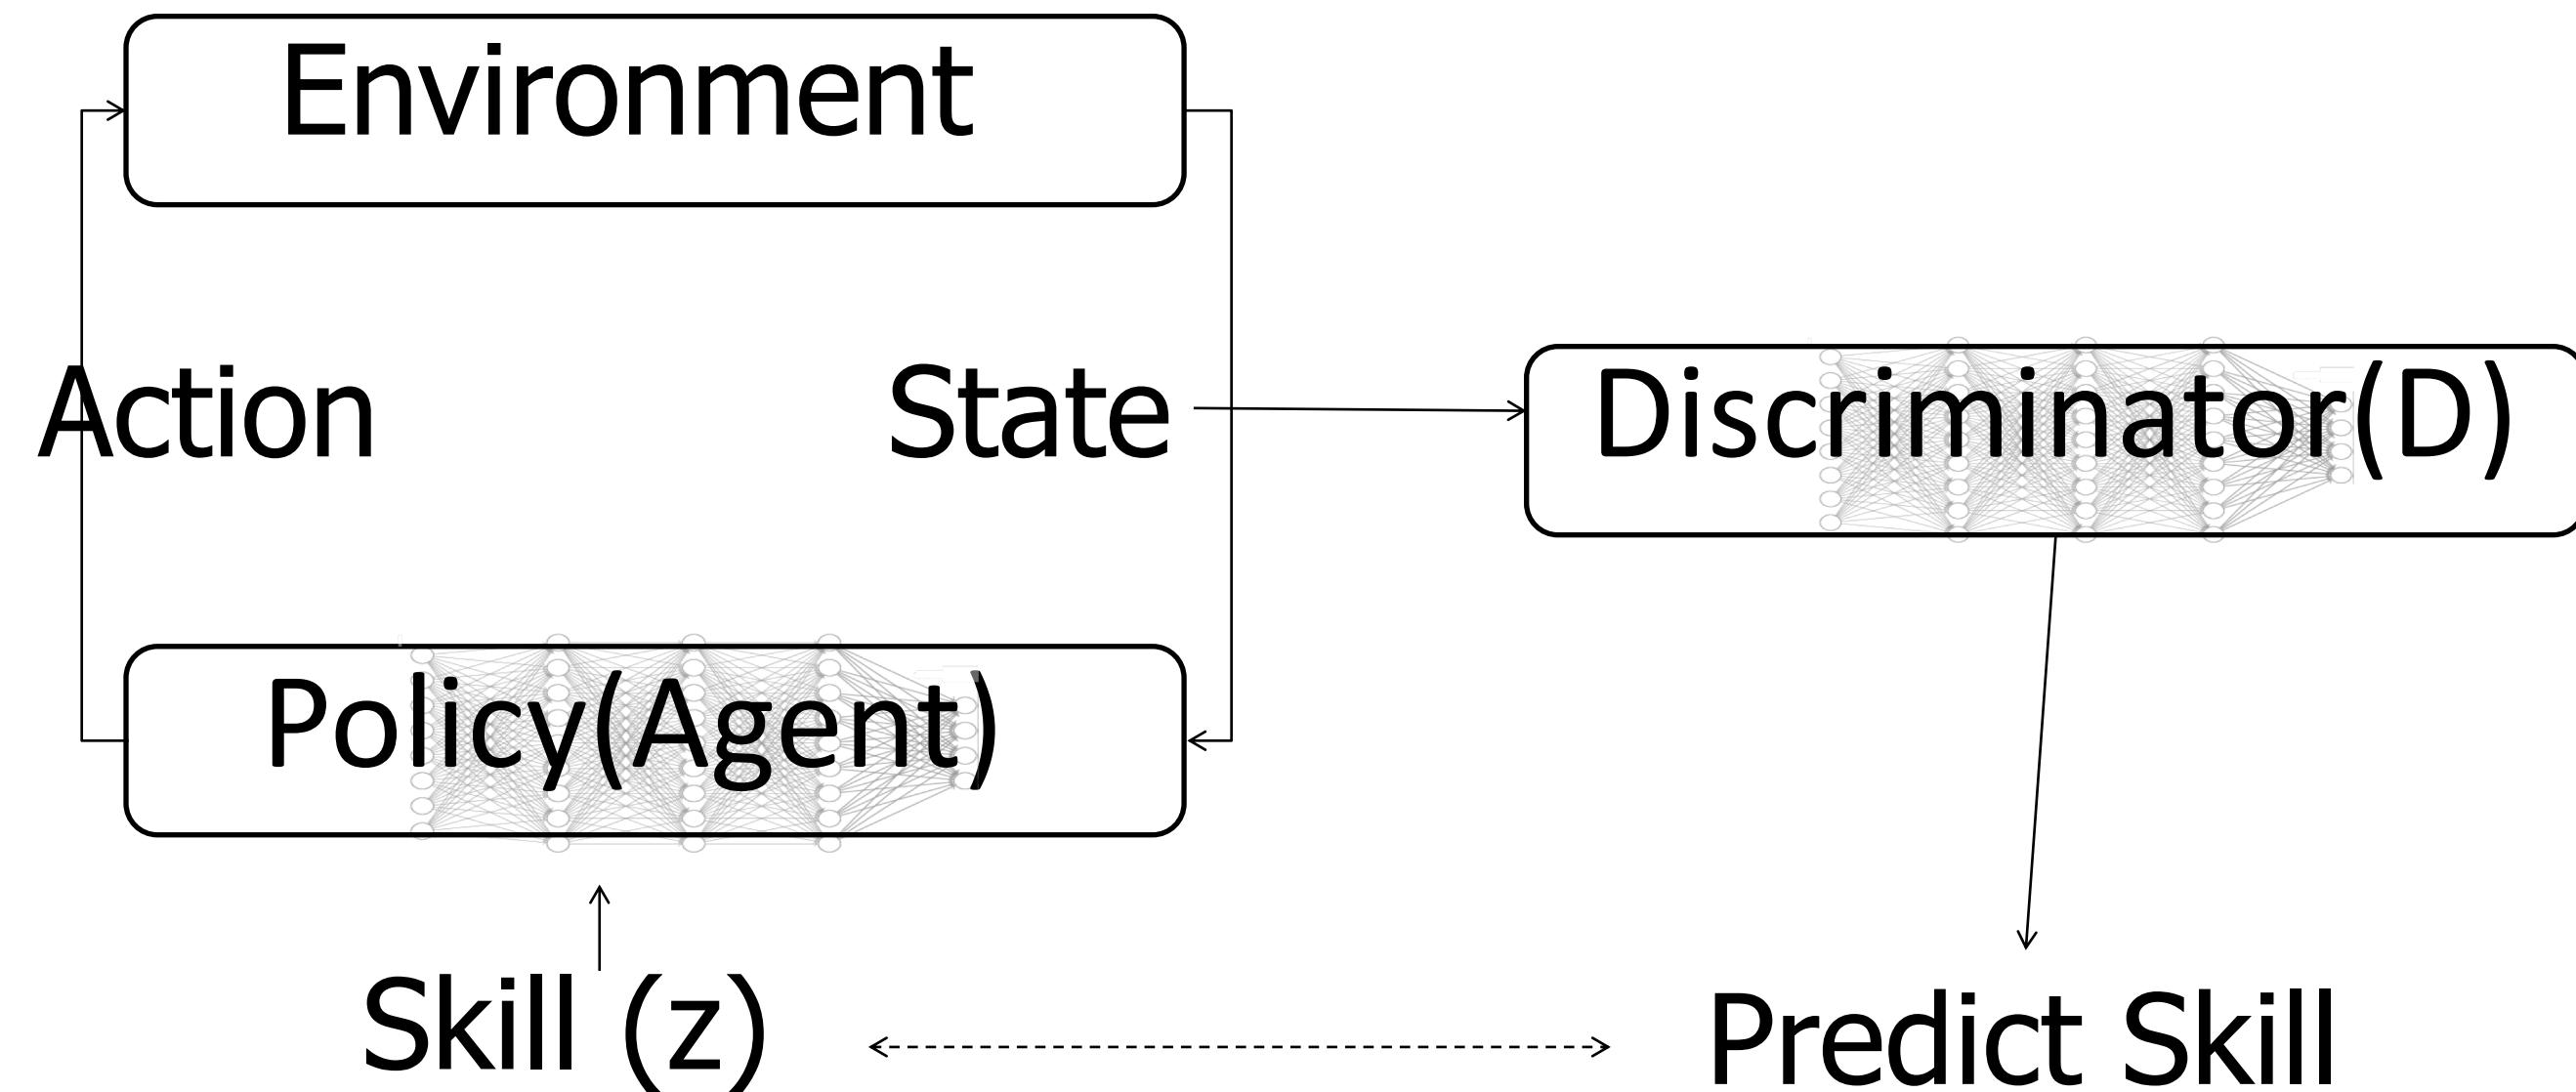


- Important: Random functions over state space, **not** random policies



Random policy – exponential
Random reward – polynomial

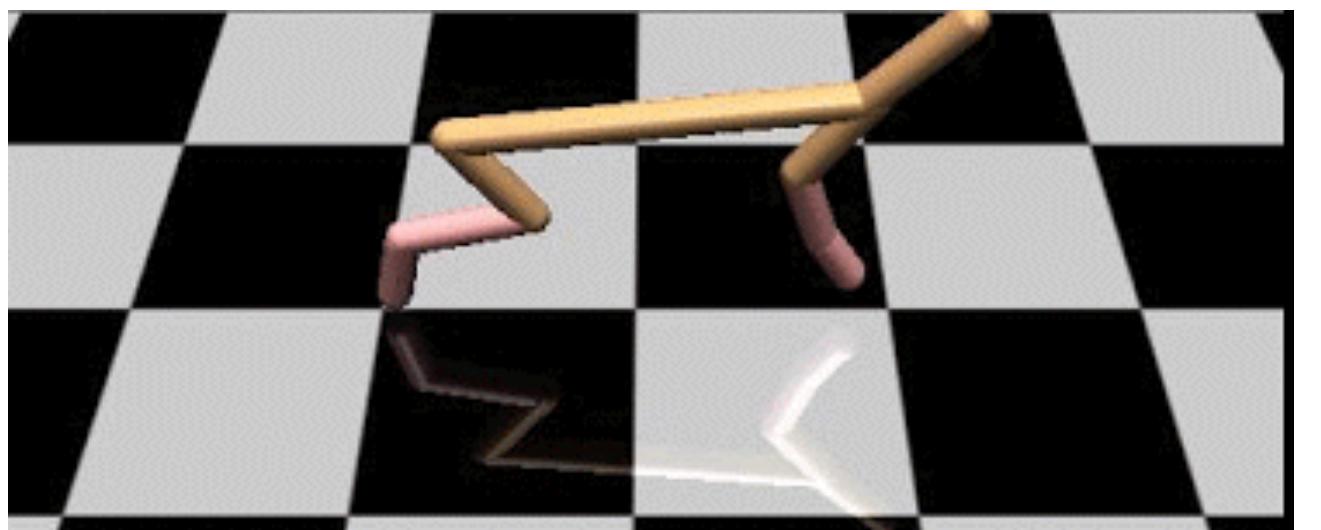
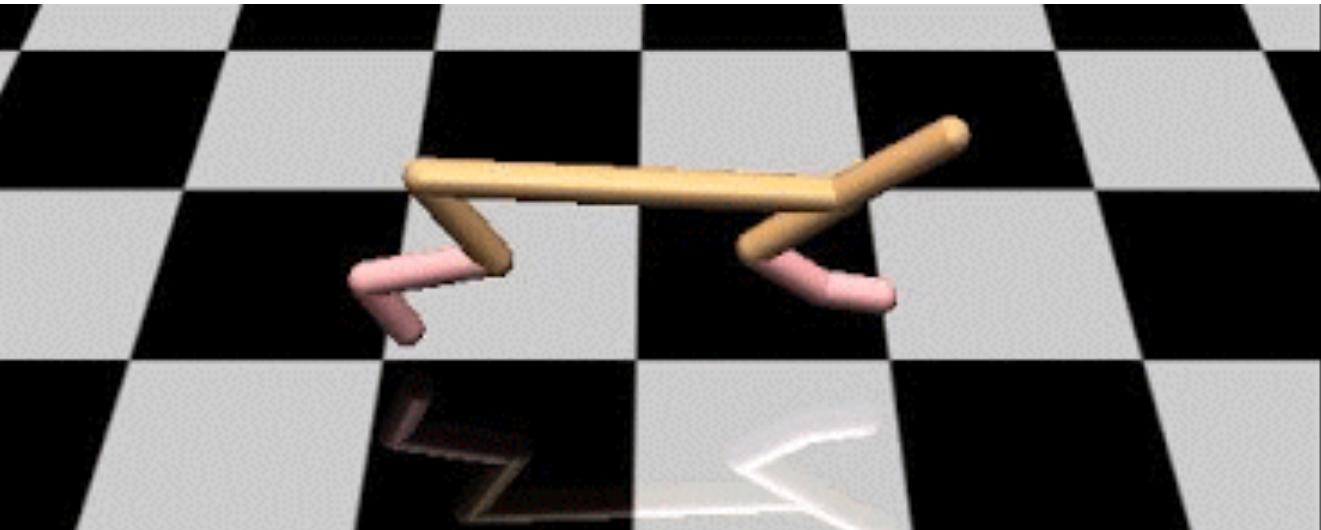
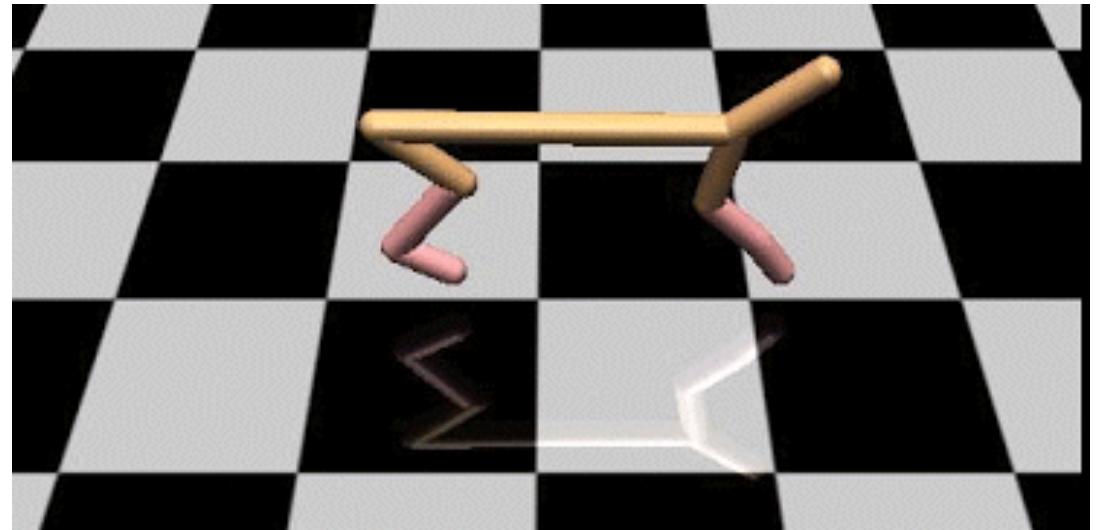
Diversity-Driven Proposals



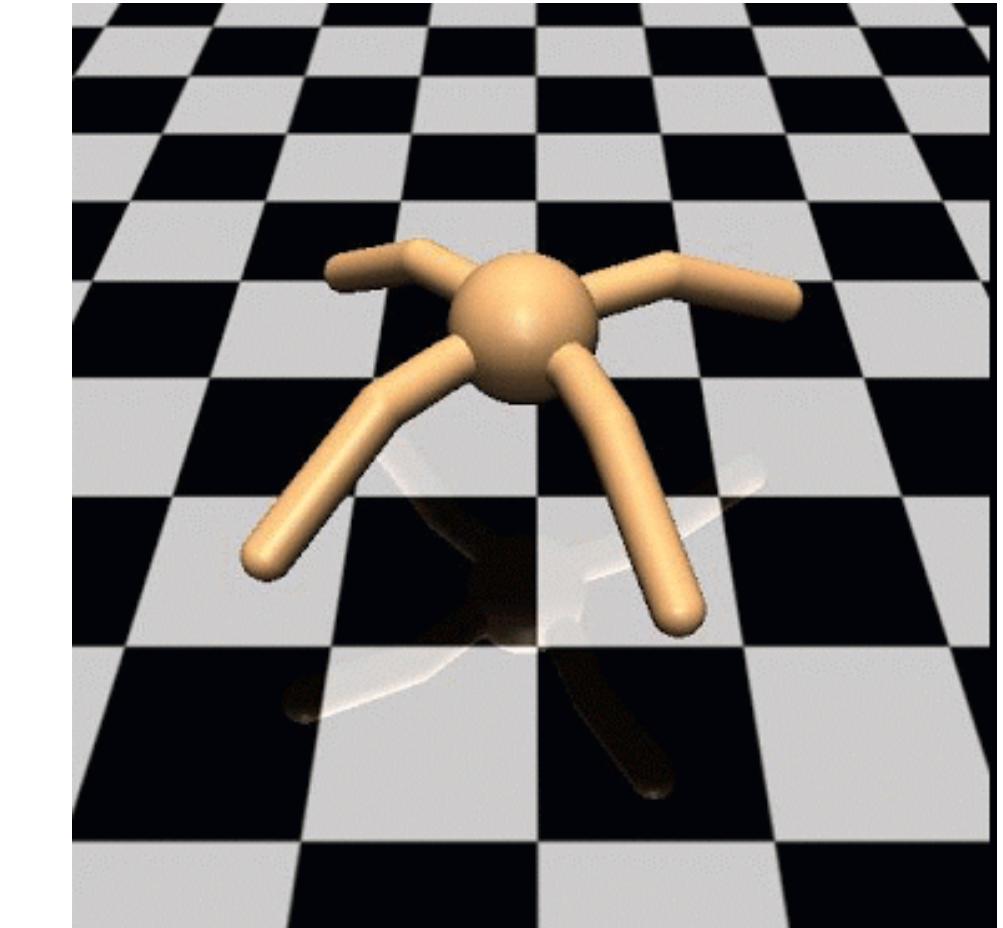
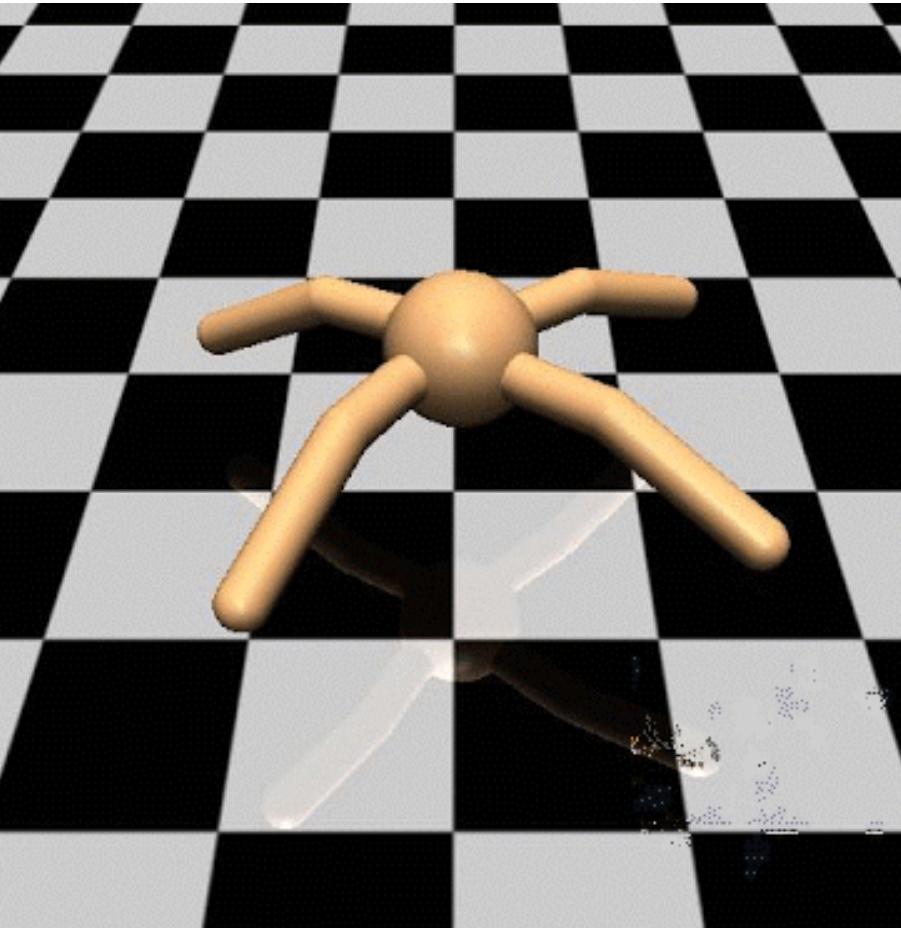
- Policy → visit states which are discriminable
- Discriminator → predict skill from state

Task Reward for UML: $R(s, z) = \log p_D(z|s)$

Examples of Acquired Tasks



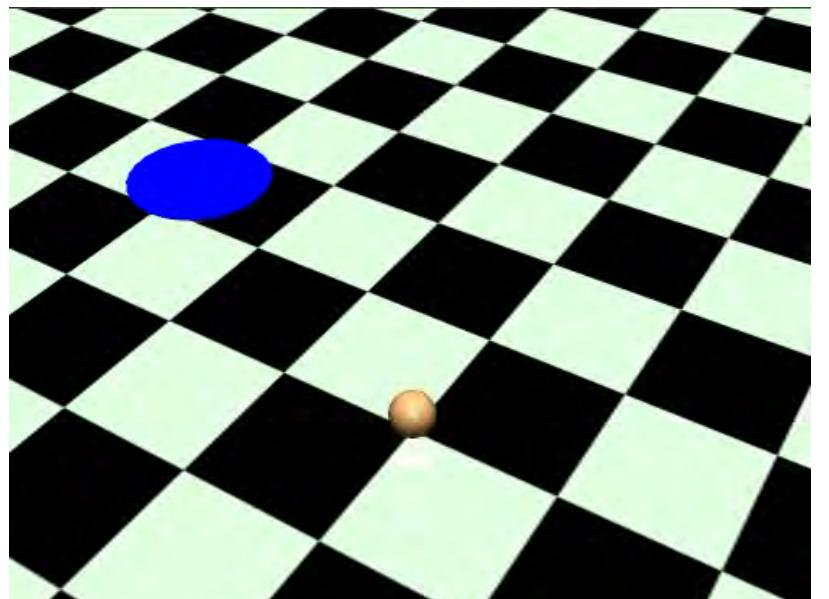
Cheetah



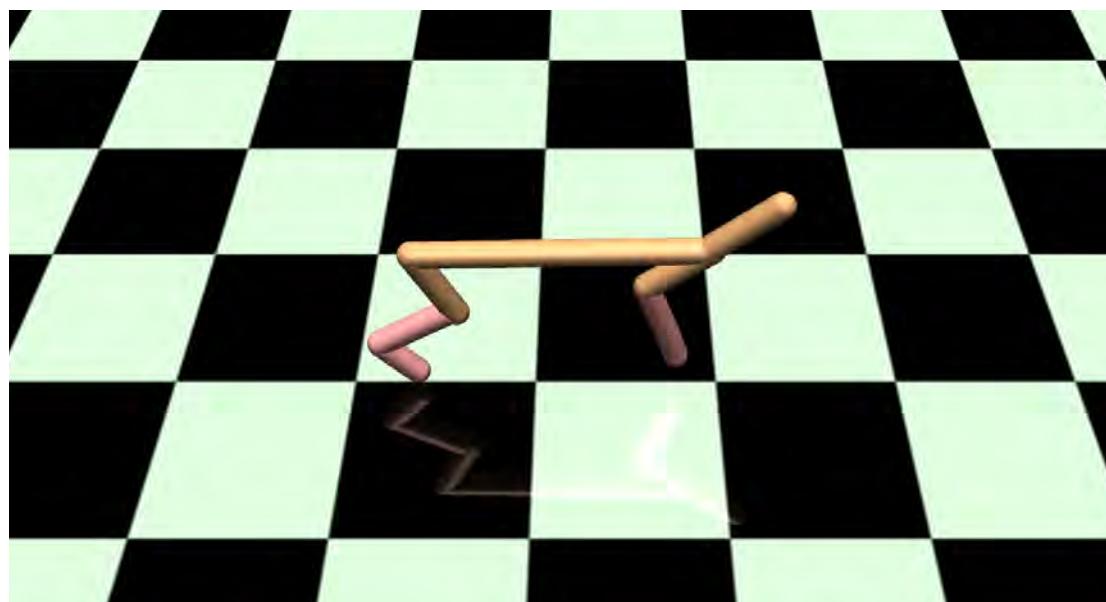
Ant

Does it work?

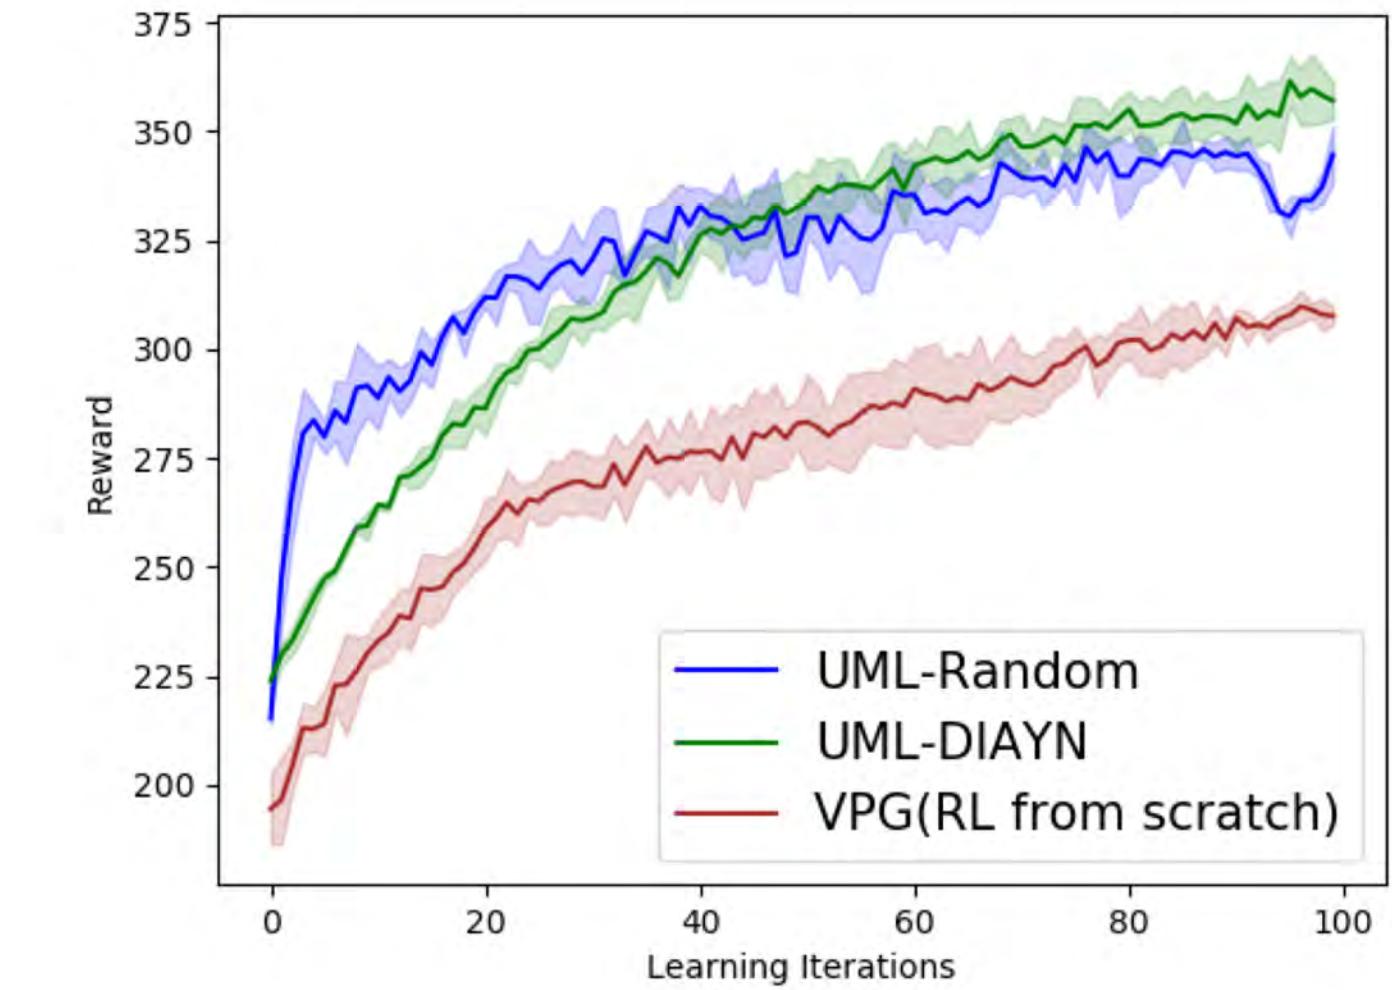
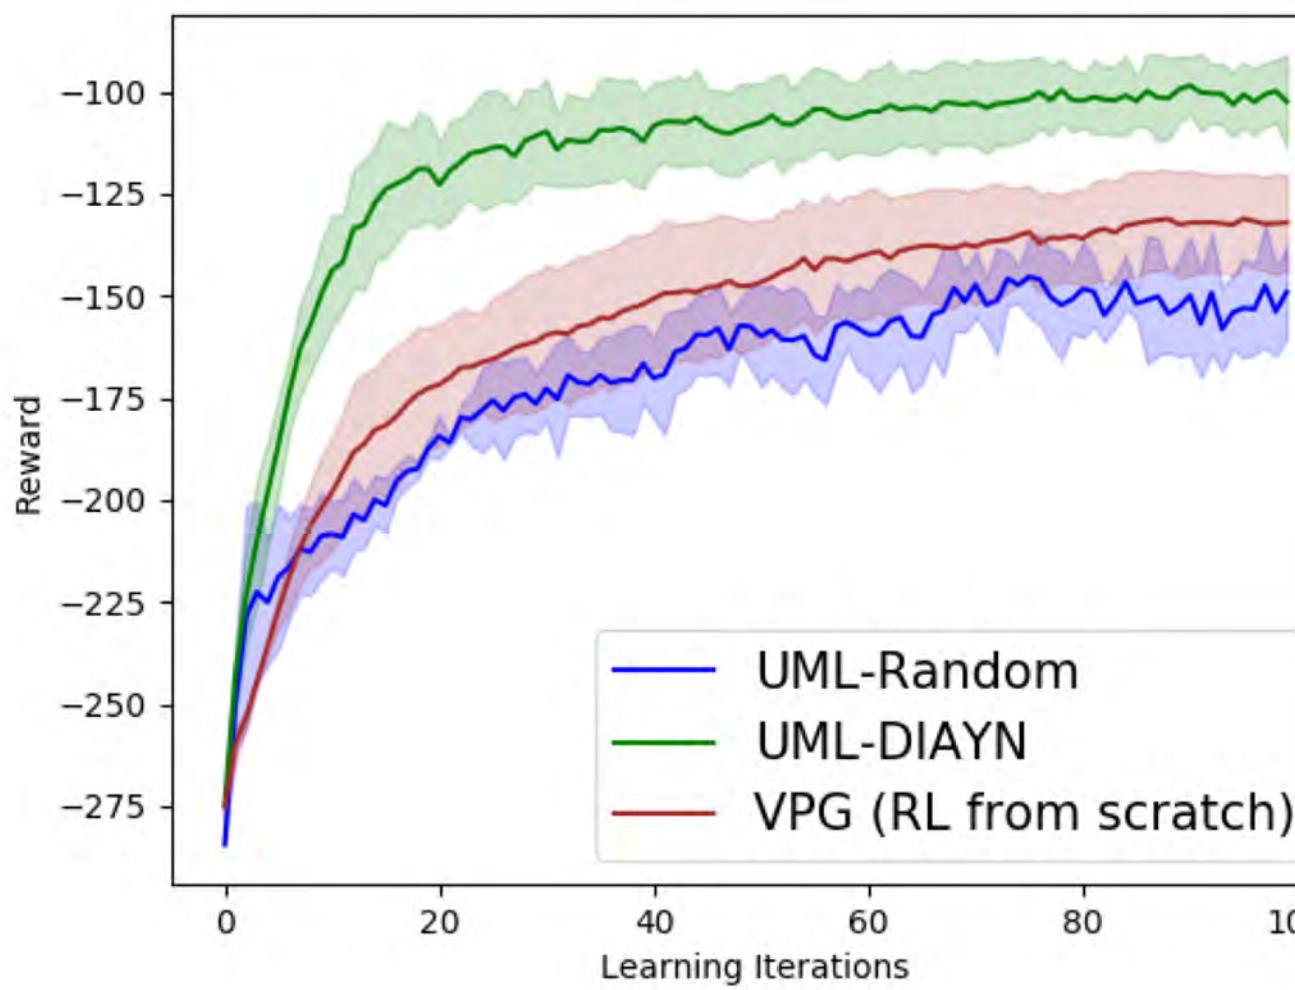
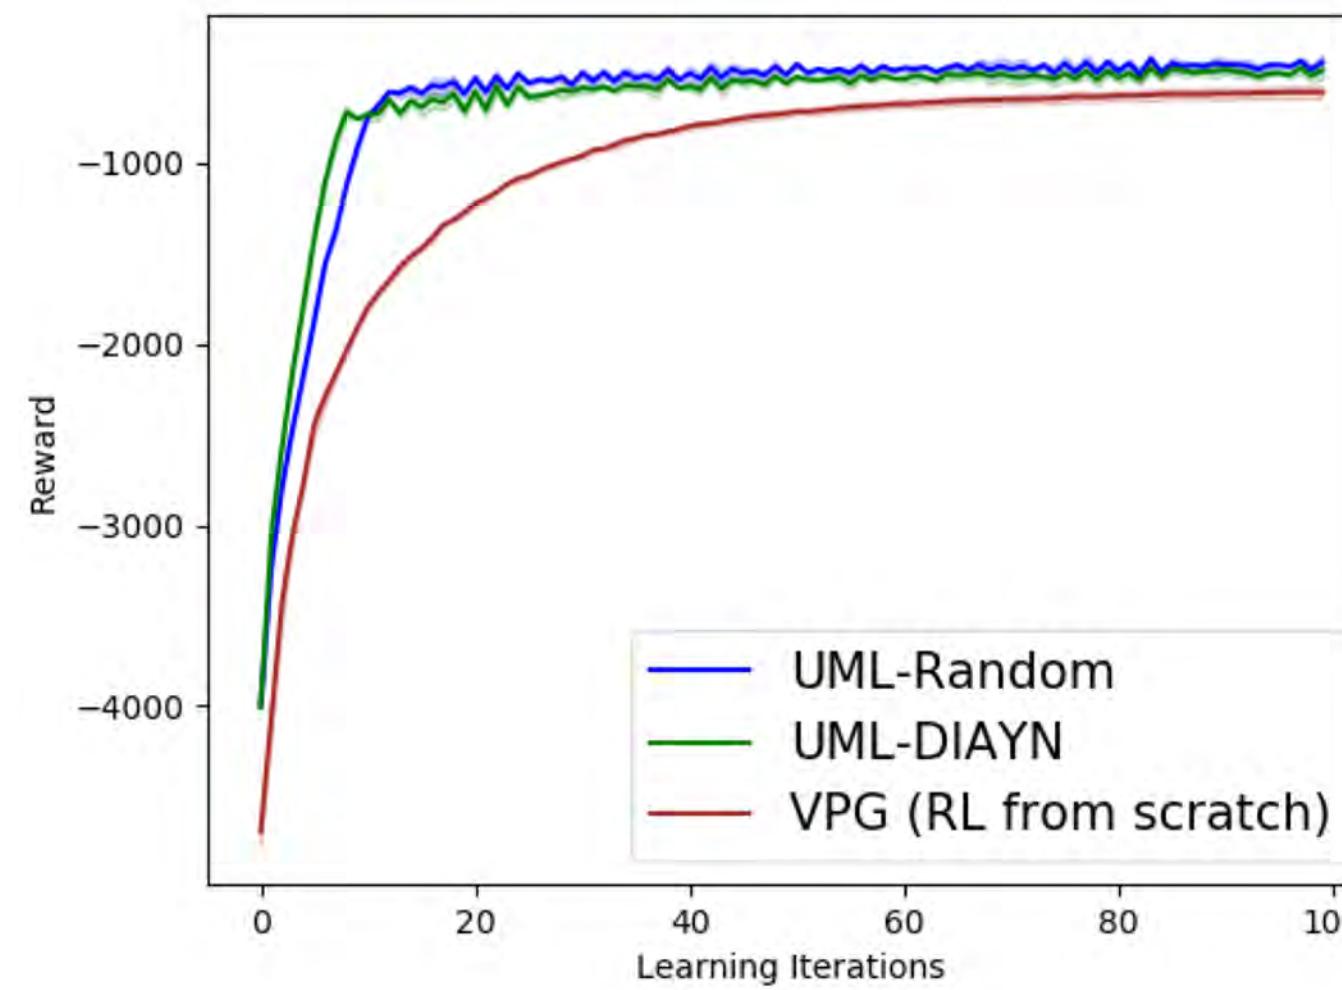
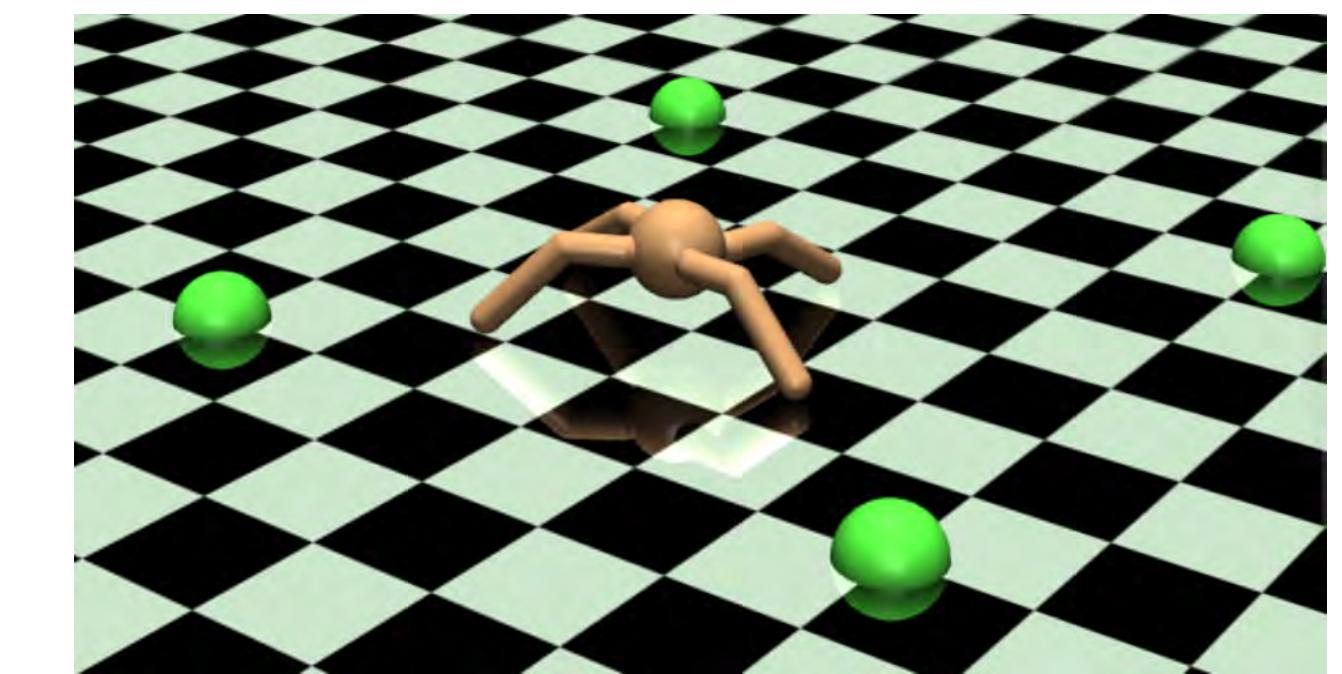
2D Navigation



Cheetah



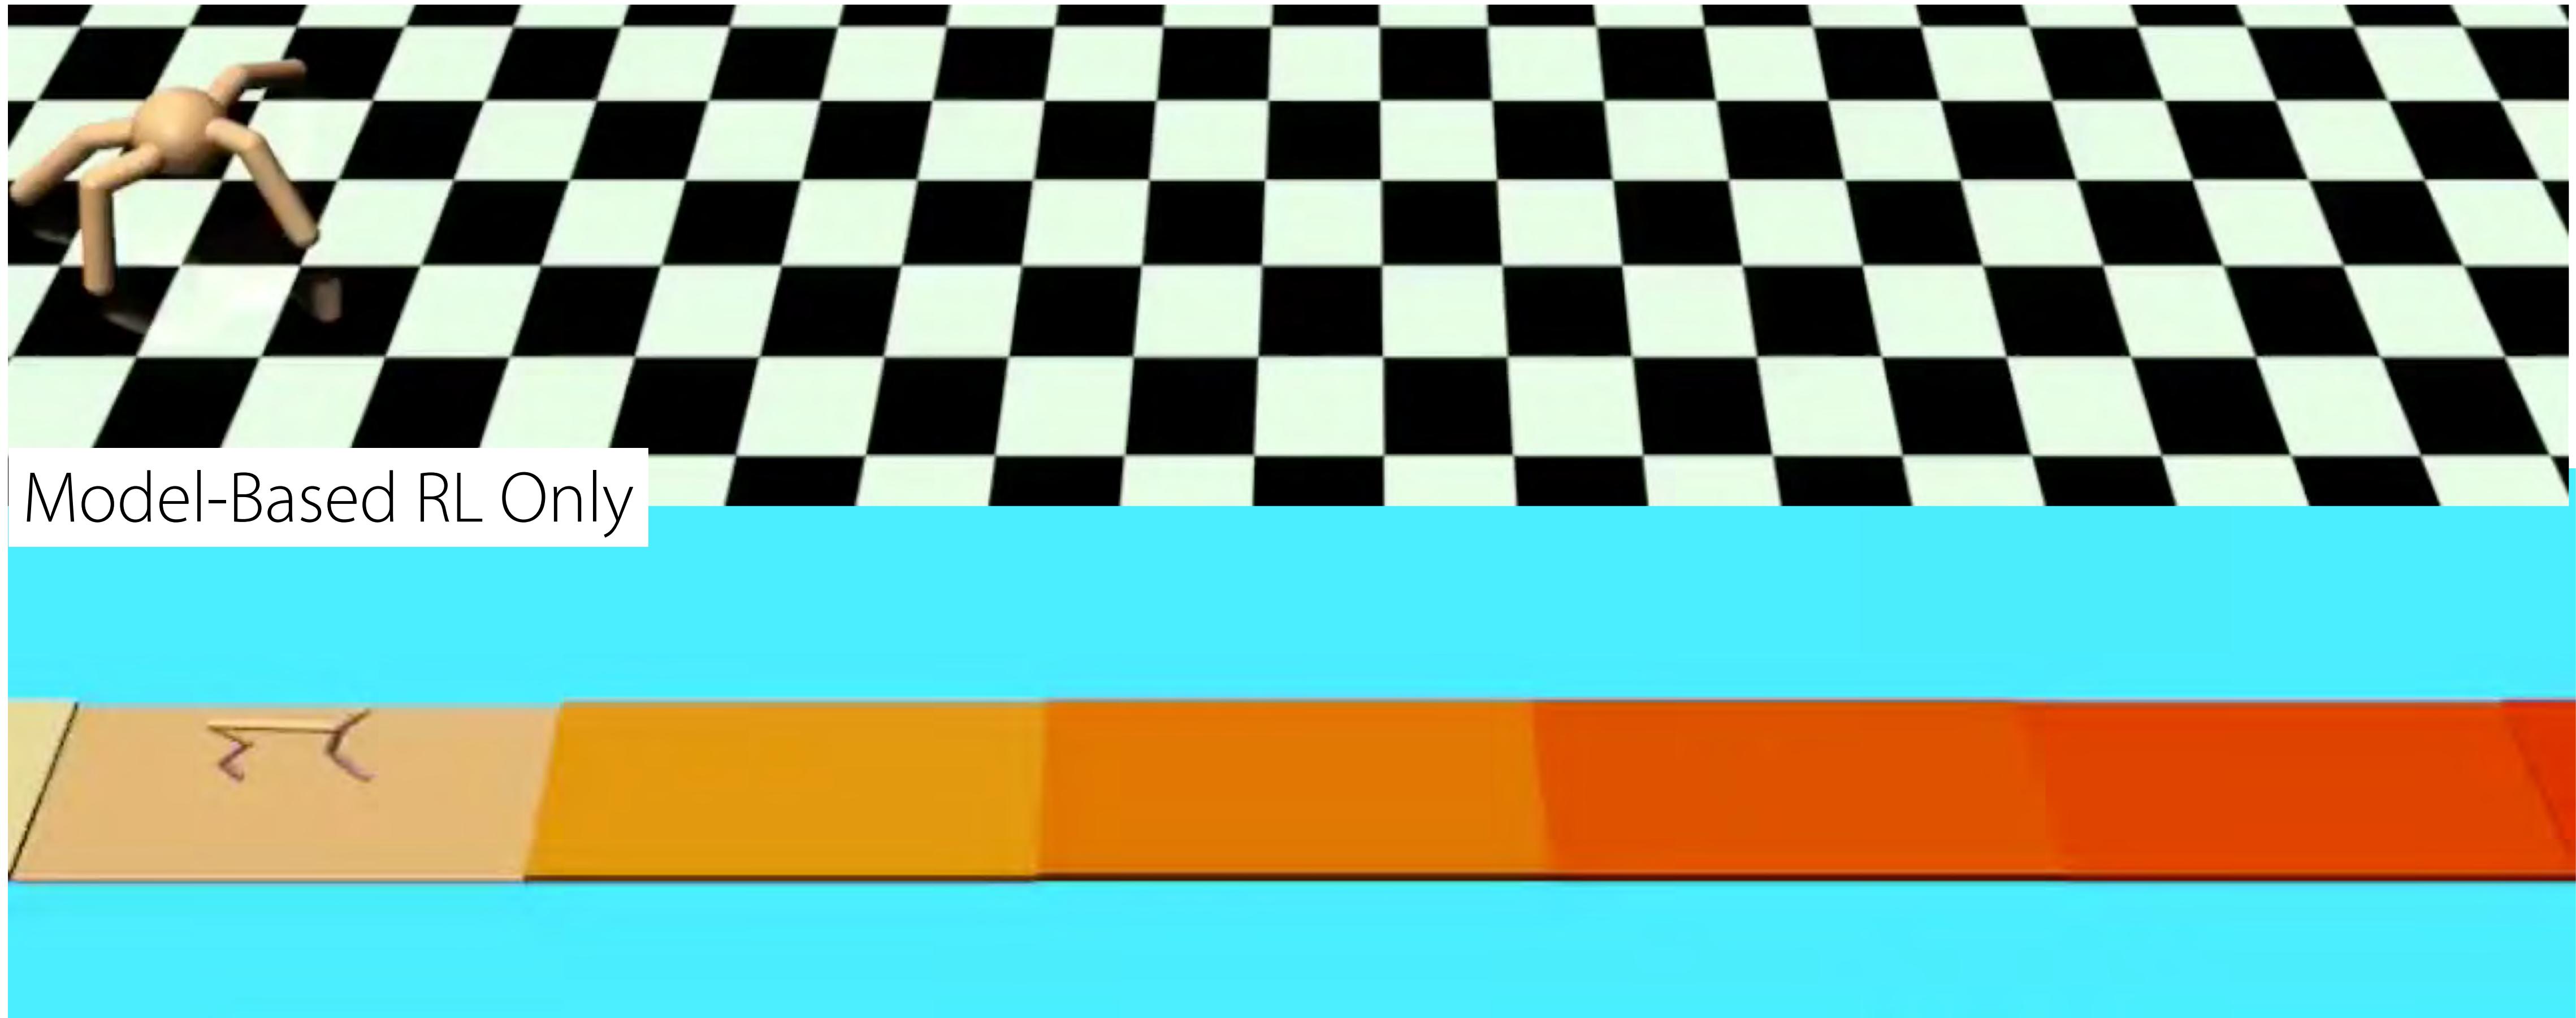
Ant



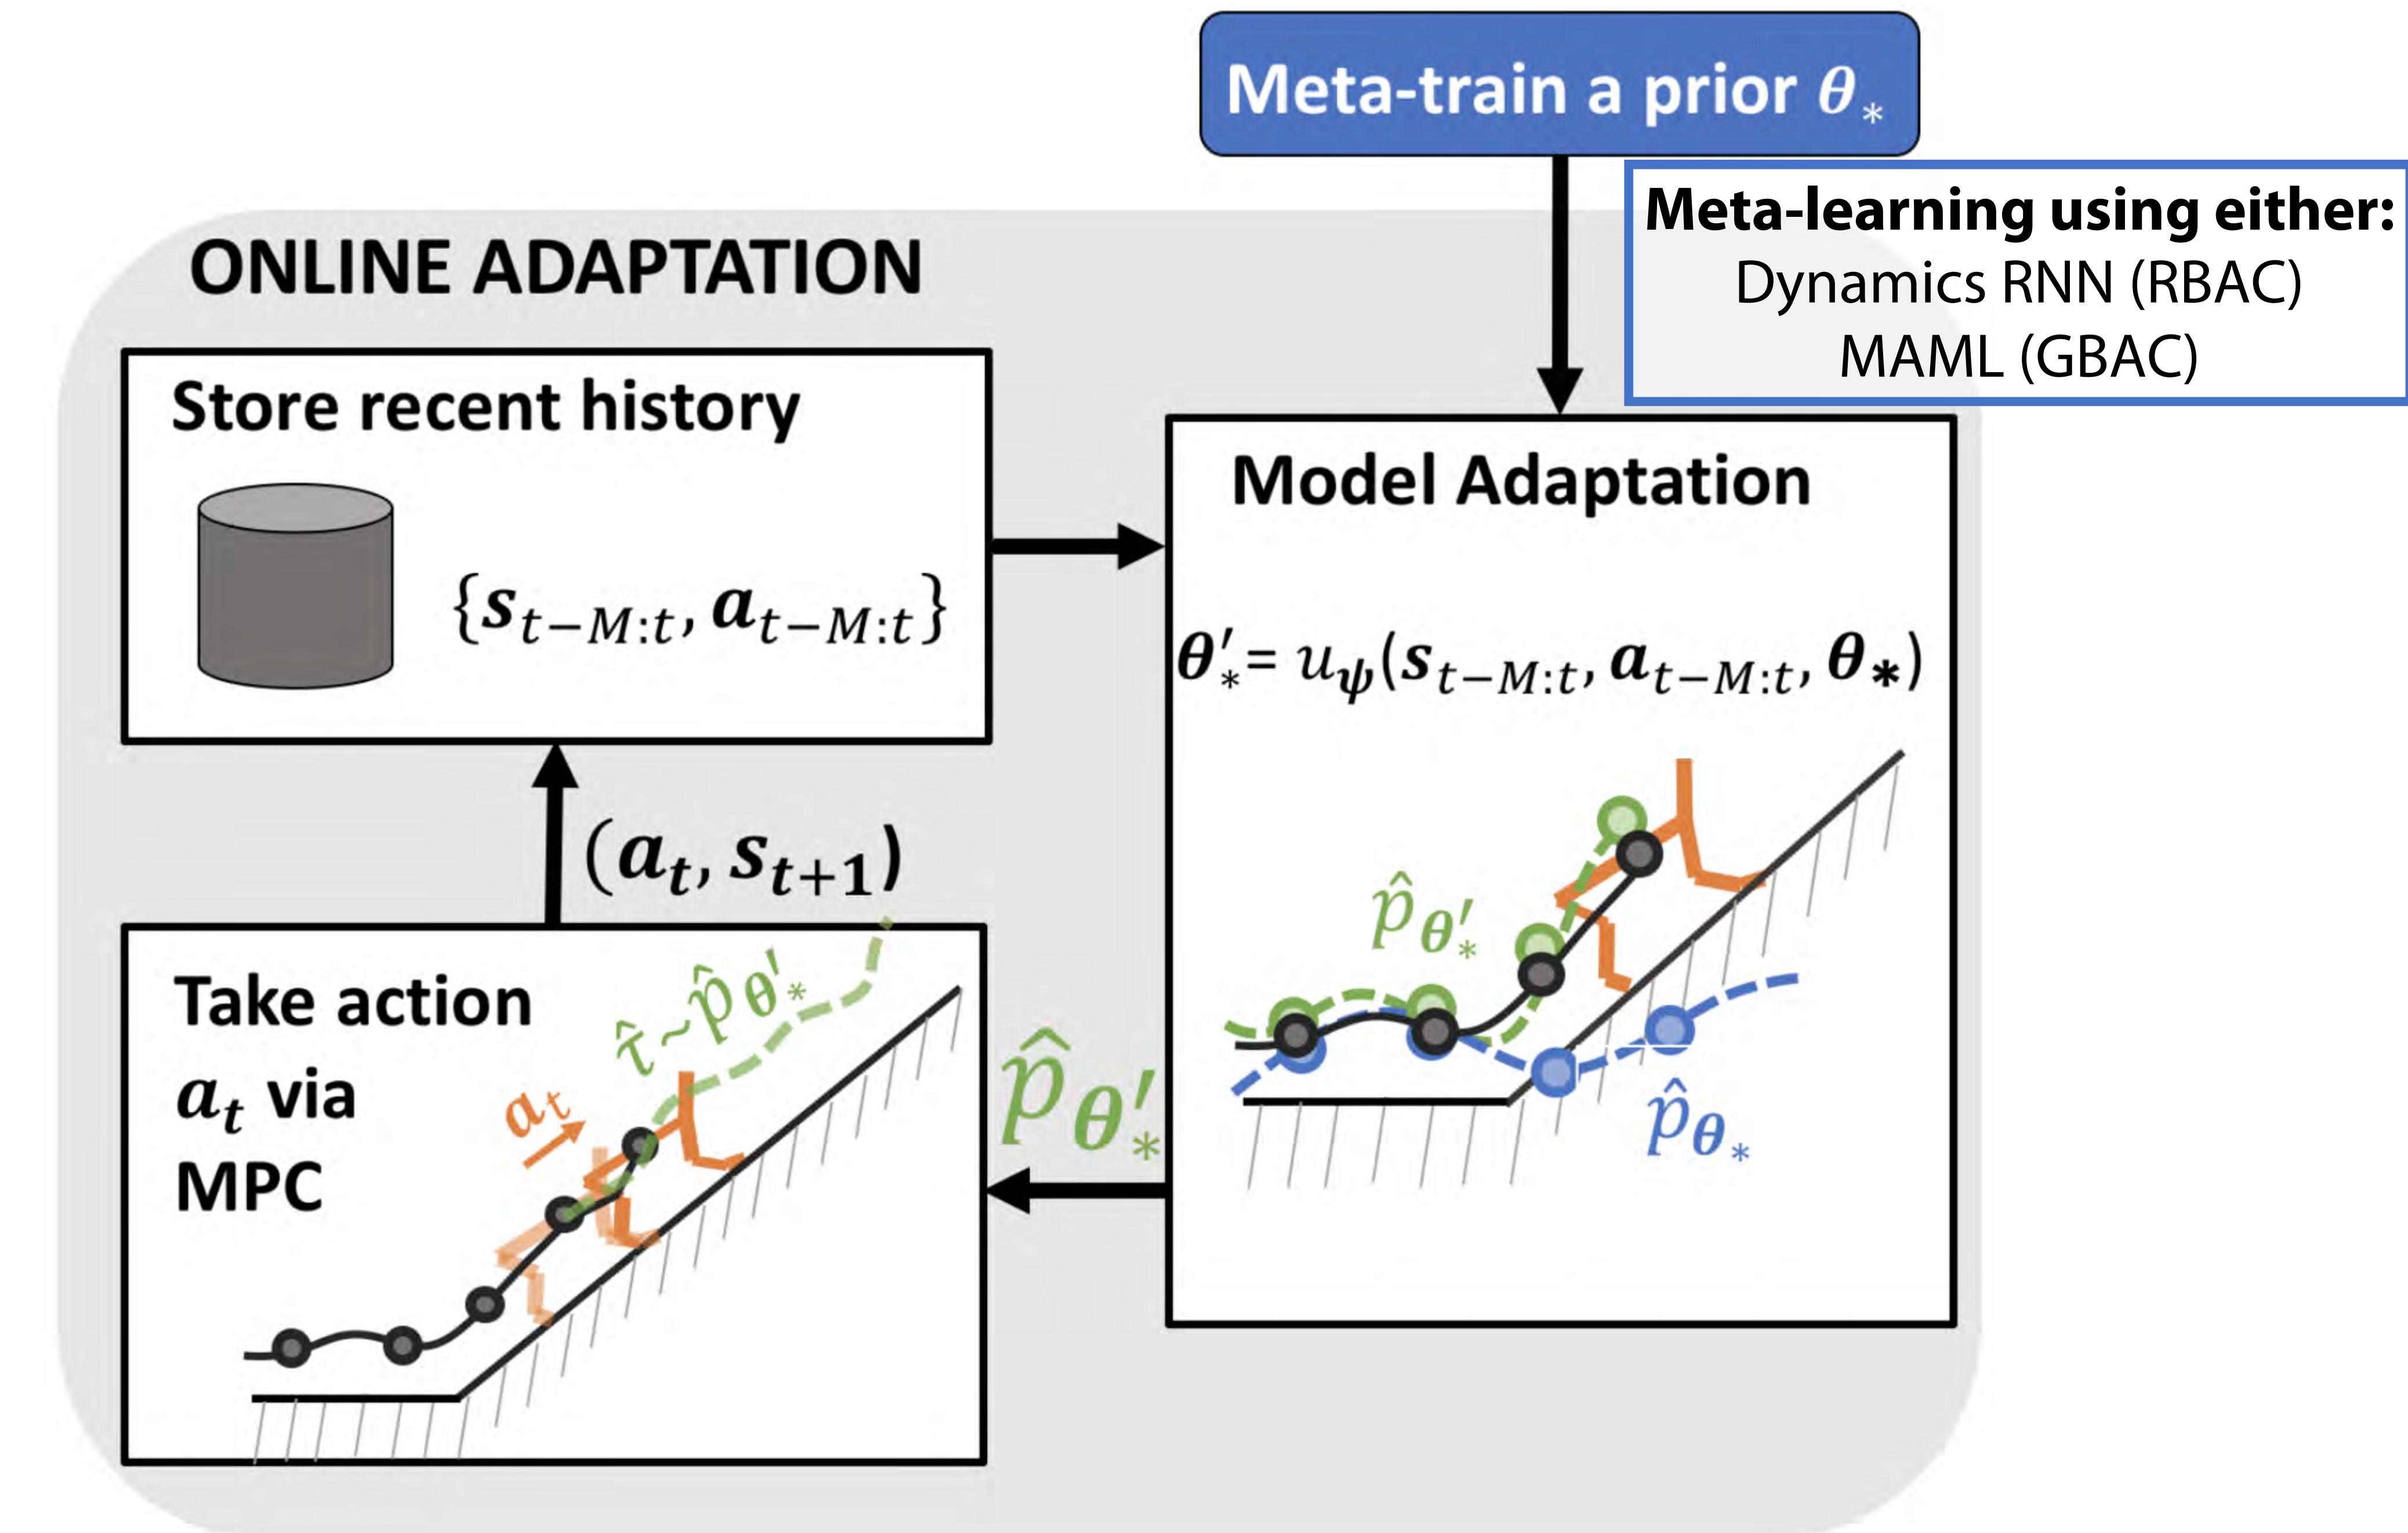
Meta-test performance with rewards

What about varying transition models?

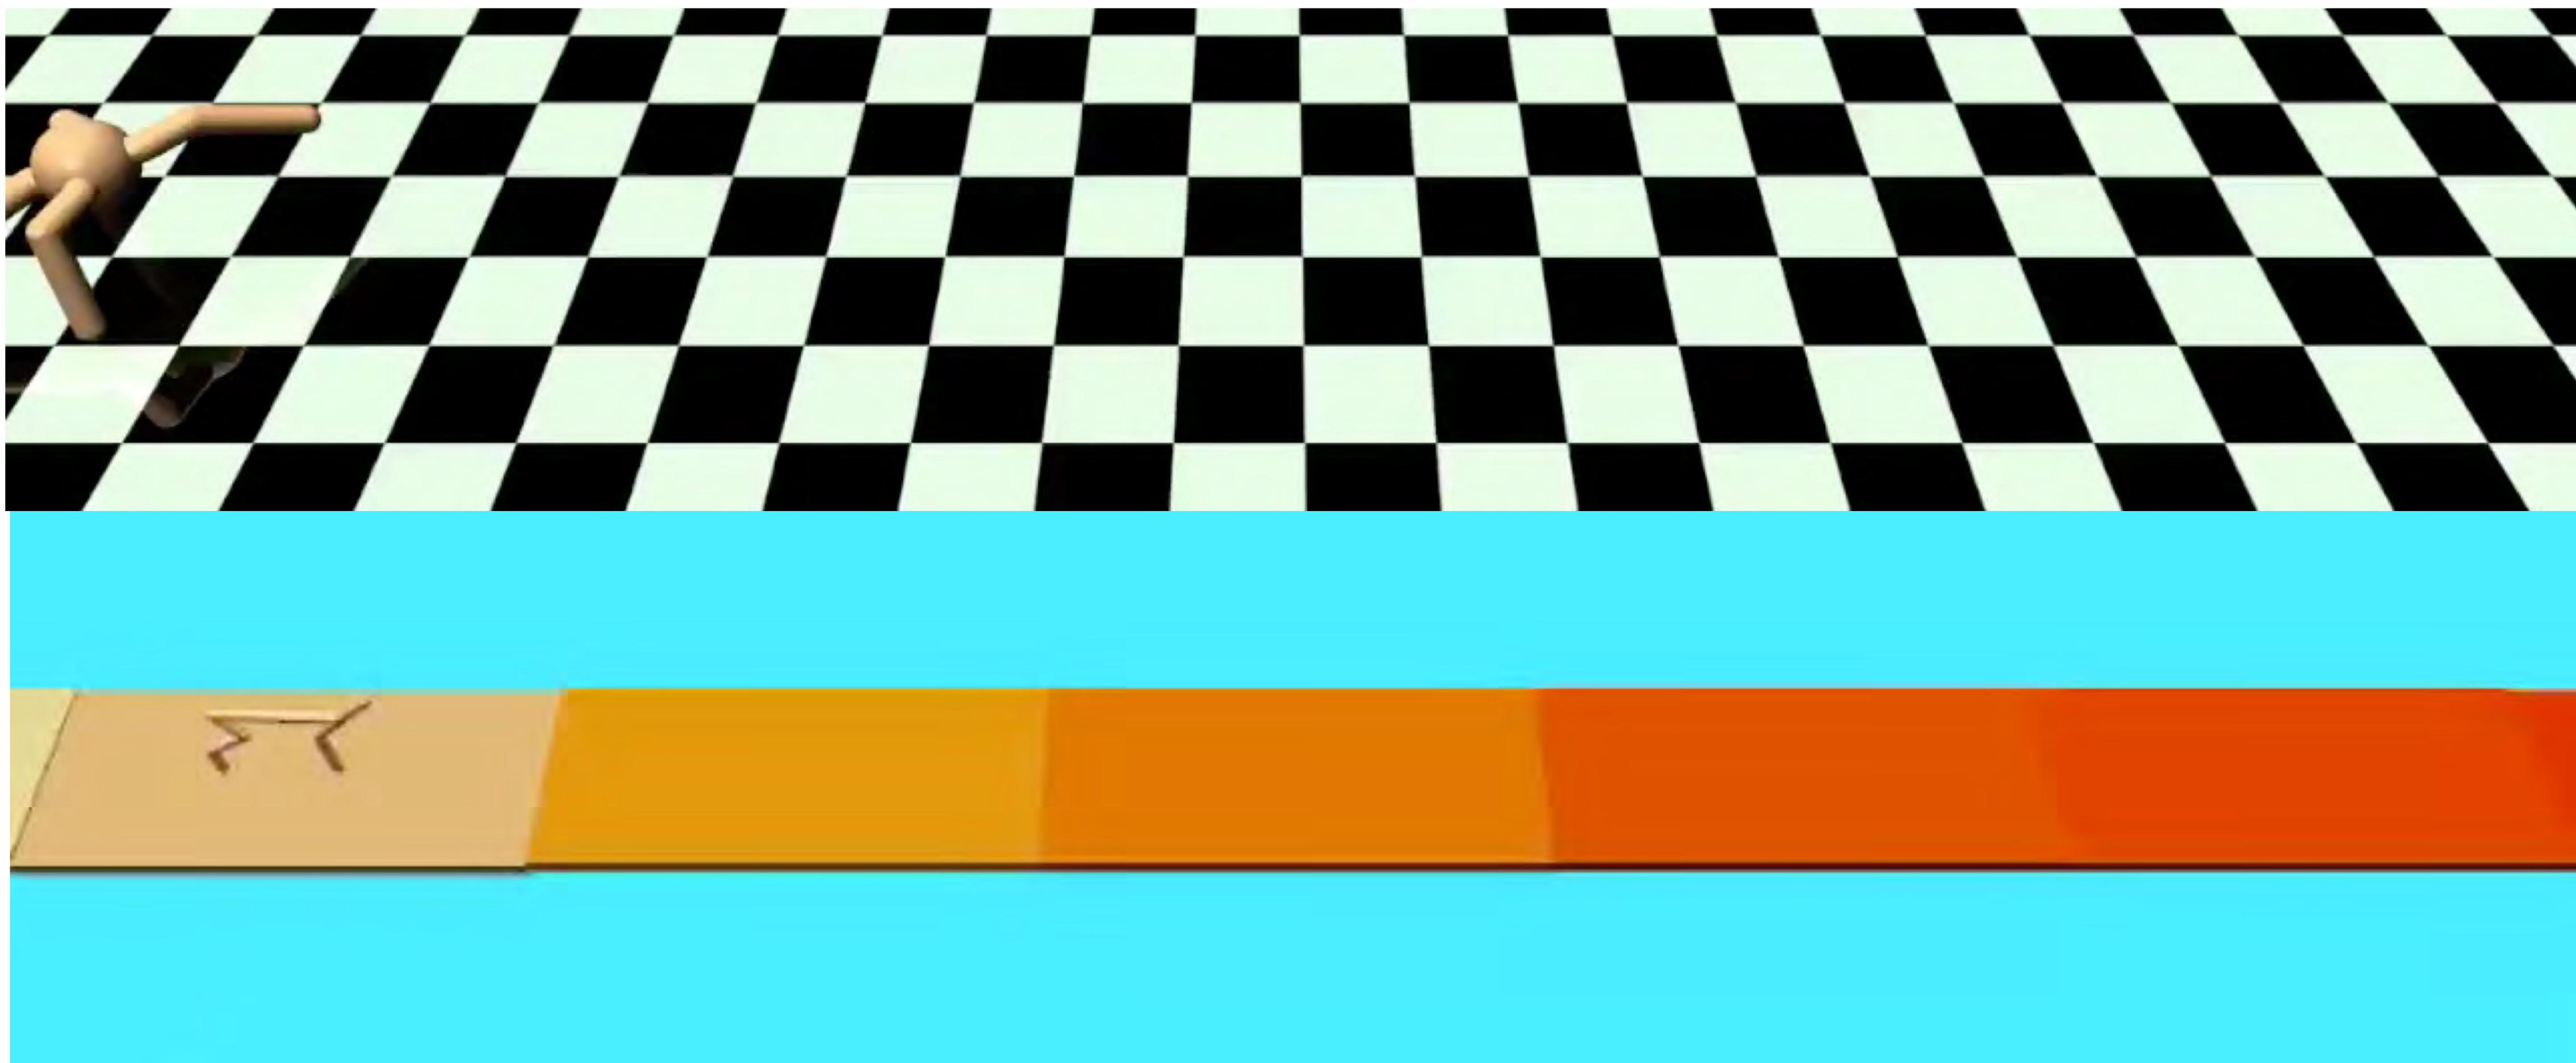
Task distribution is more-or-less free!



online adaptation = few-shot learning
tasks are **temporal**



Online Adaptation via Meta-Learning with MAML



Outline

- Meta-Learning Problem Statement
- Model-Agnostic Meta-Learning (MAML)
- Probabilistic Interpretation of MAML
- Meta-Learning with Automated Task Proposals
- **Extensions to Robot Imitation & Intent Inference**

One-Shot Visual Imitation Learning

Goal: Given one visual demonstration of a new task, learn a policy

Visual imitation is expensive.

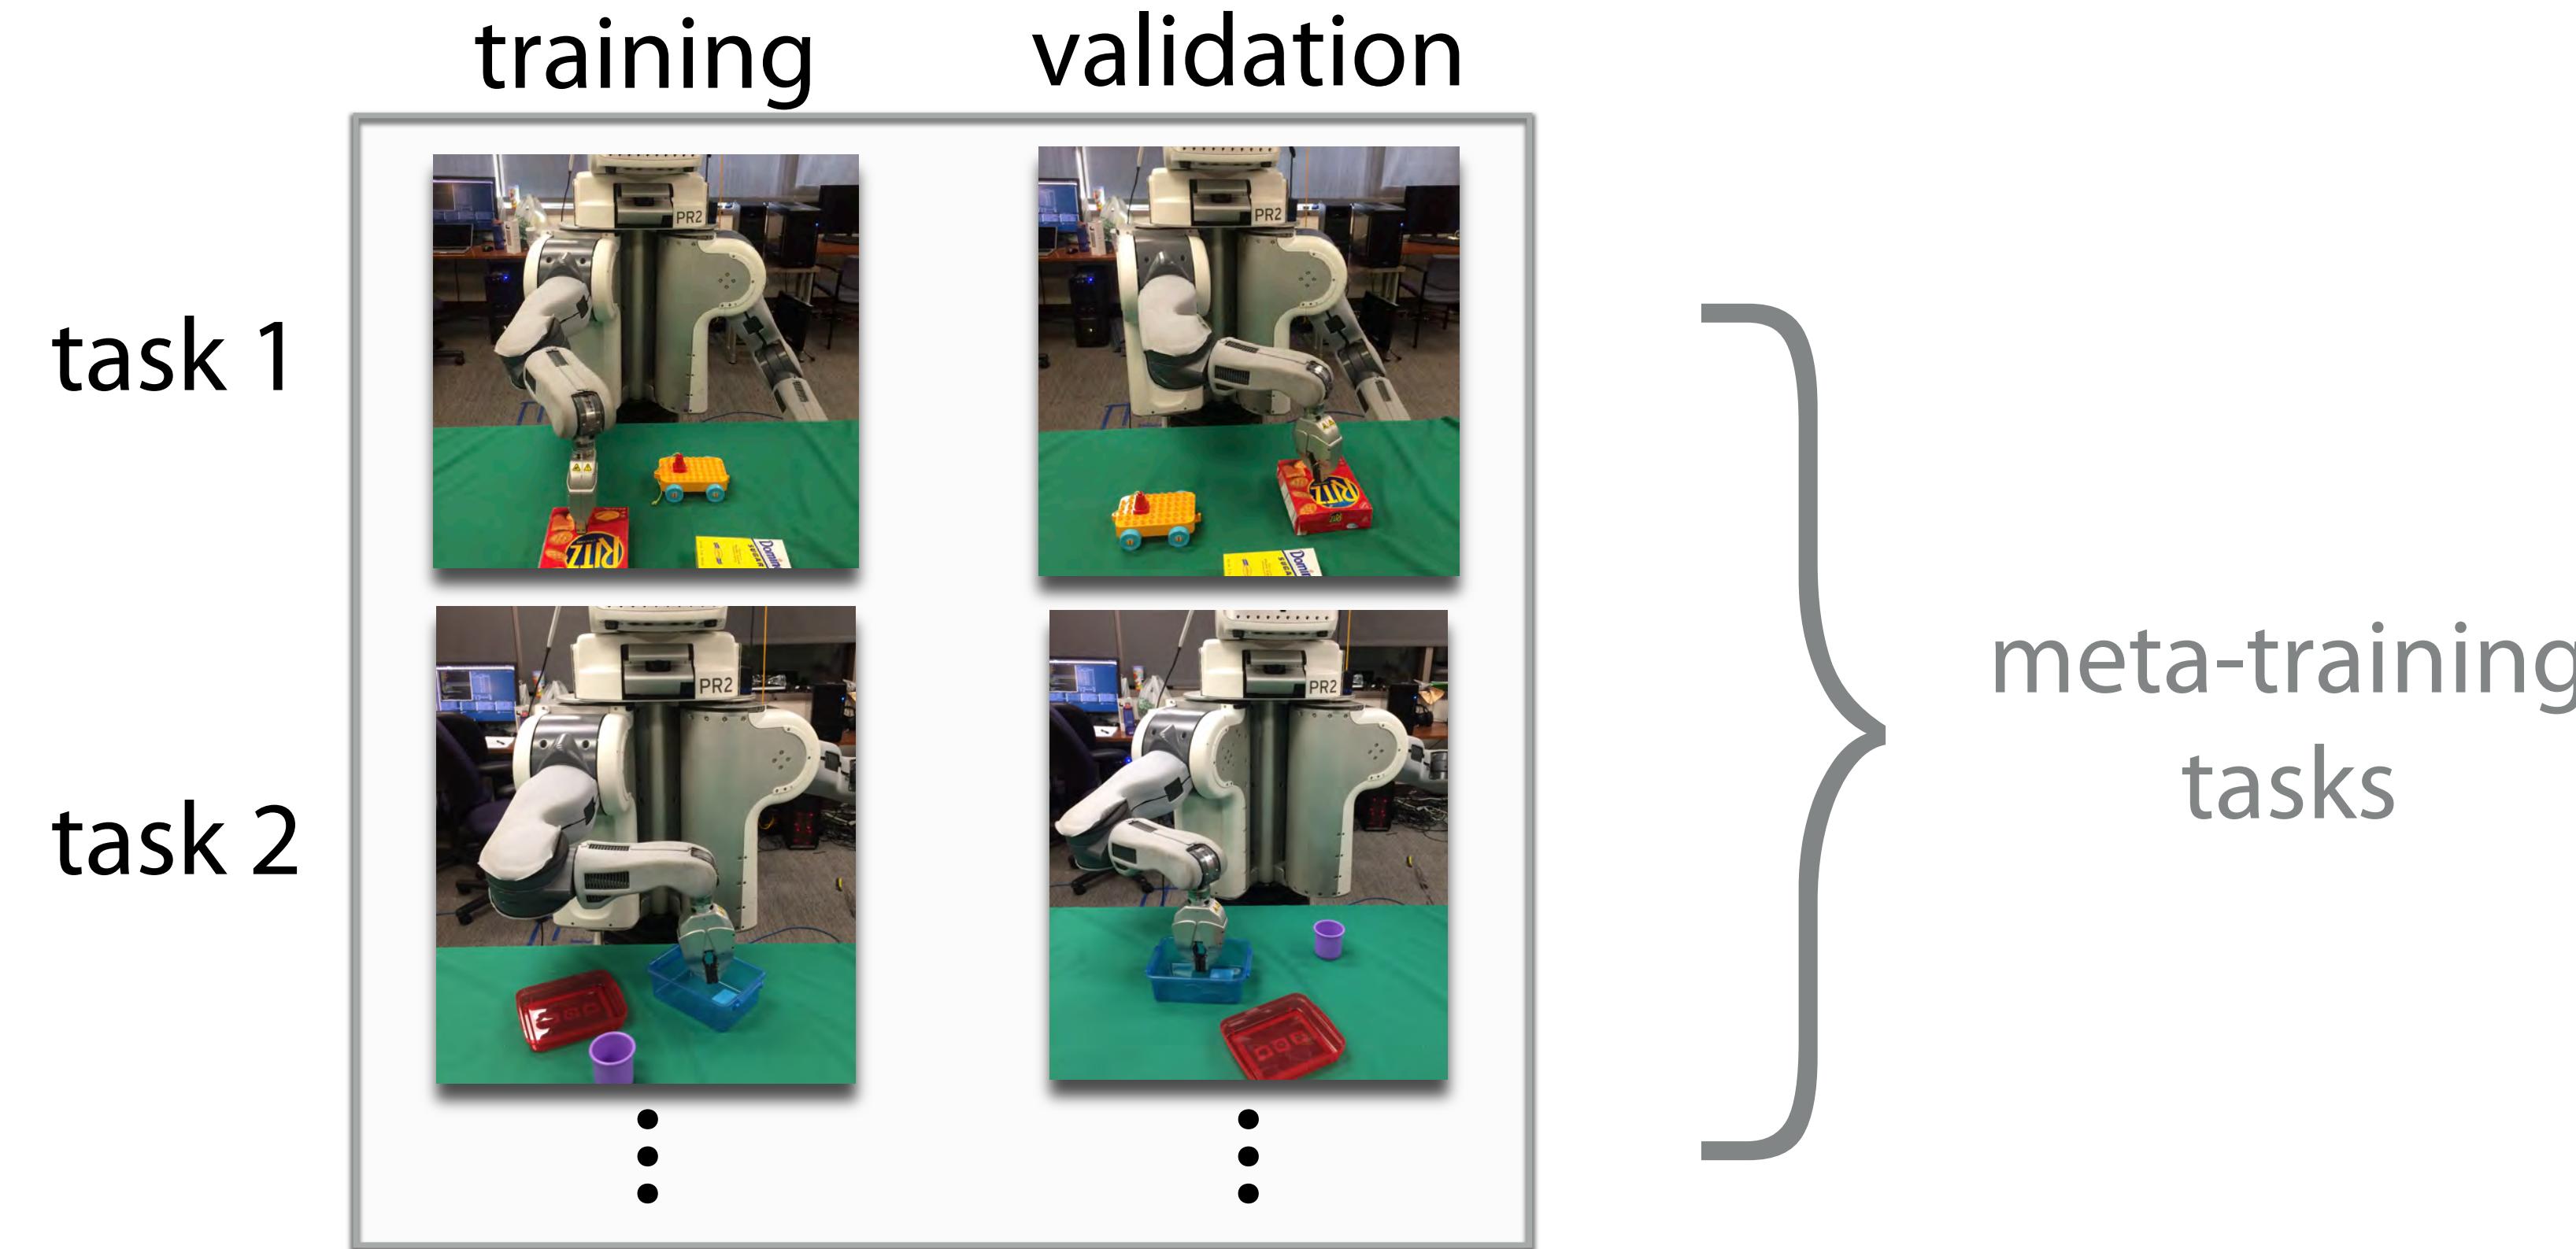


Rahmanizadeh et al. '17 Zhang et al. '17
learns from raw pixels,
but requires many demonstrations

Through meta-learning: reuse data from other tasks/objects/environments

One-Shot Visual Imitation Learning

How? Learn to learn the policy from one demonstration



Meta-test time:



Learn policy from one demo.

real-world placing *from pixels*

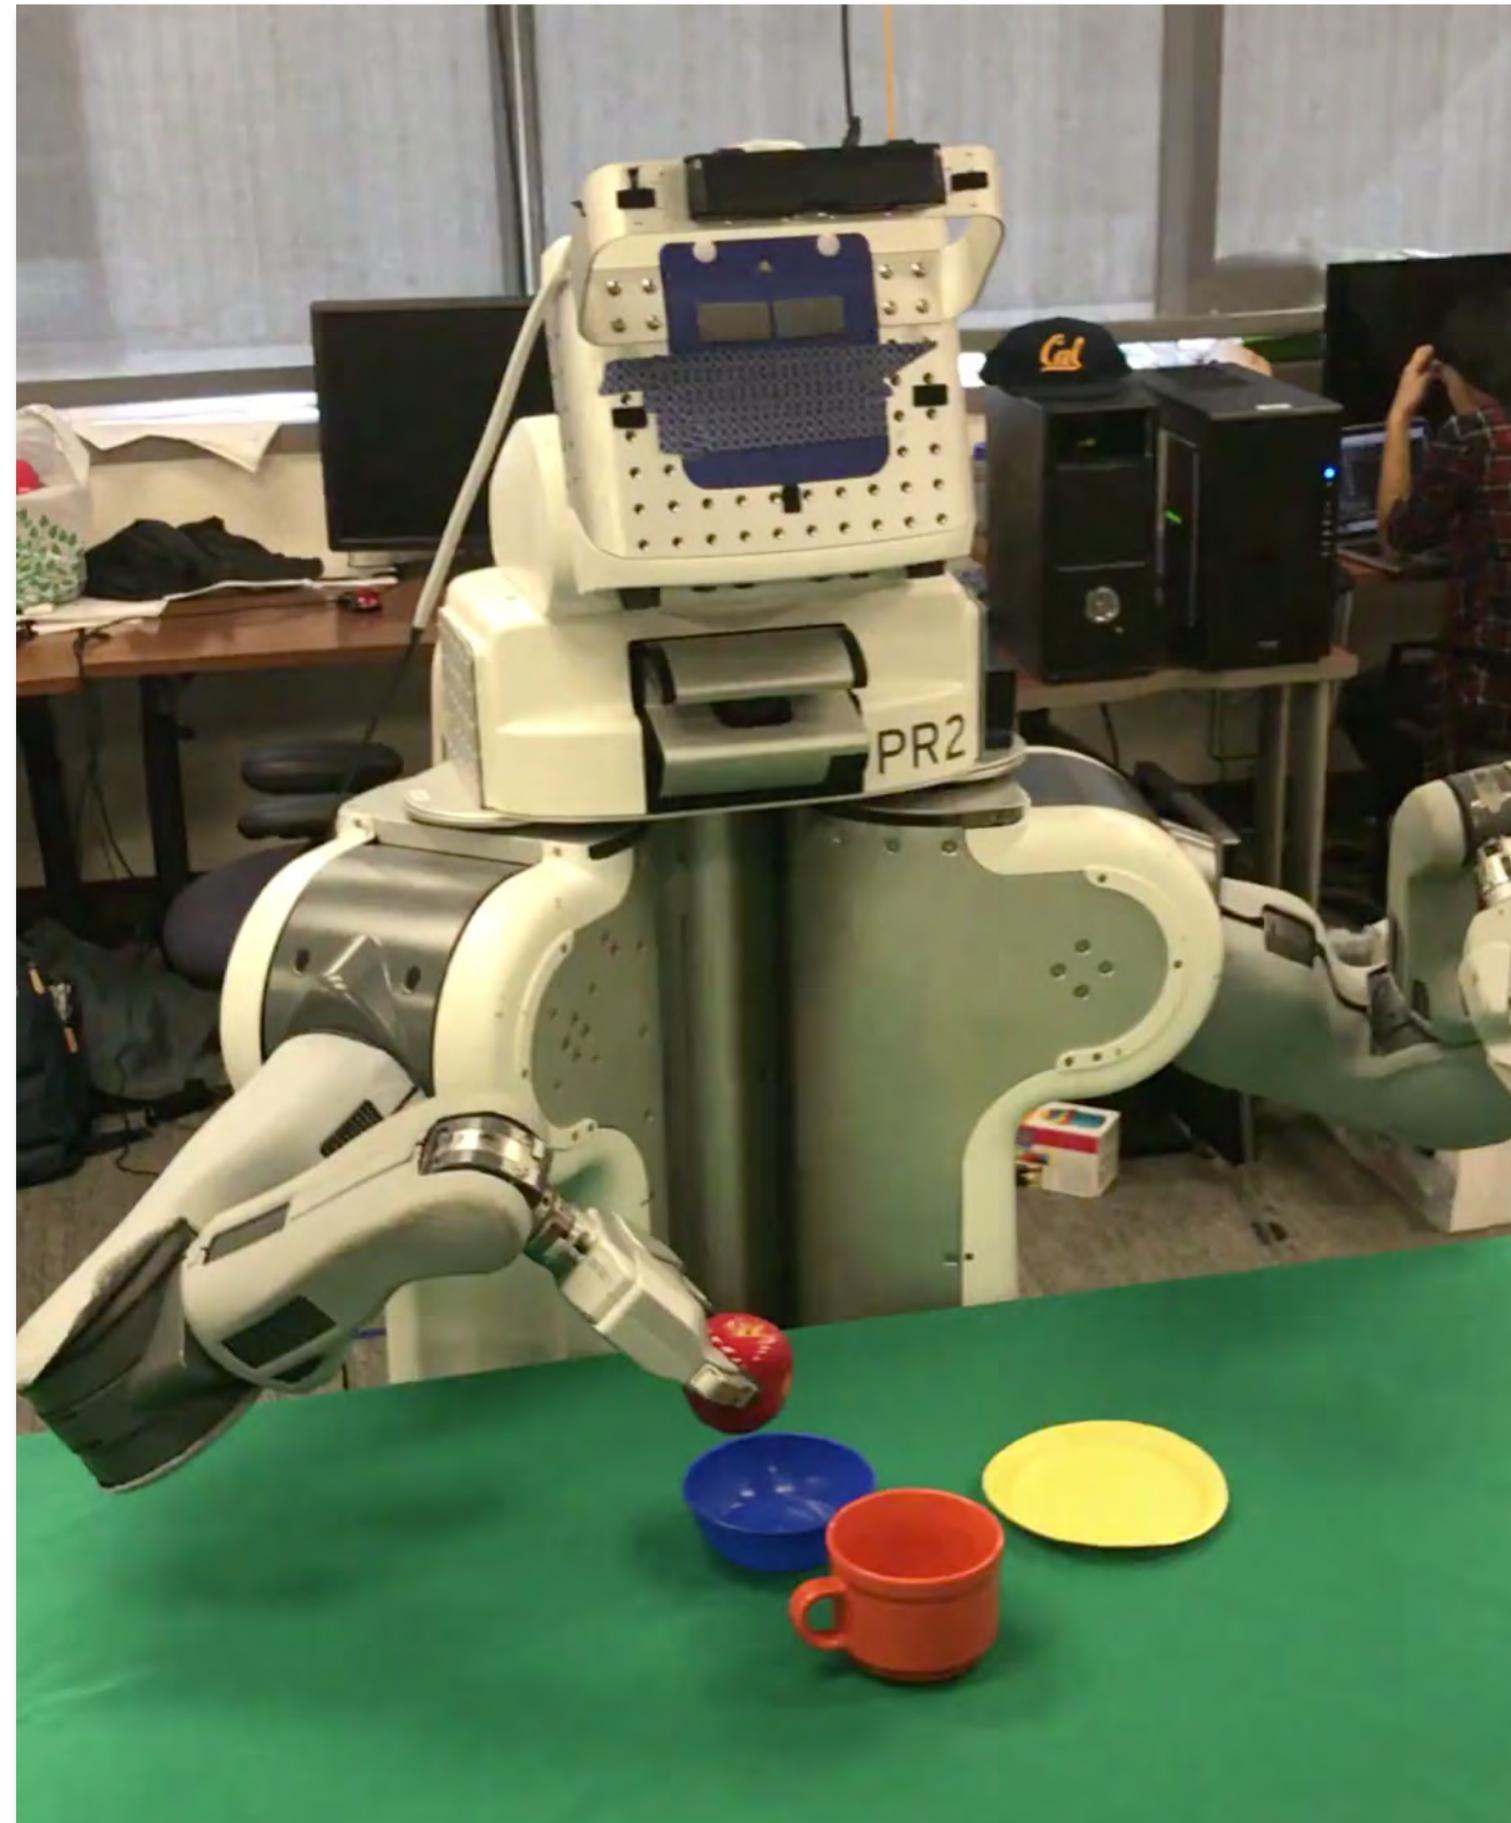


subset of
training objects

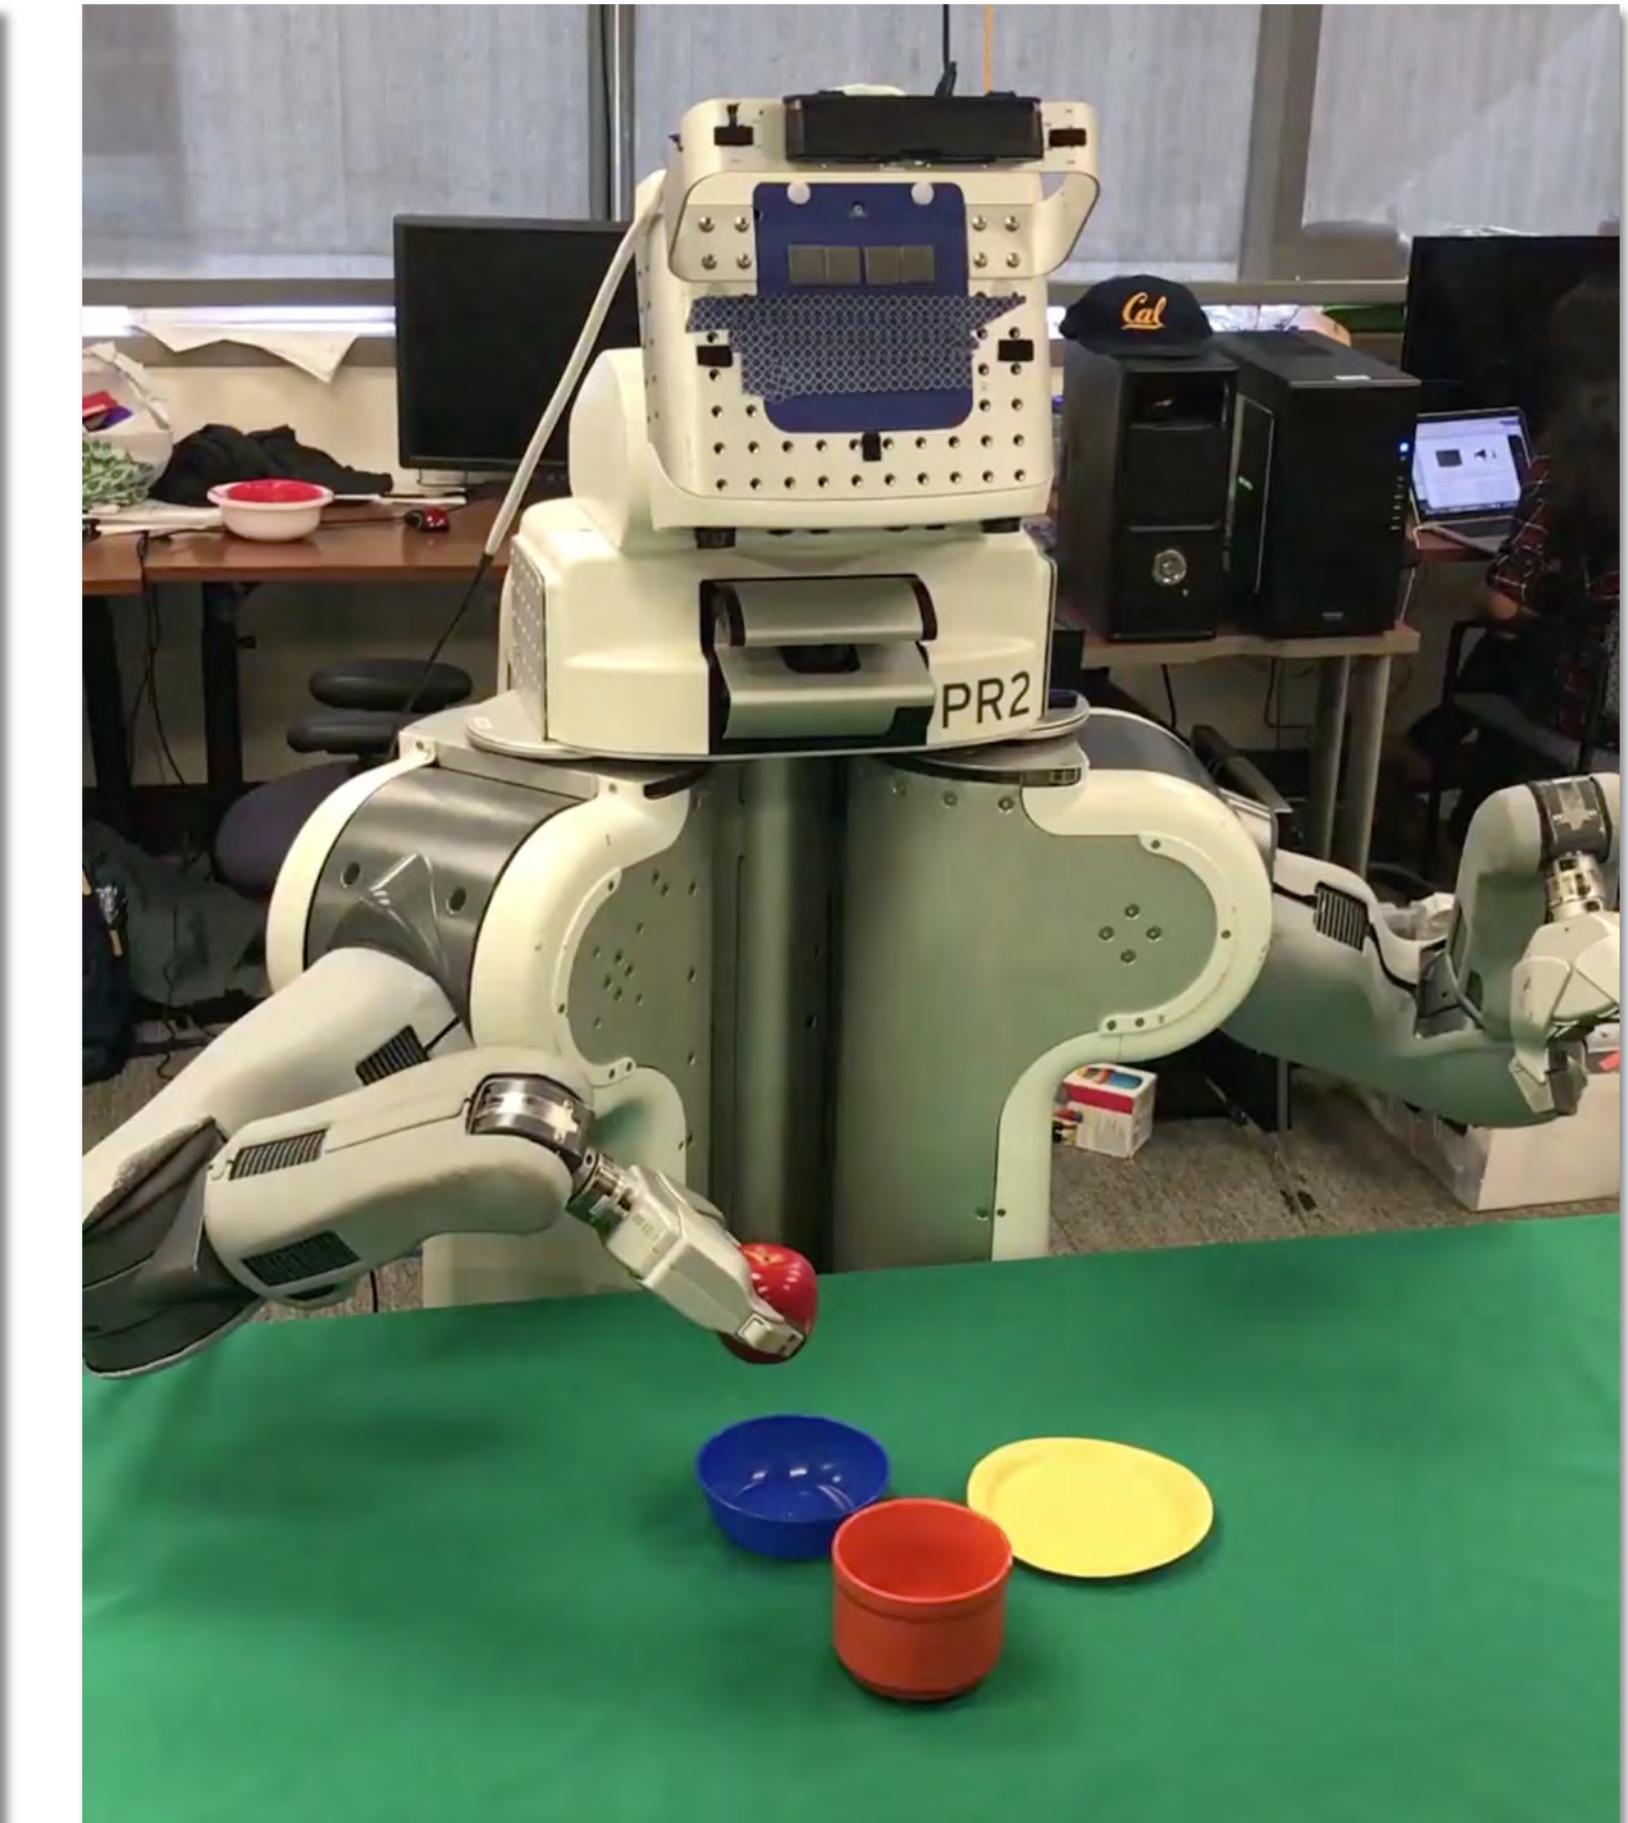


held-out test objects

input demo



resulting policy

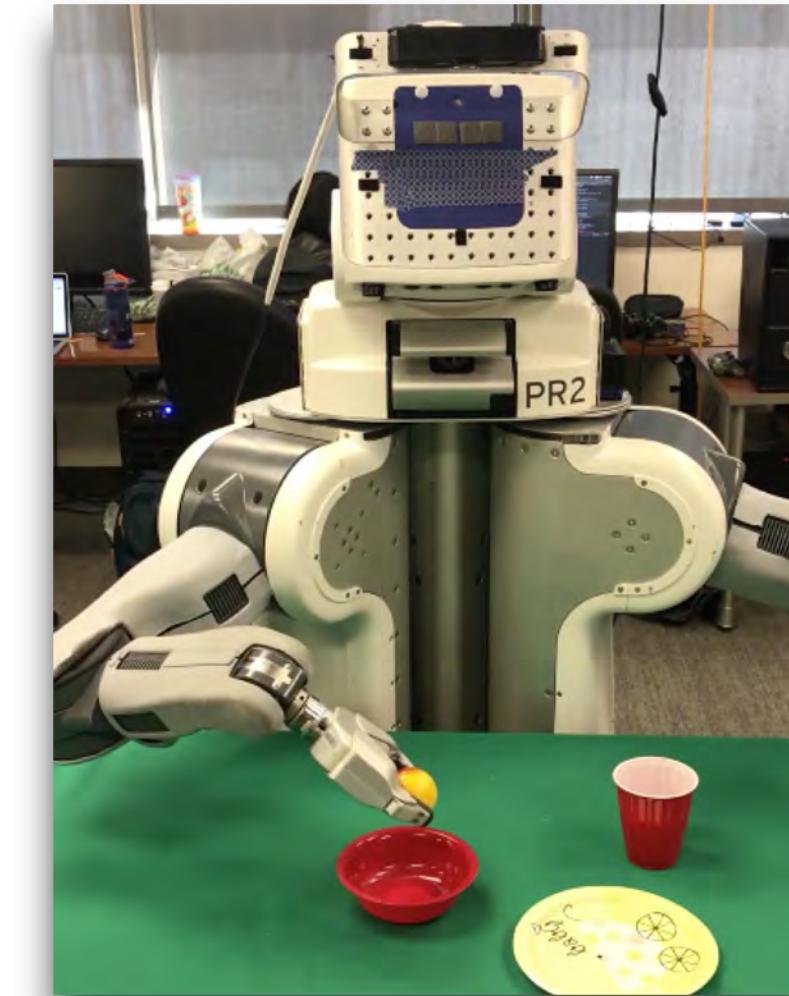


Few-Shot Learning *from Weak Supervision*

Given one teleoperated demonstration:
Given a video of a human:



Learn a policy.



Yu*, Finn*, Xie, Dasari, Zhang,
Abbeel, Levine RSS '18

Given 1 example of 5 classes:



Classify new examples



Grant, Finn, Peterson, Abbott, Levine,
Darrell, Griffiths NIPS CIAI Workshop '17

Learning to Learn from Weak Supervision

meta-training

$$\min_{\theta} \sum_{\text{task } i} \mathcal{L}_{\text{test}}^i(\theta - \alpha \nabla_{\theta} \mathcal{L}_{\text{train}}^i(\theta))$$

fully supervised

meta-test

$$\theta' \leftarrow \theta - \alpha \nabla_{\theta} \mathcal{L}(\theta)$$

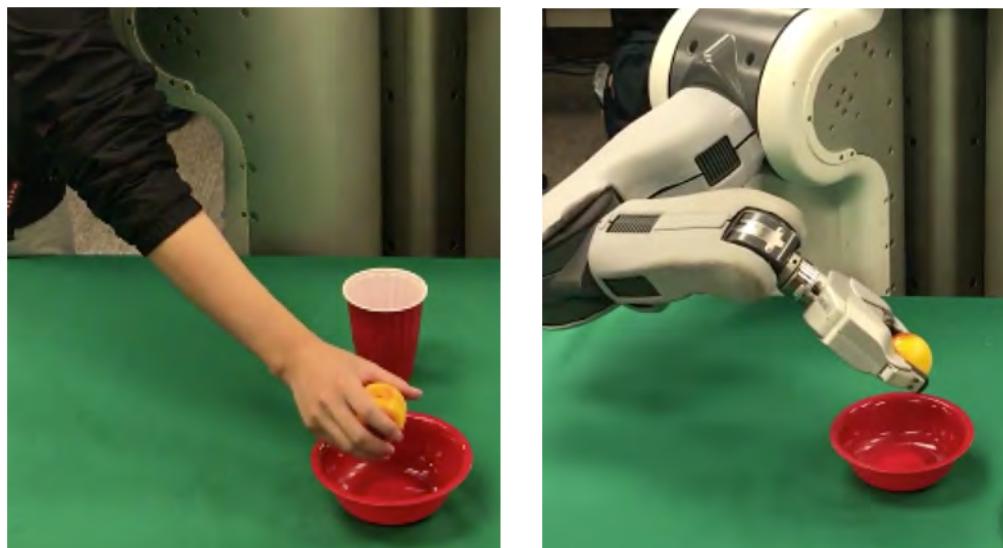
weakly supervised

weakly supervised

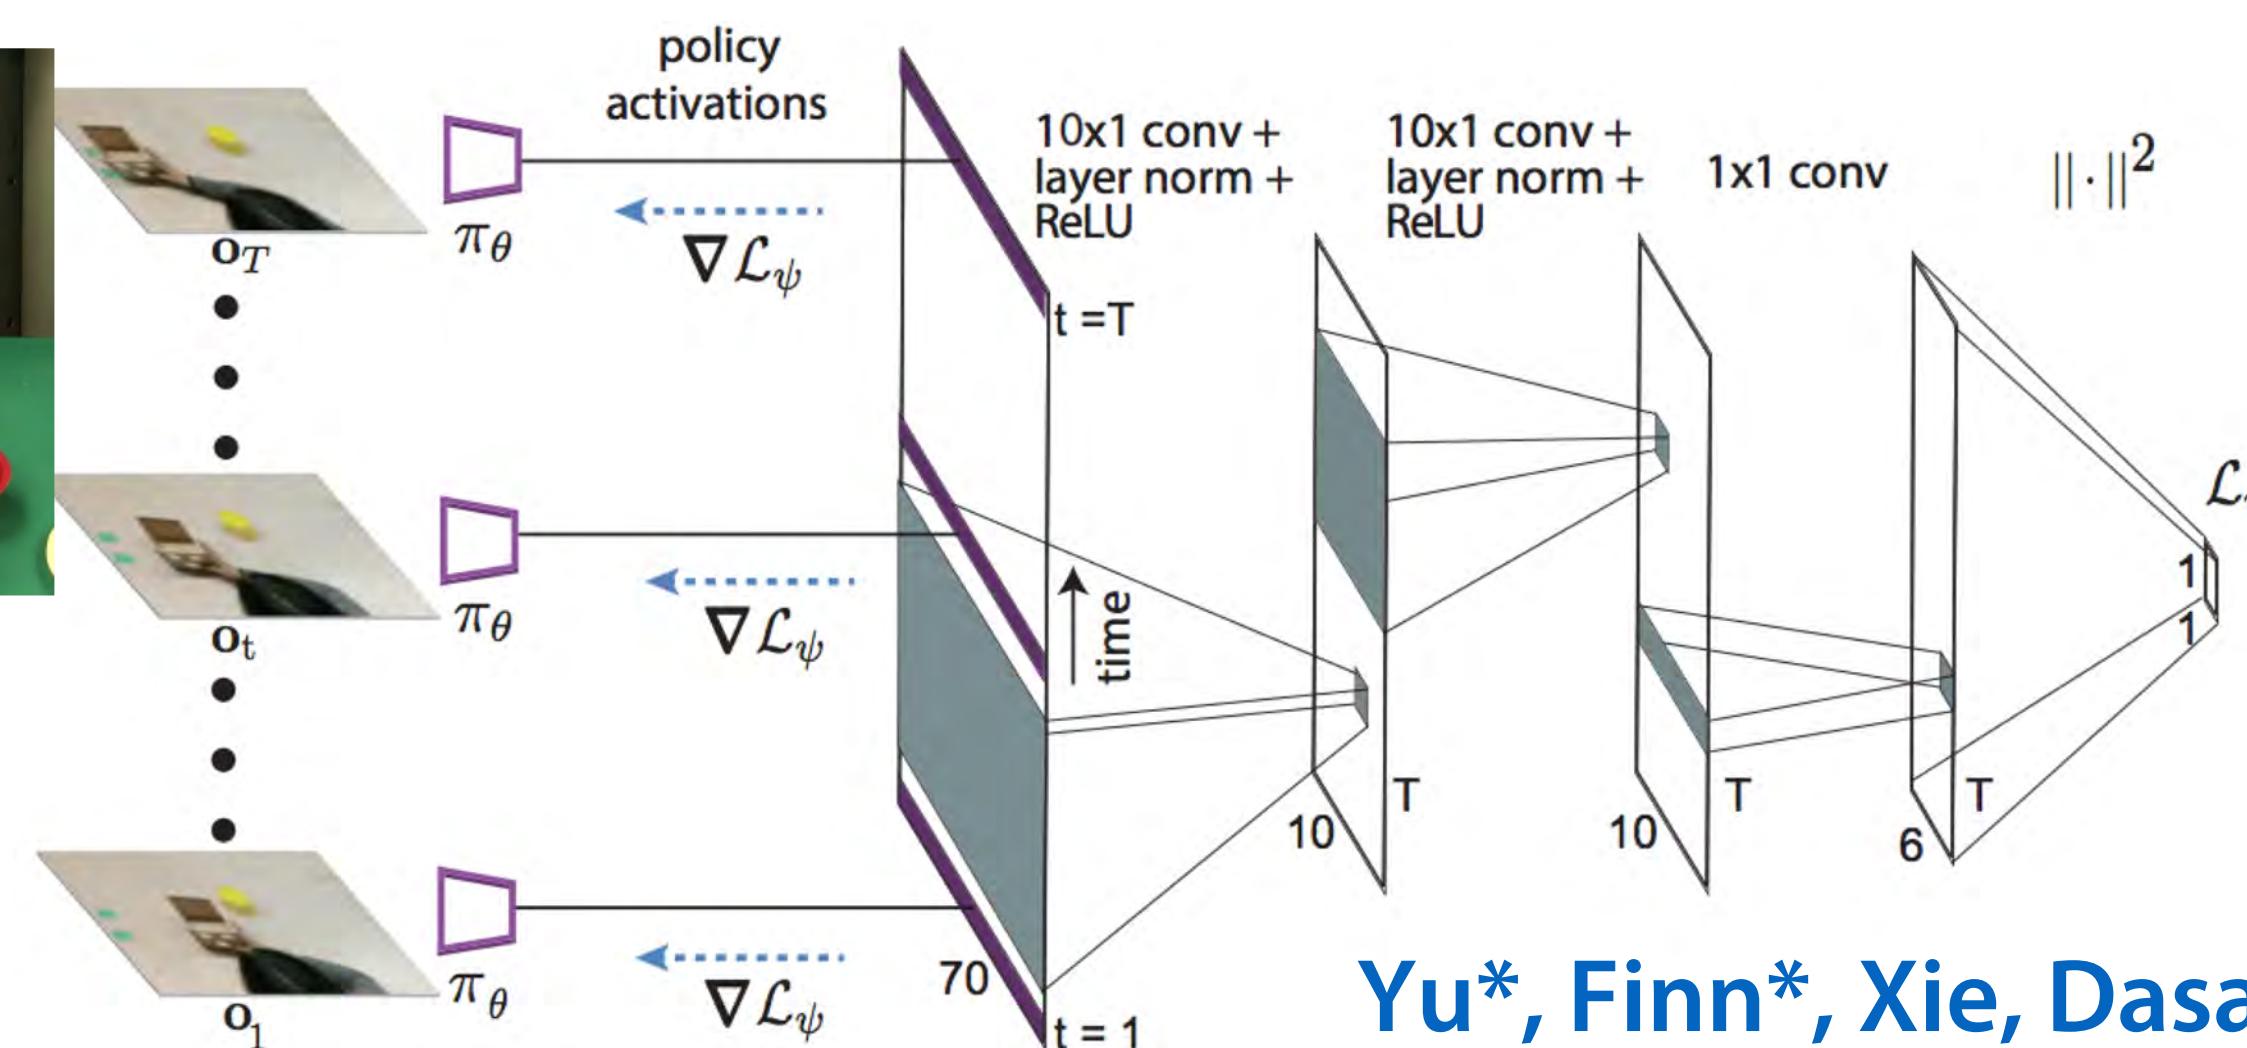
What if the weakly supervised loss is unavailable?

imitation loss

$$\mathcal{L} = \sum_t \|\pi_{\theta}(o_t) - a_t^*\|^2$$



$$\min_{\theta, \psi} \sum_{\text{task } i} \mathcal{L}_{\text{test}}^i(\theta - \alpha \nabla_{\theta} \mathcal{L}_{\psi}^i(\theta))$$

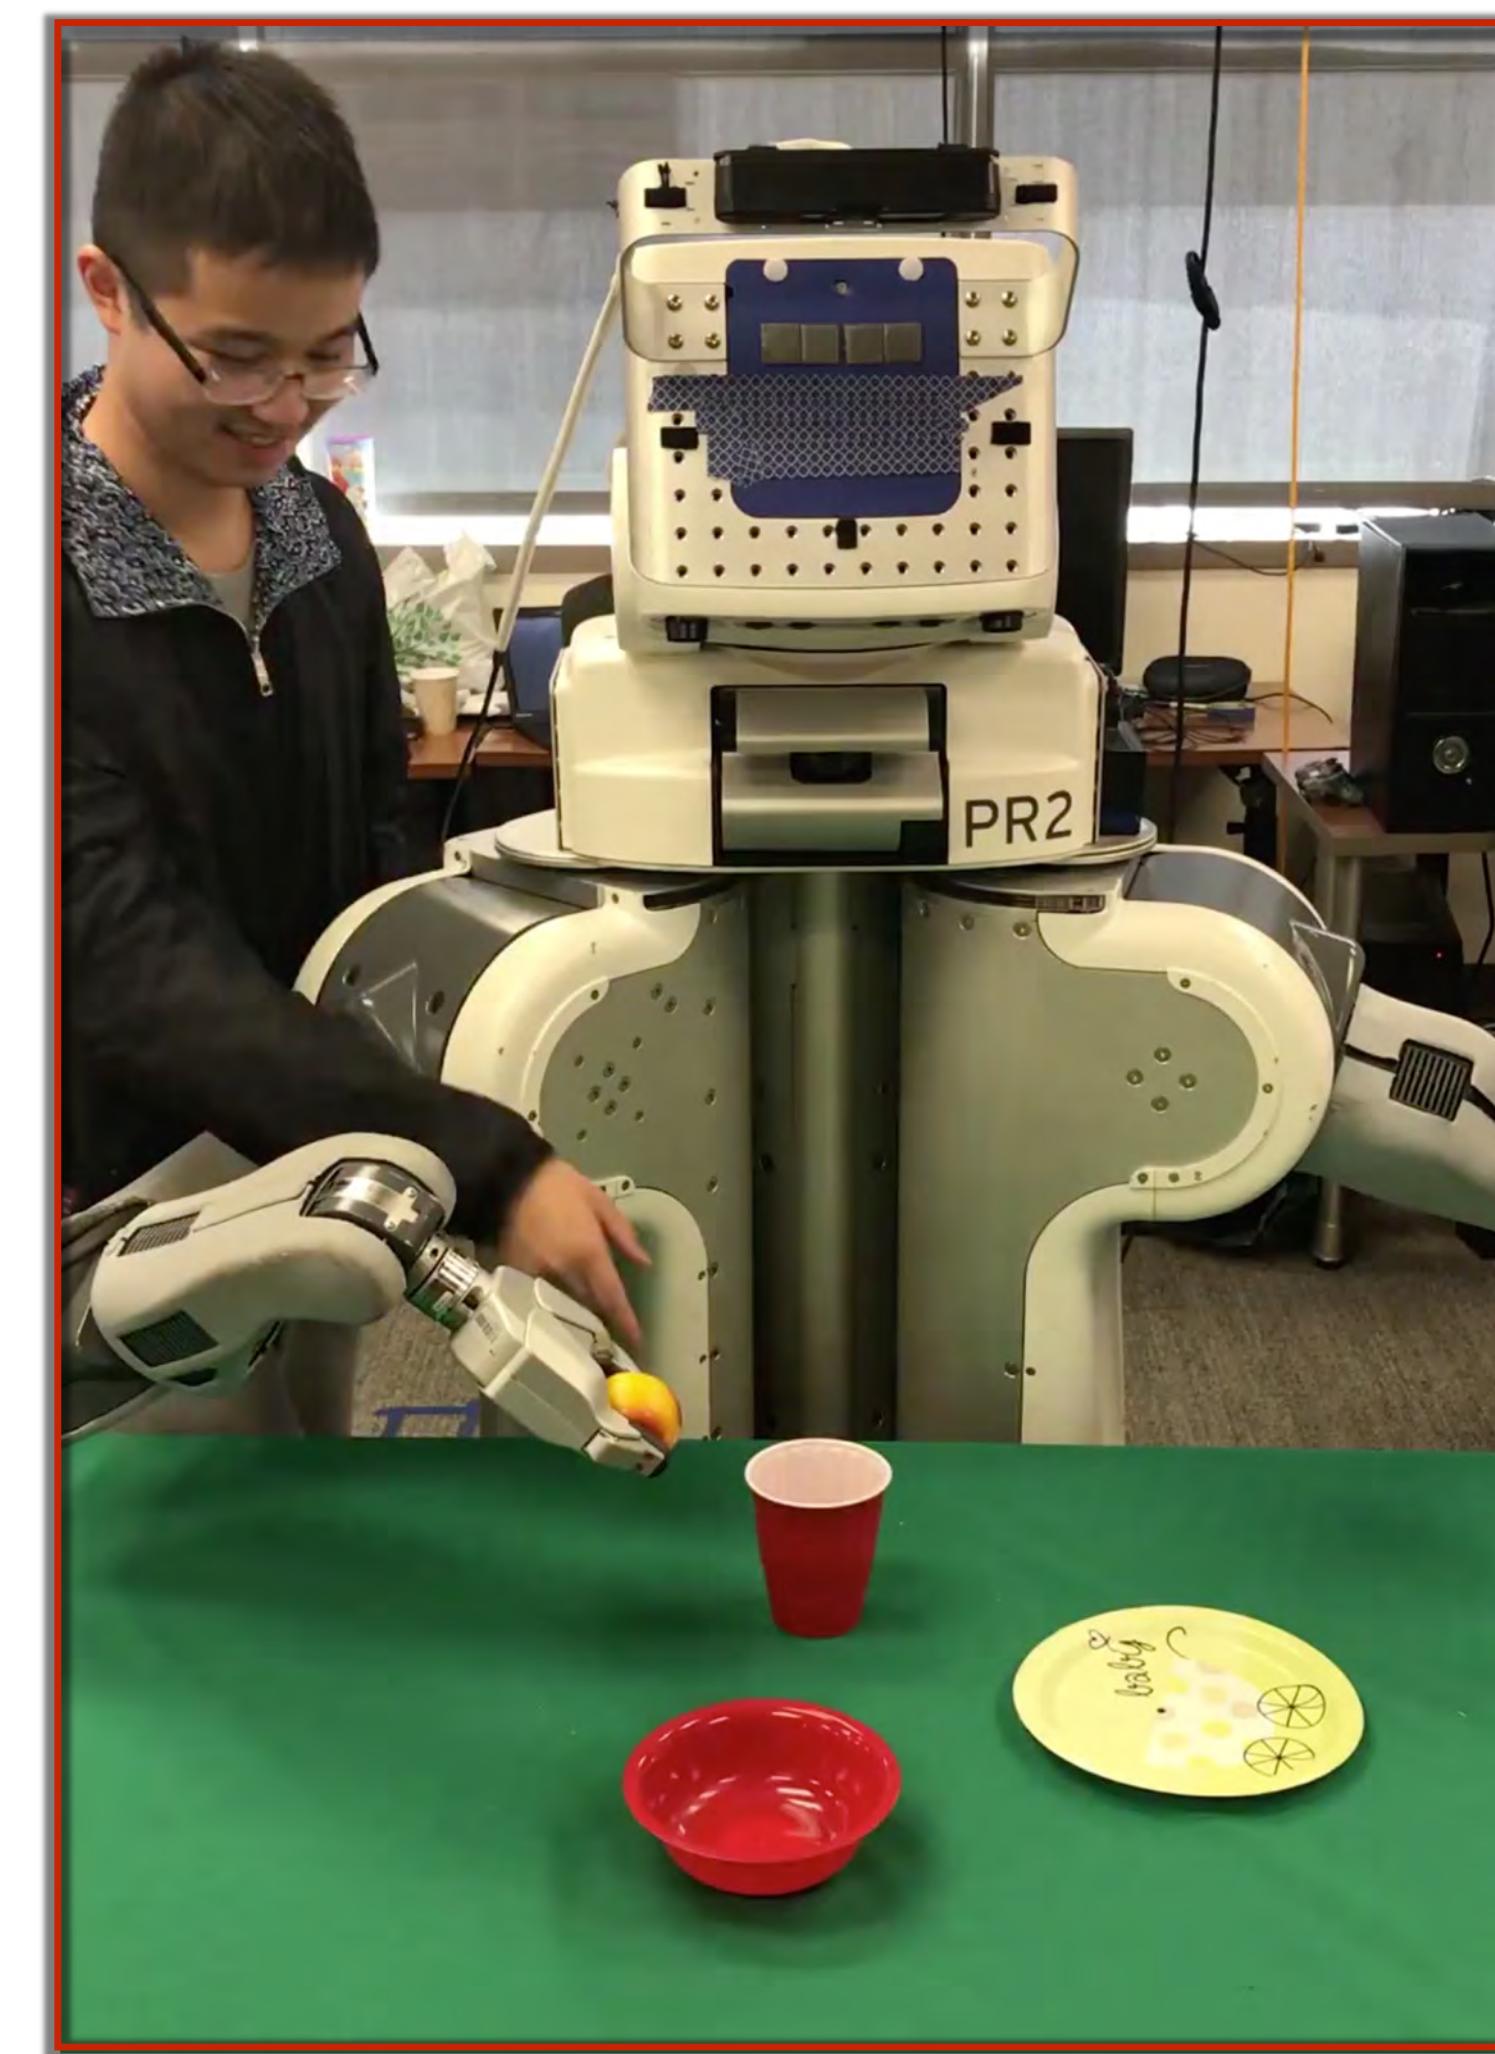


placing *from pixels*

input demo



resulting policy

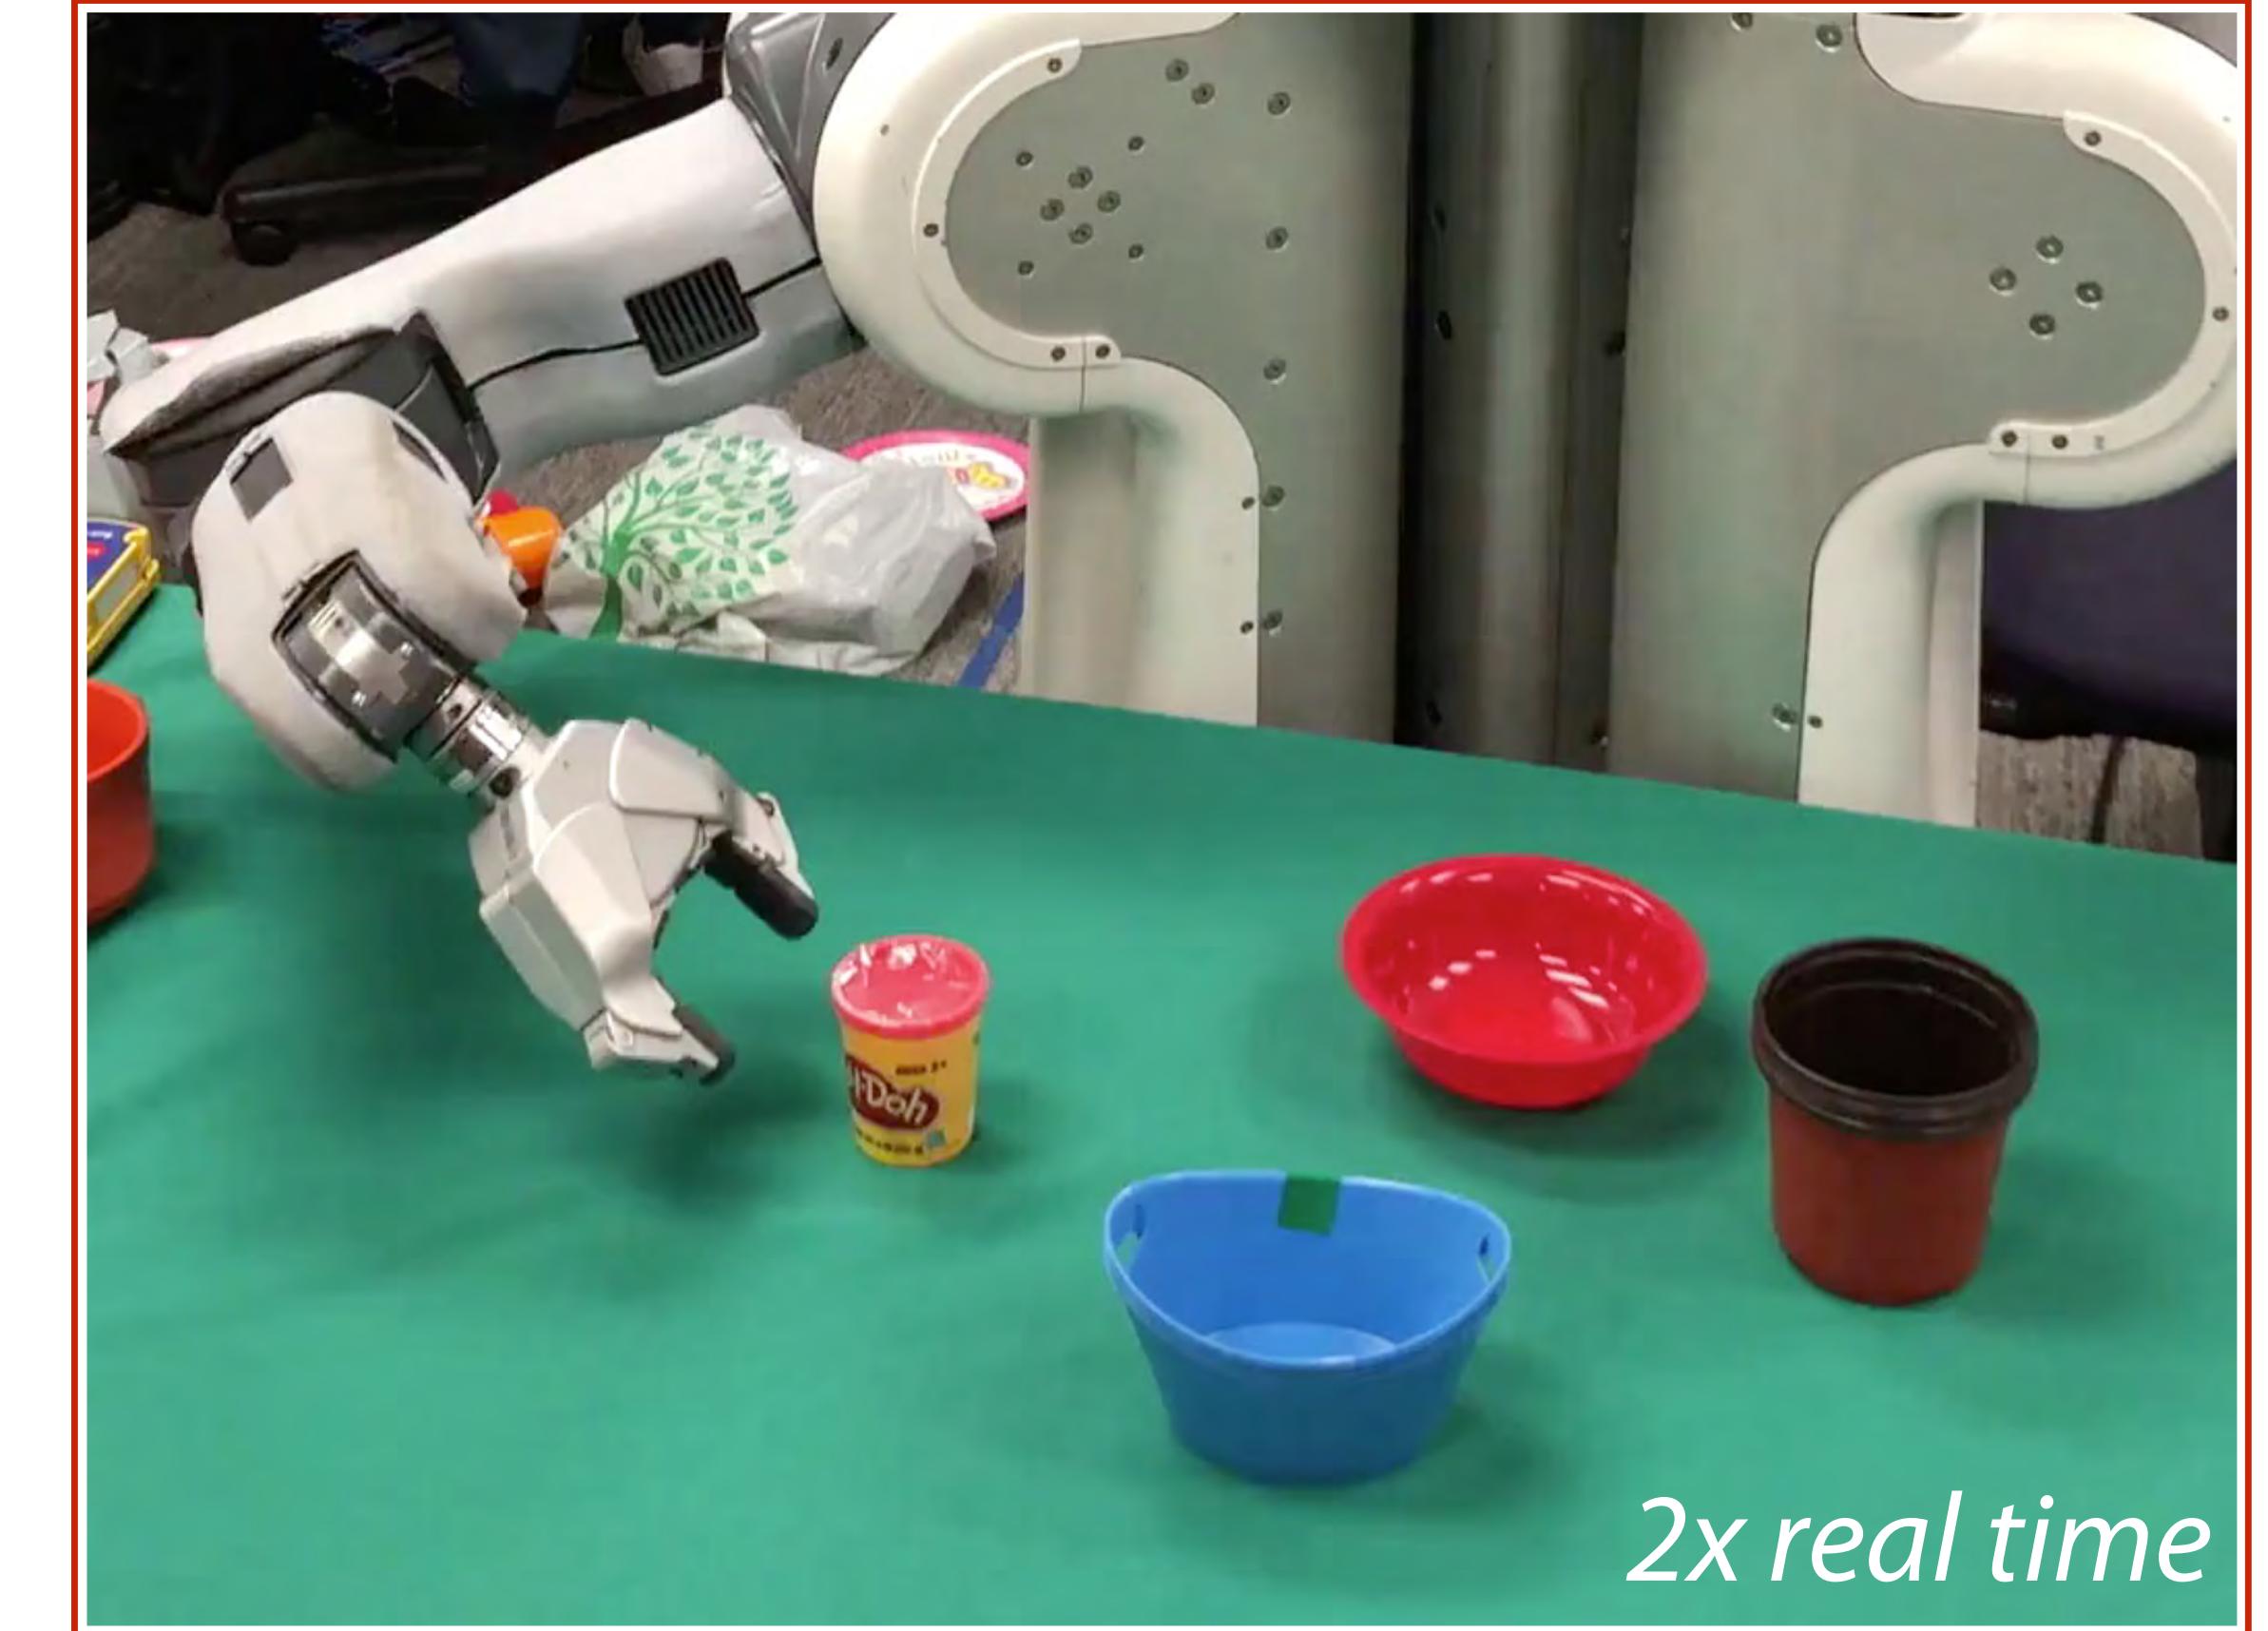


pick-and-place *from pixels*

input demo

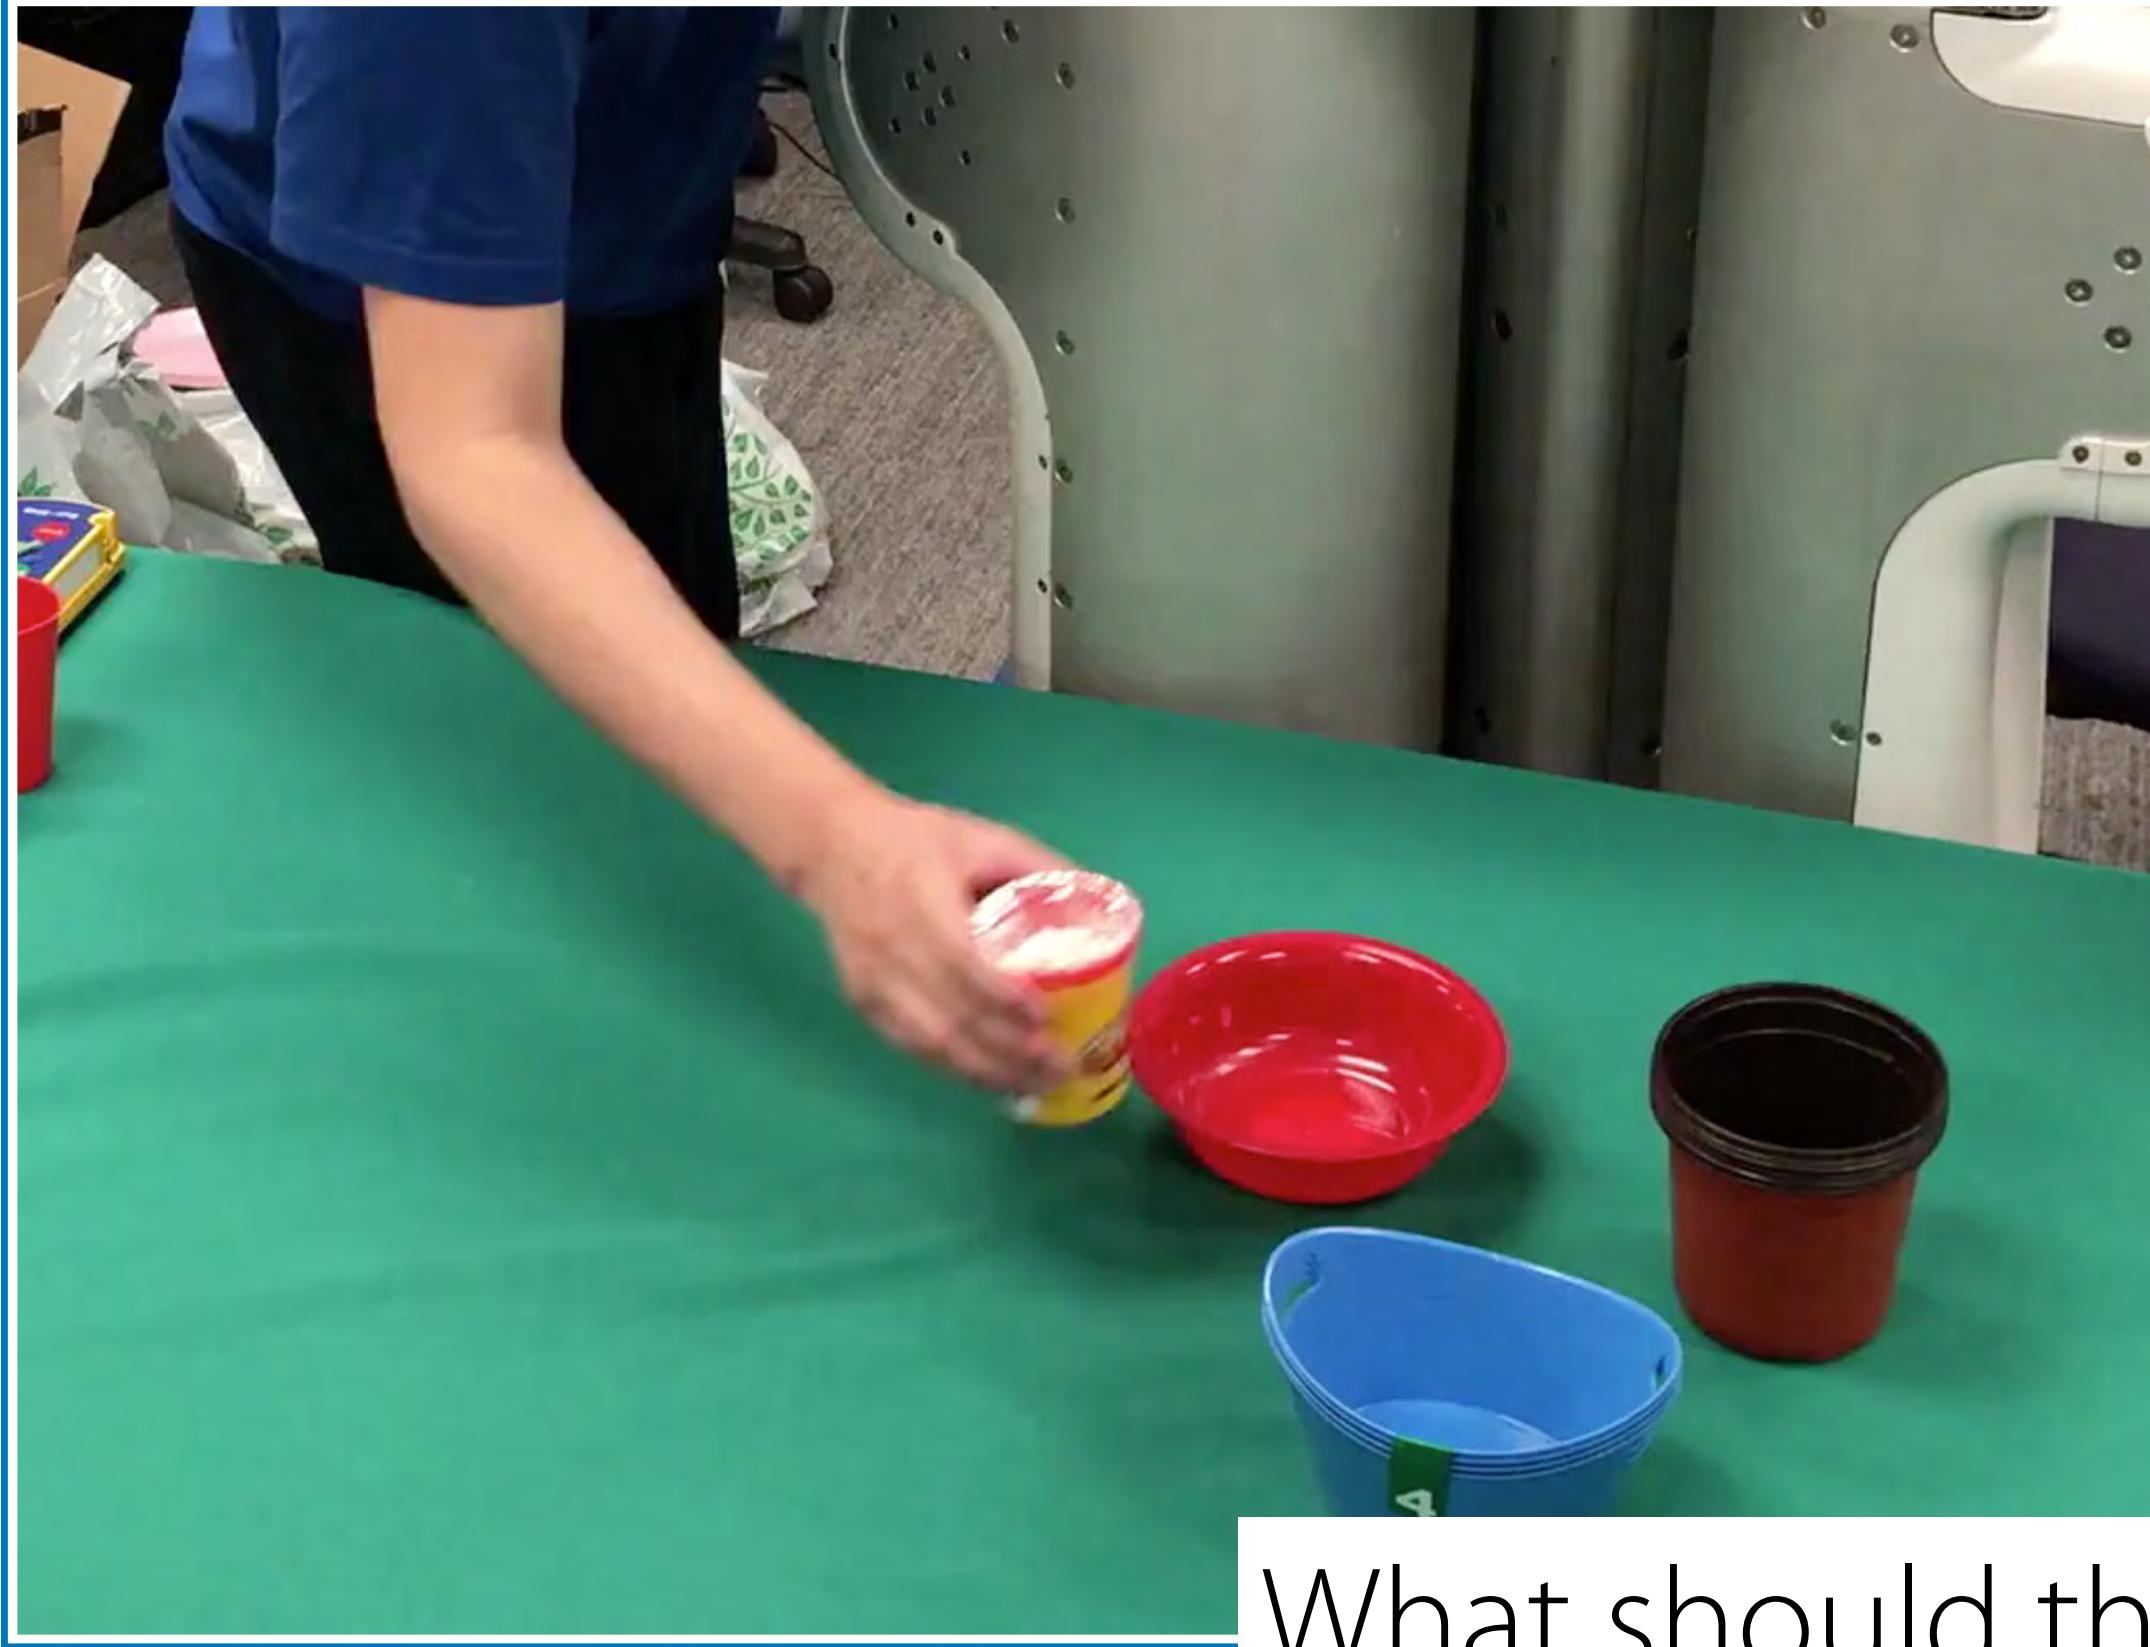


resulting policy

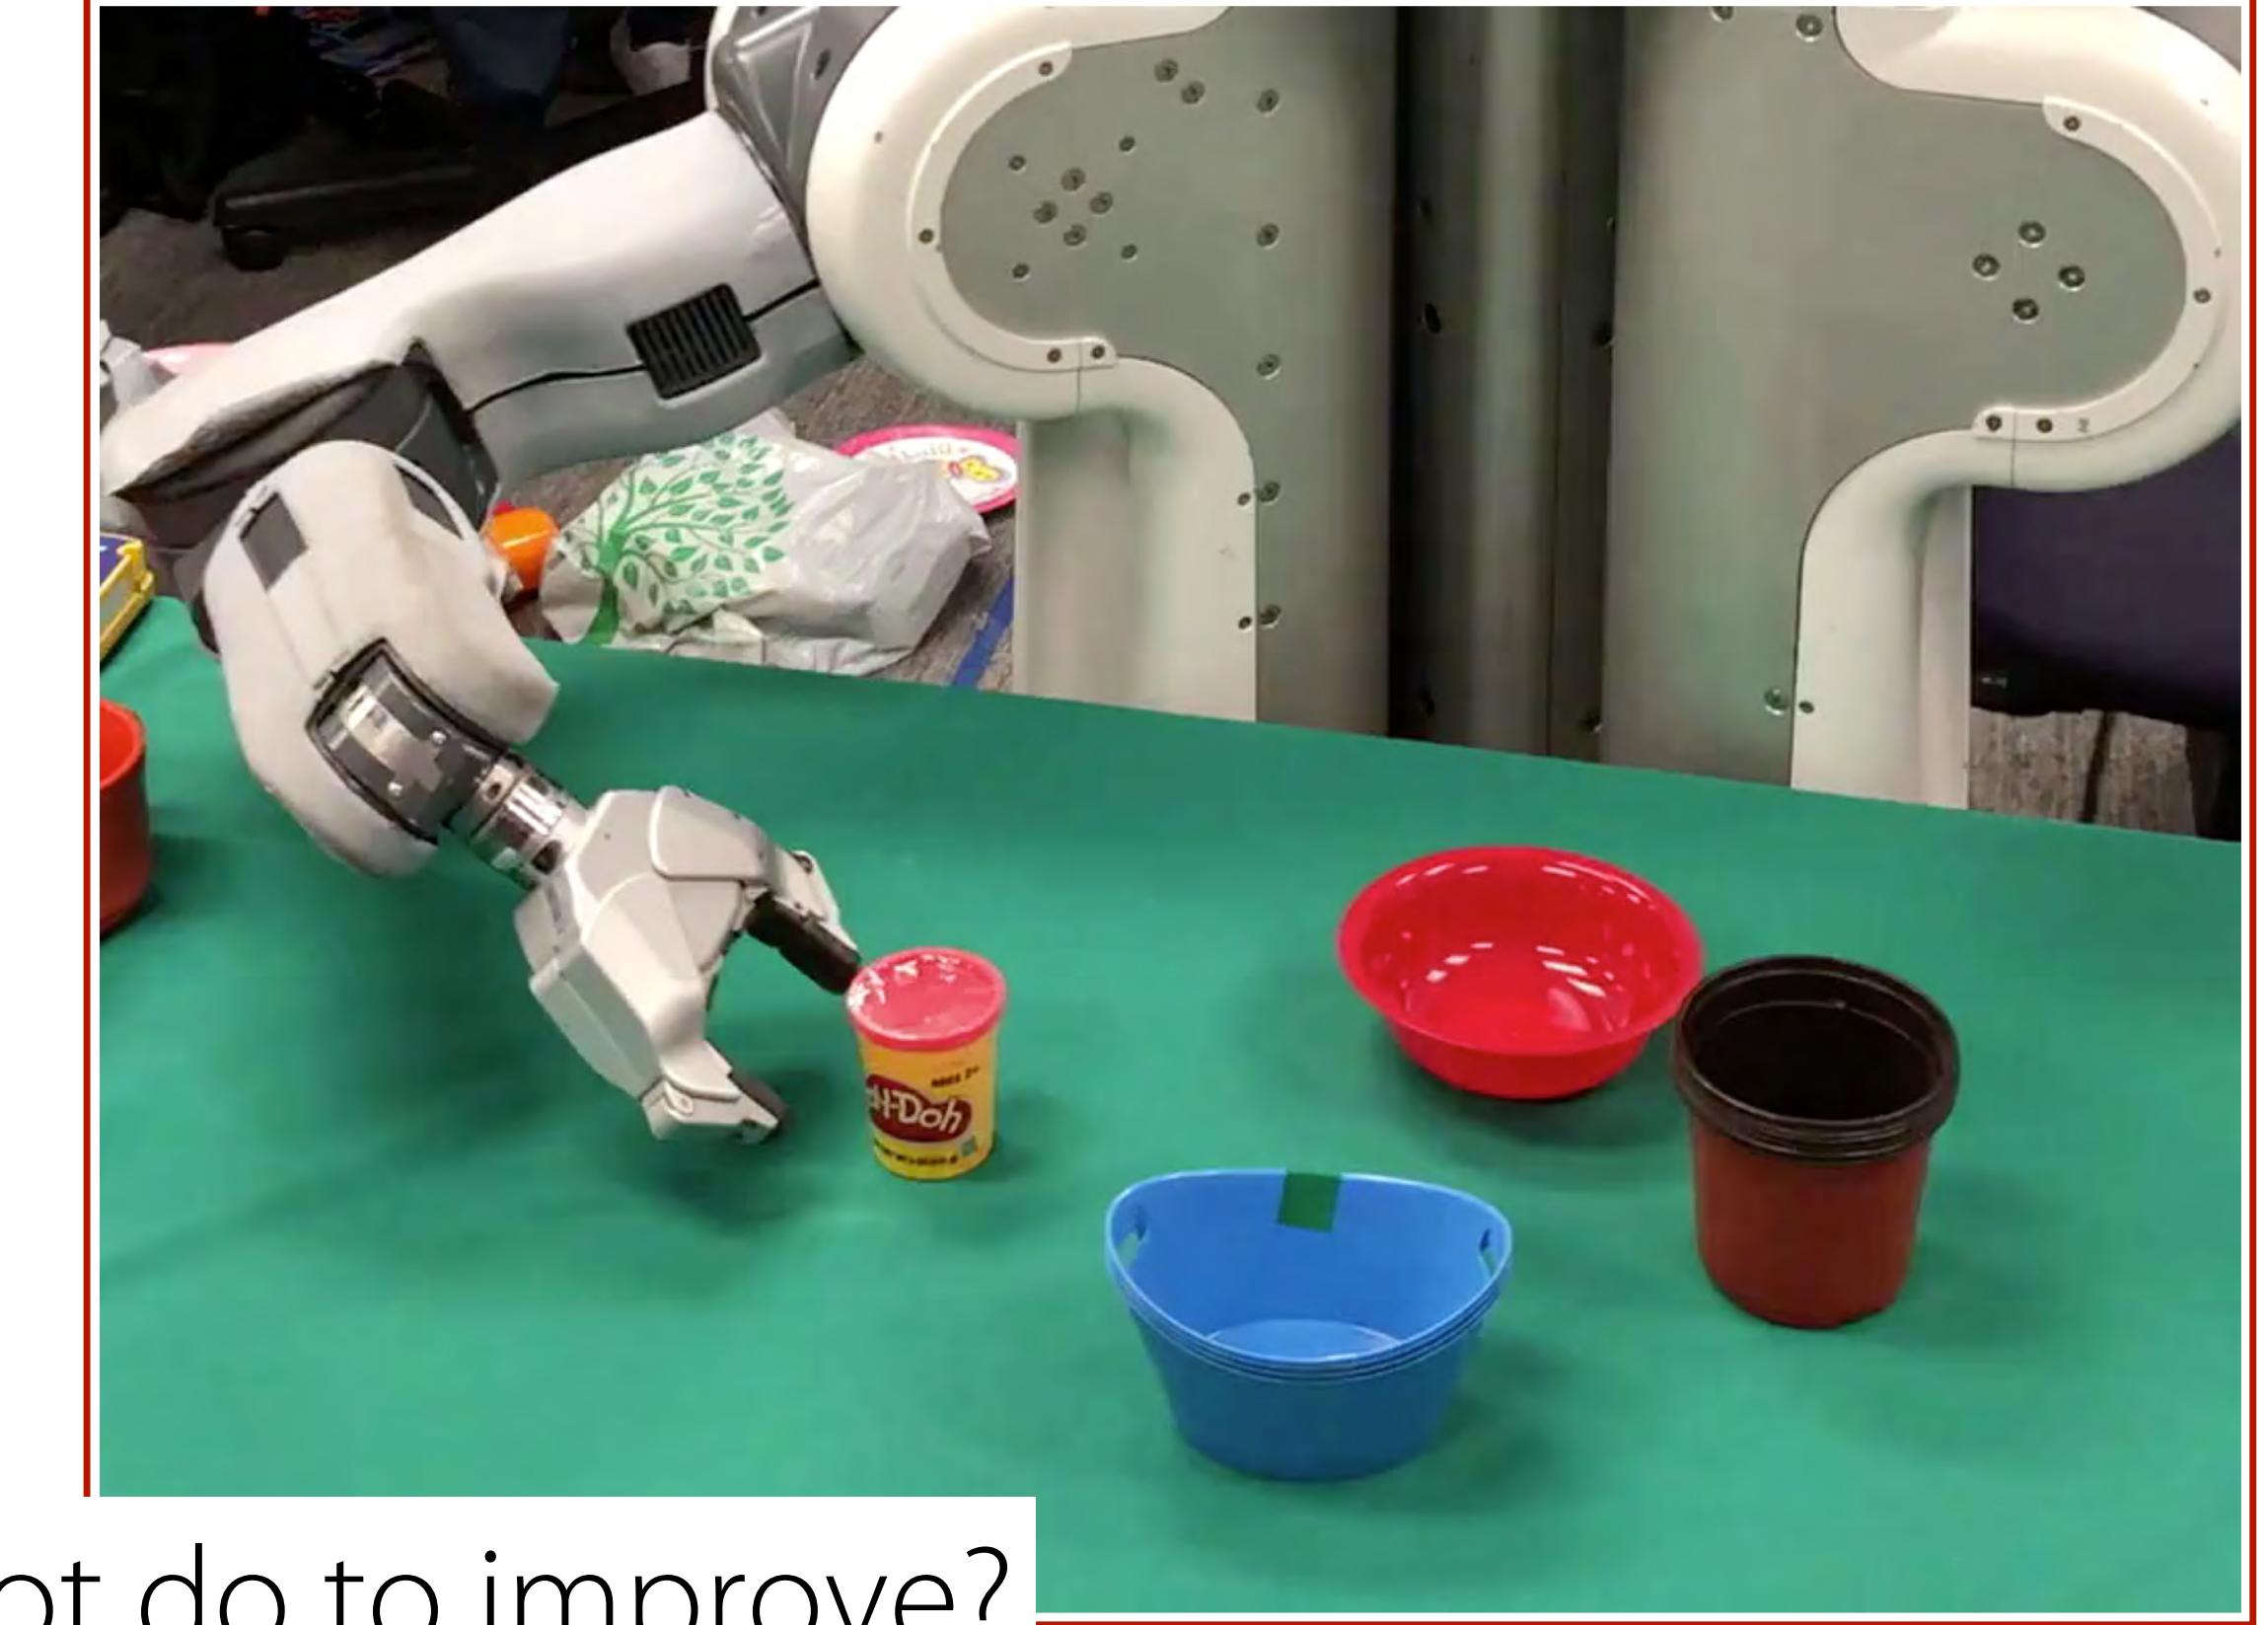


pick-and-place *from pixels*

input demo



resulting policy



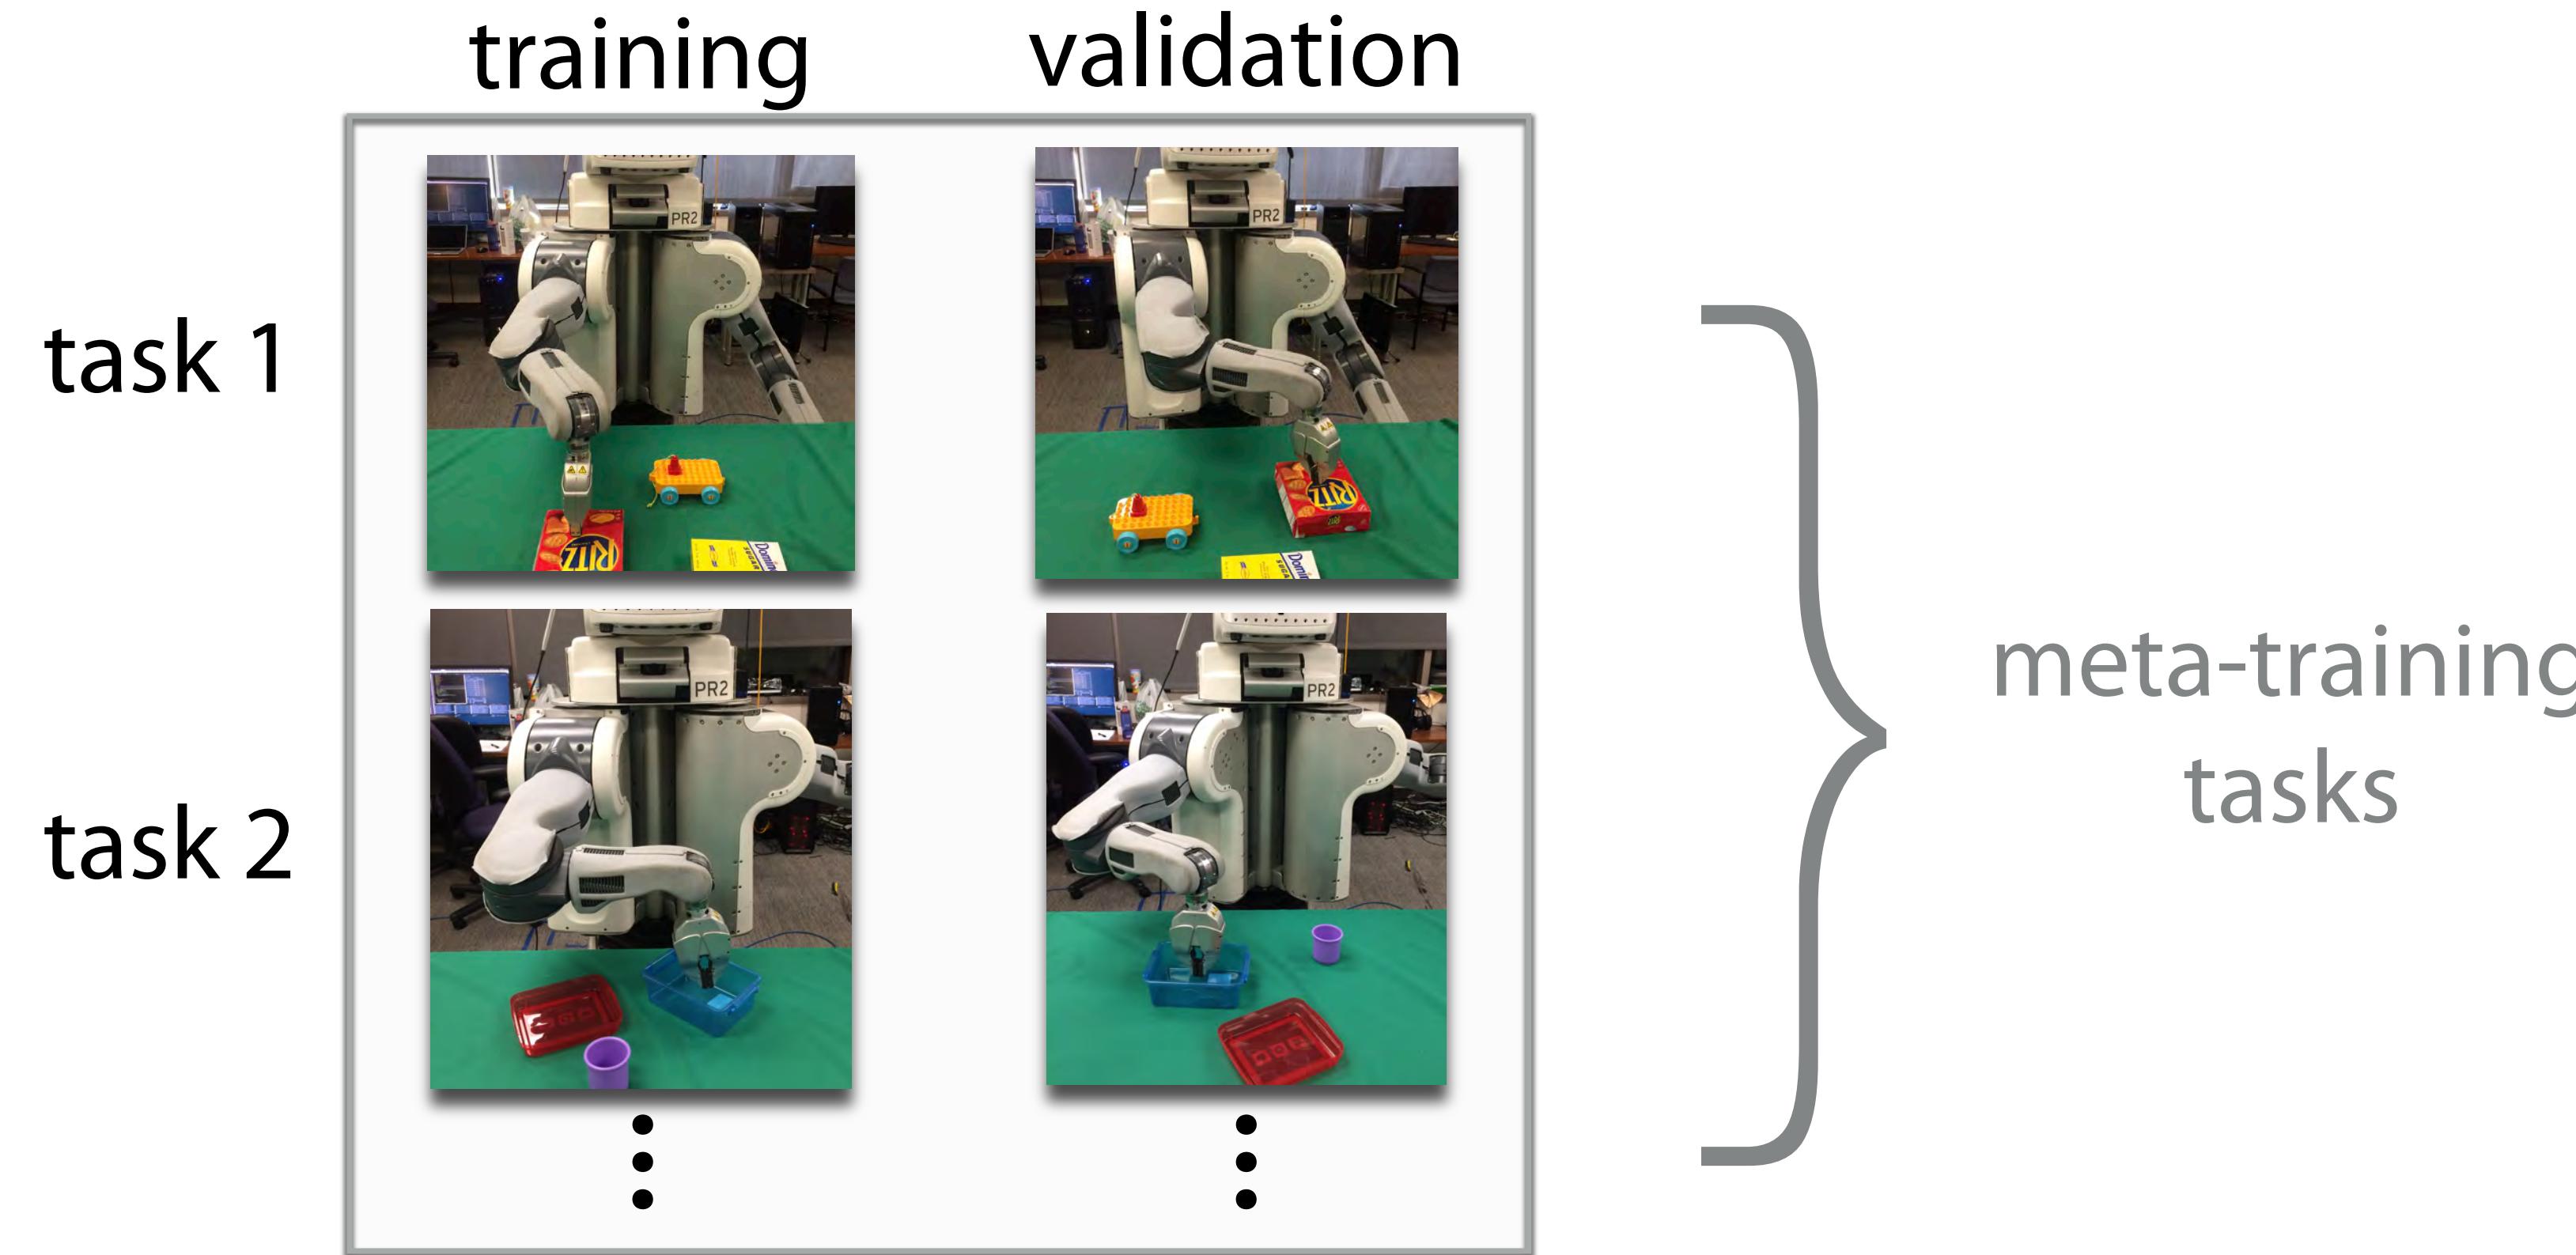
What should the robot do to improve?

Error in **inferring the goal** or **executing on it**?

Can we decouple **goal inference** & **policy learning**?

Meta-Learning Approach

How? Learn to learn the policy from one demonstration



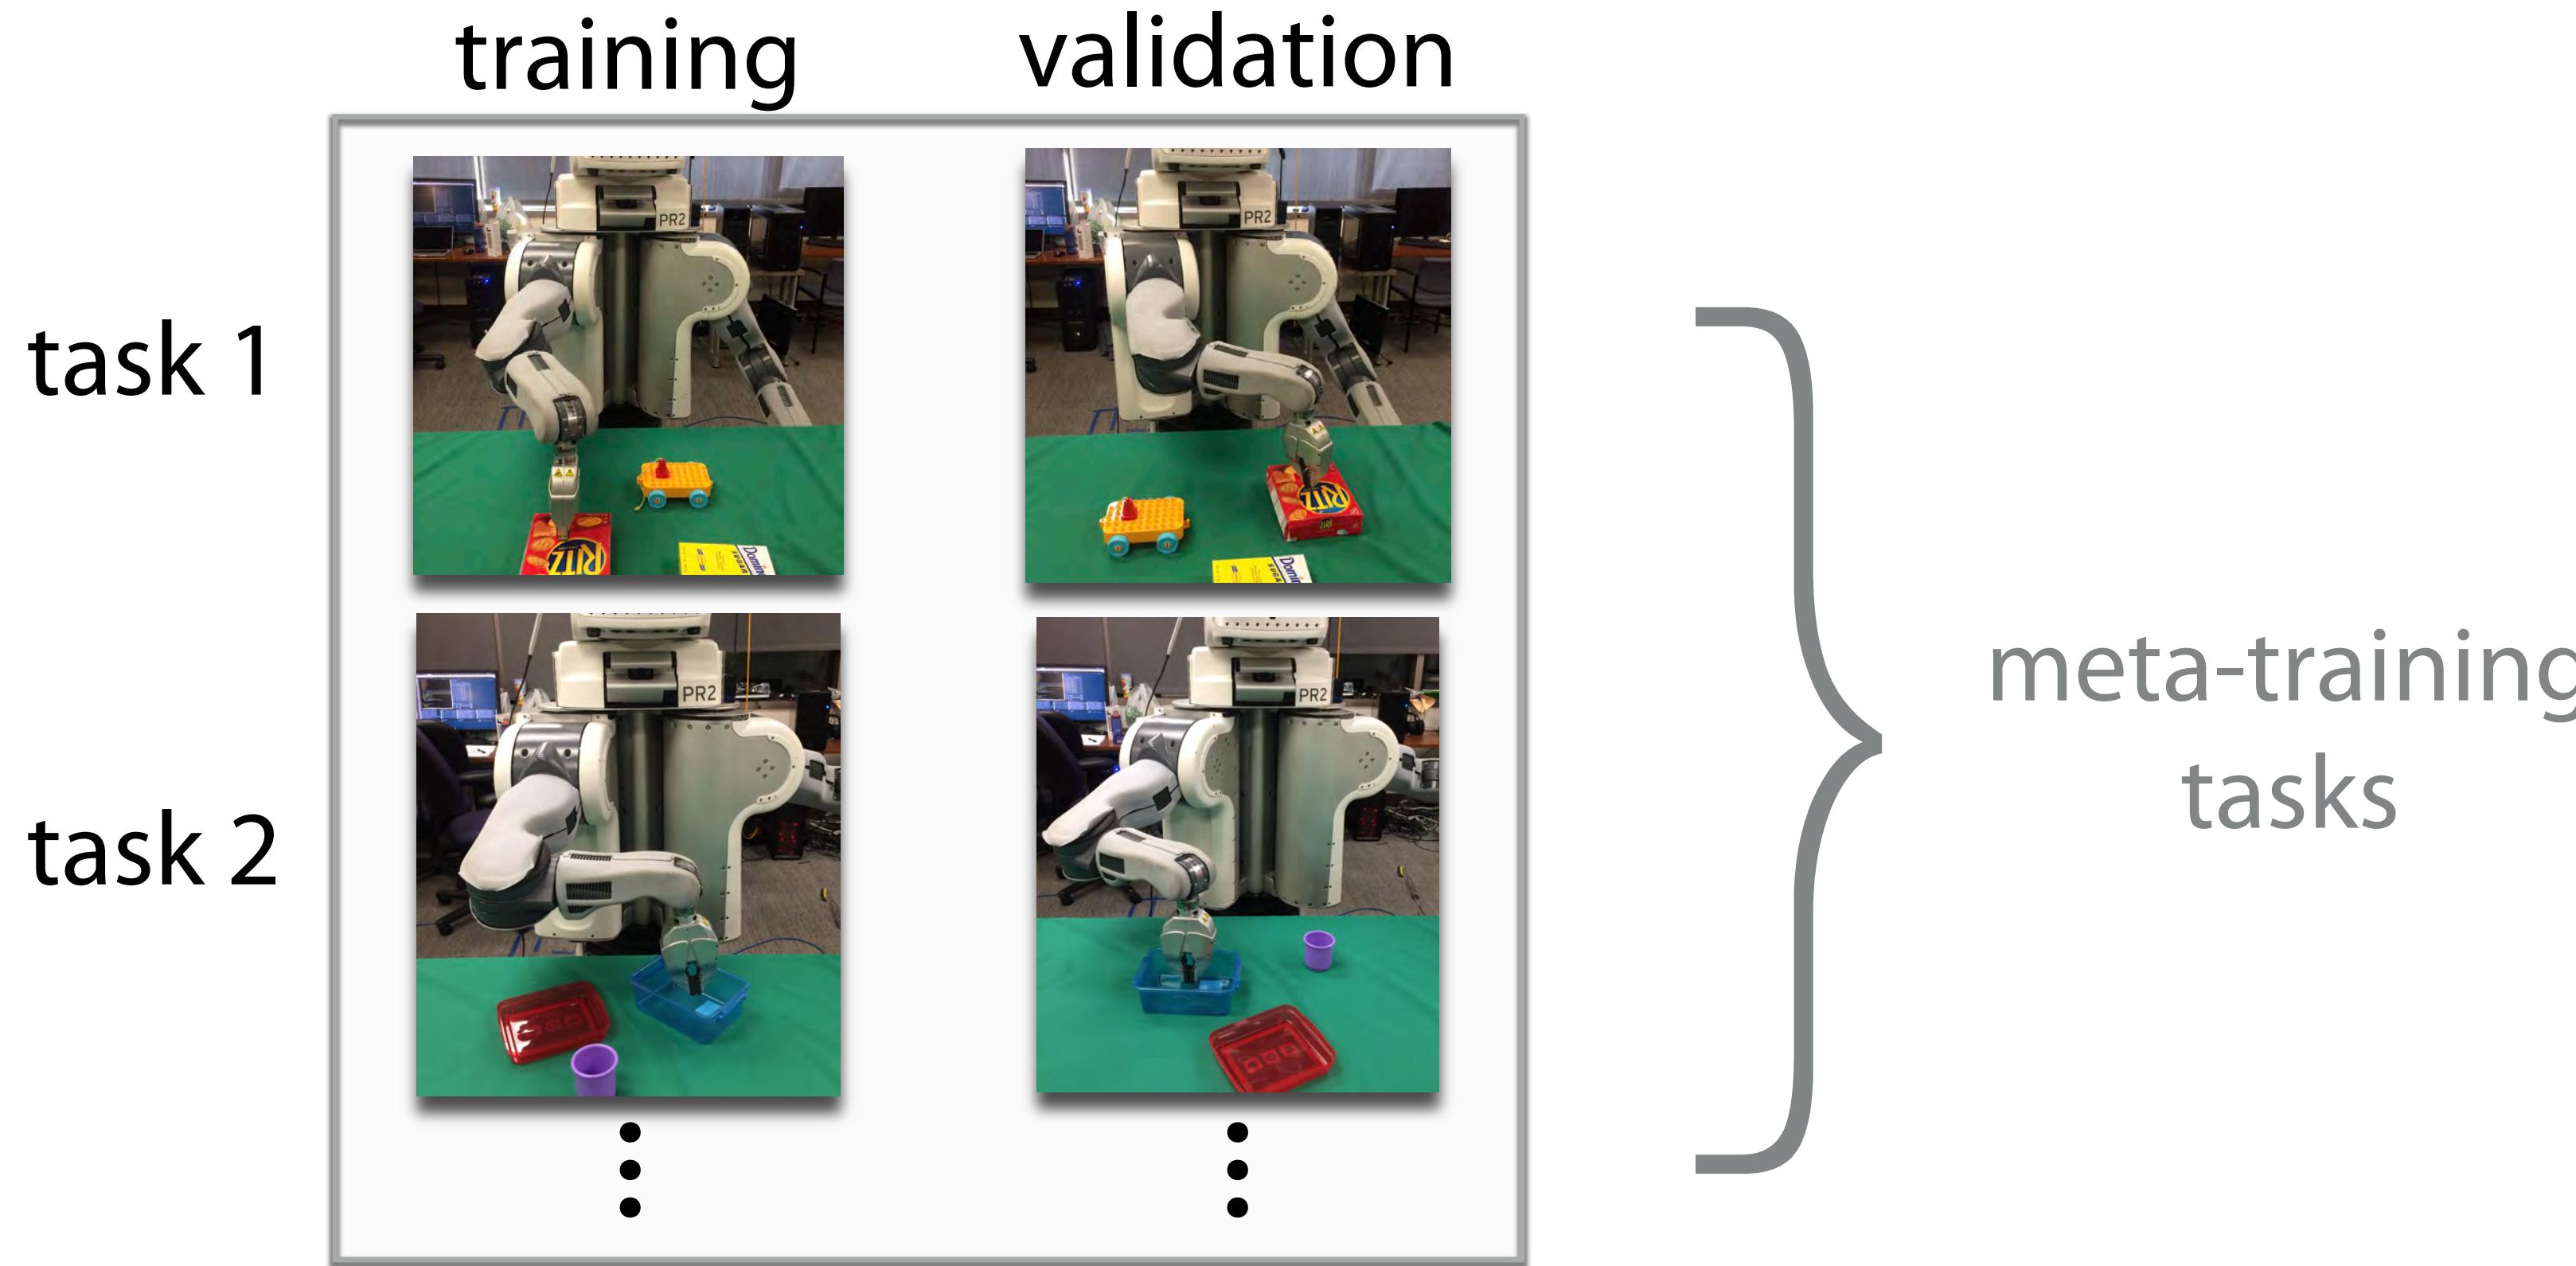
Meta-test time:



Learn policy from one demo.

Meta-Learning Approach

How? Learn to **infer the objective** from one demonstration



Meta-test time:

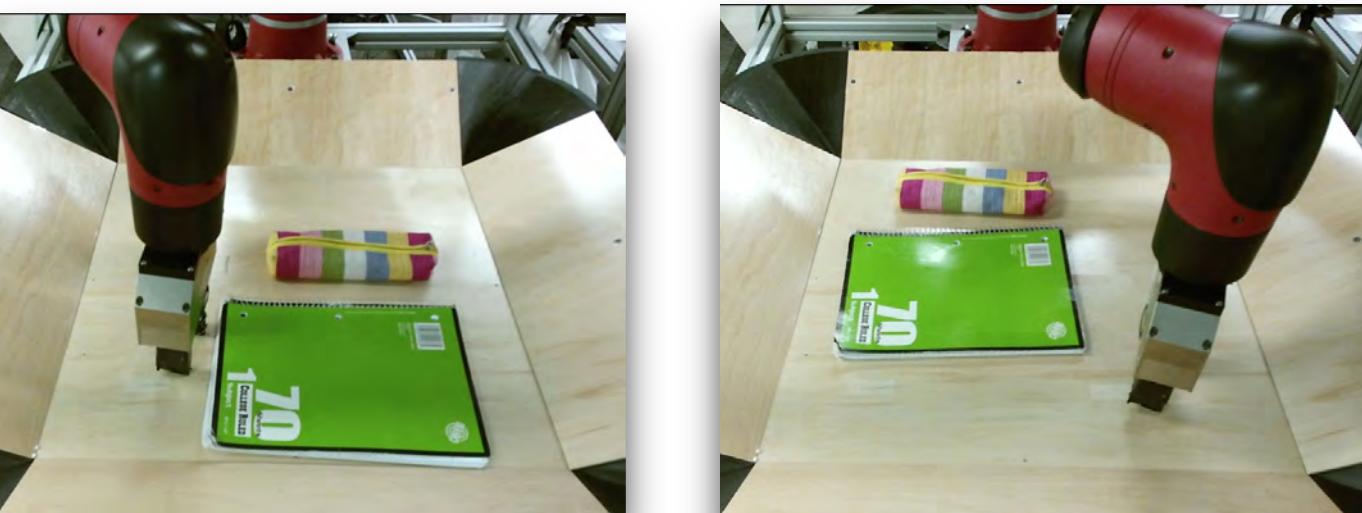


Learn **objective** from one demo.

Learn policy by optimizing reward.



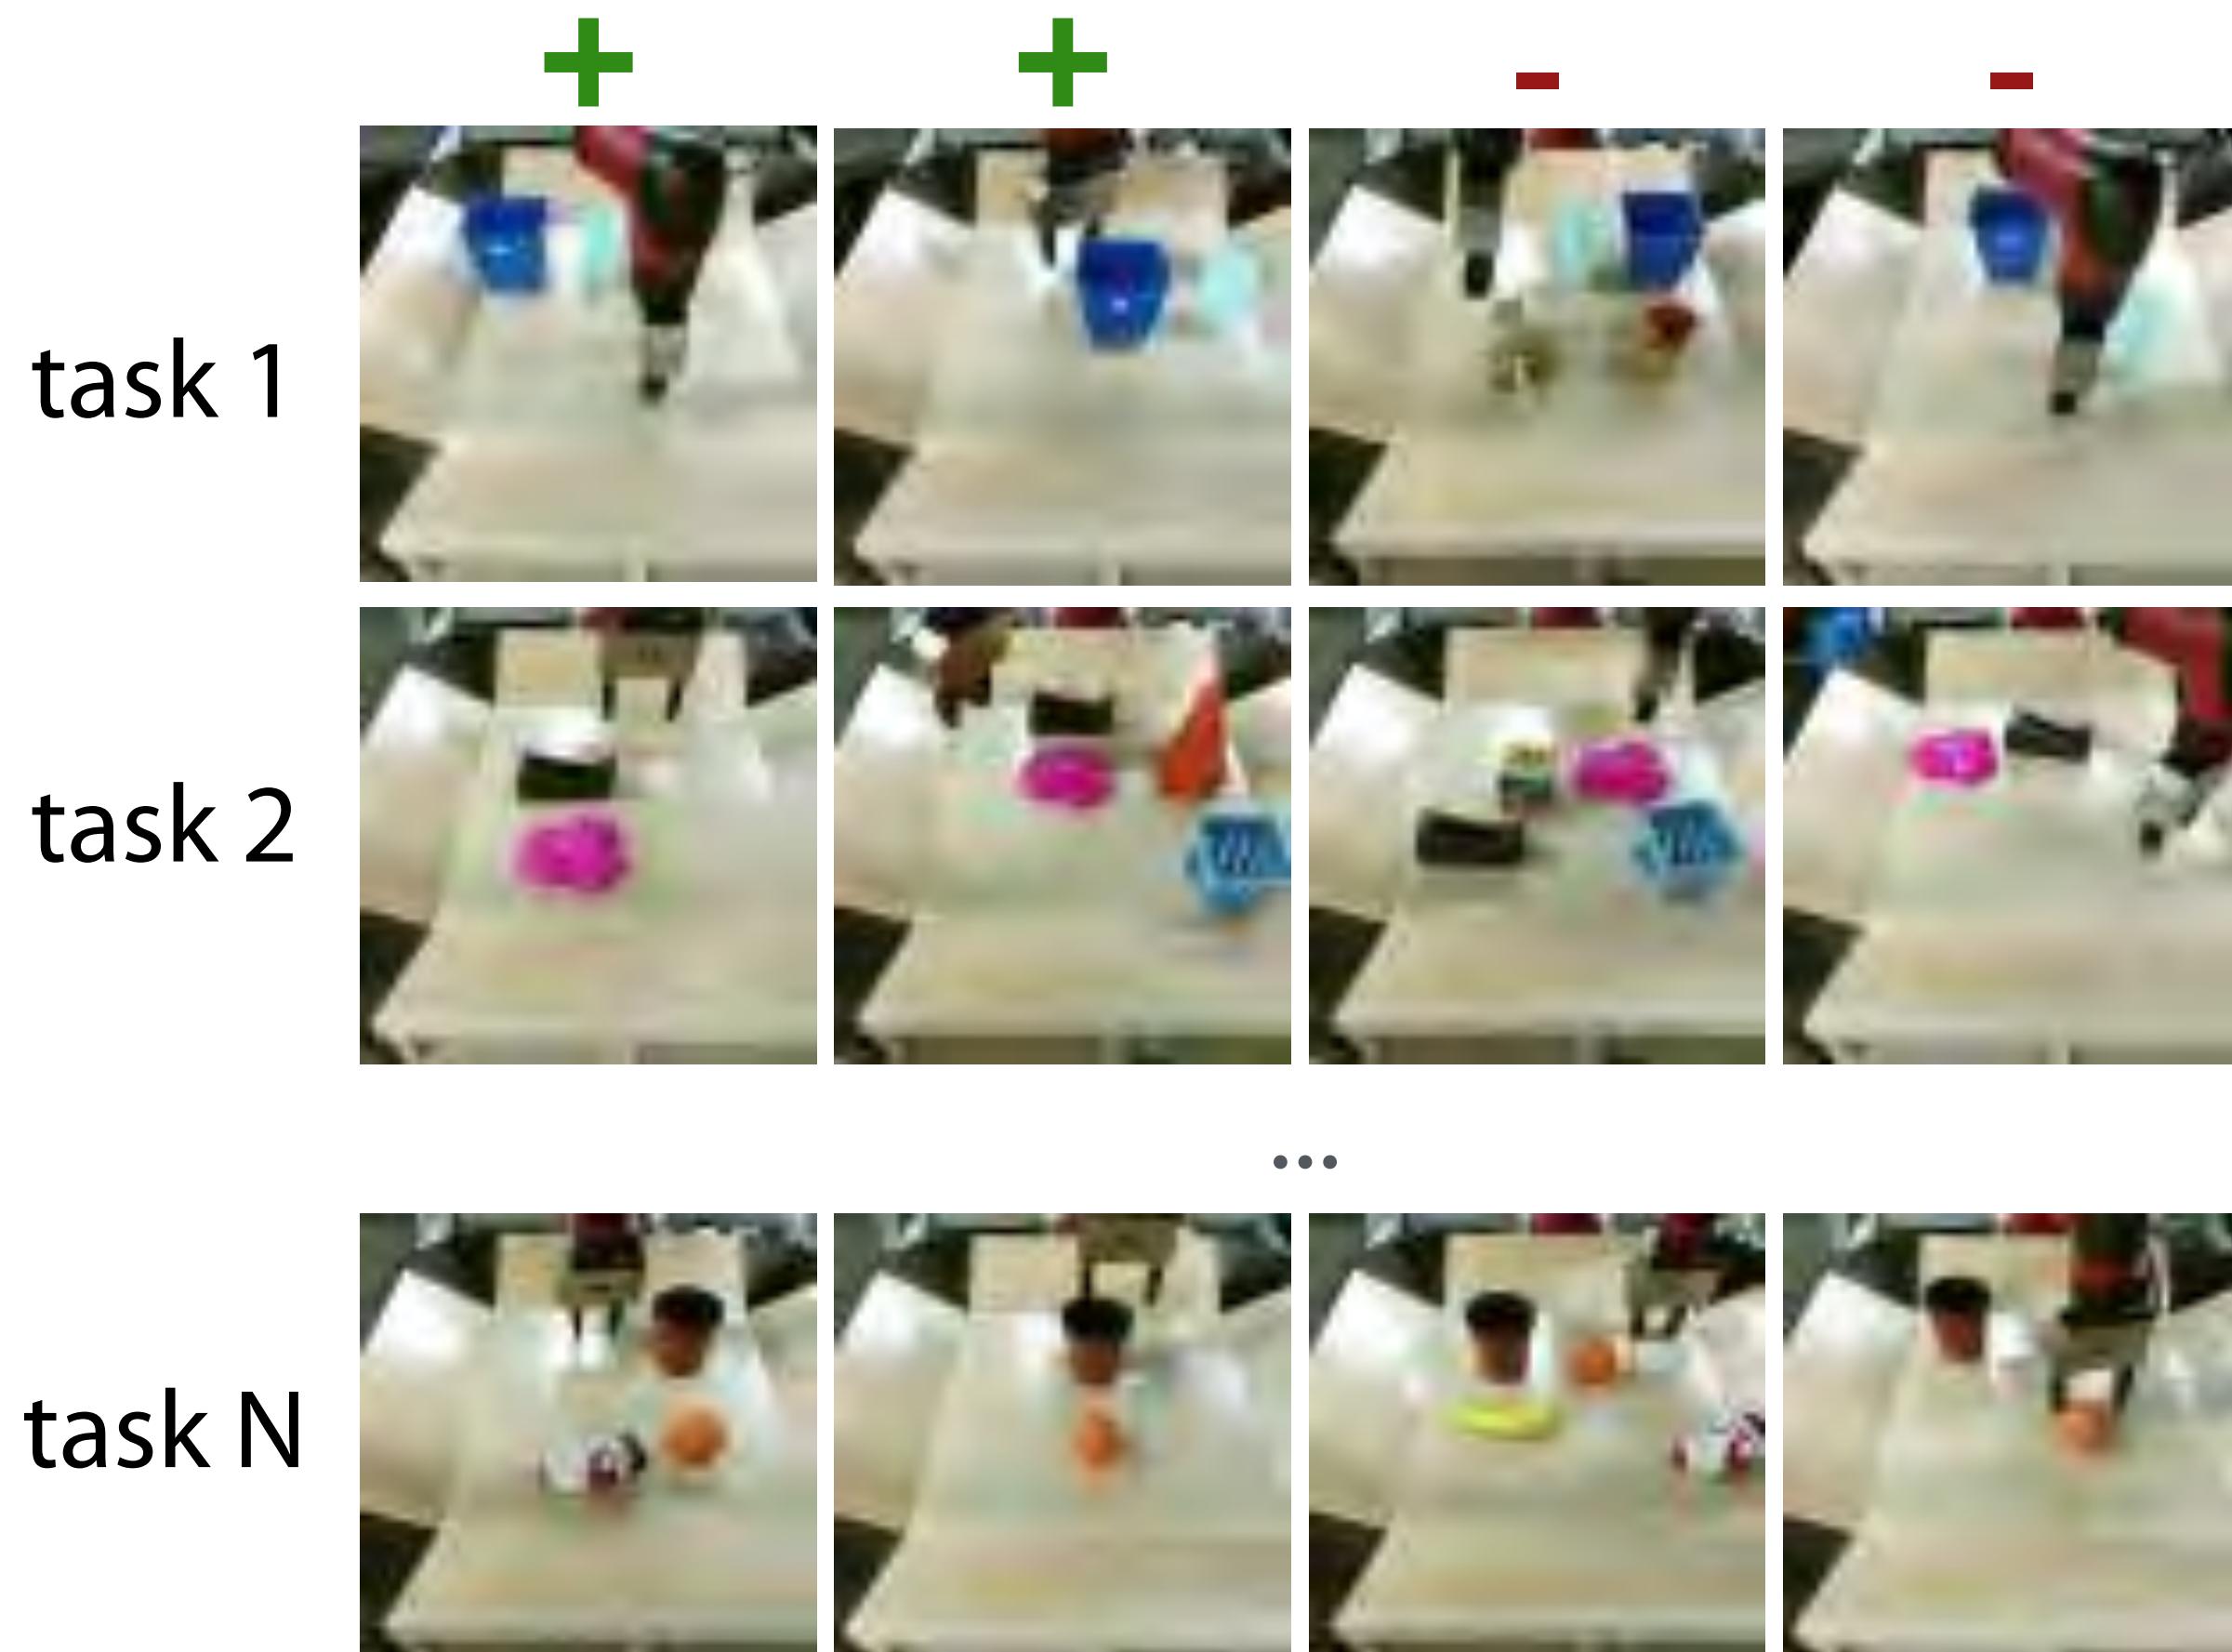
Given a few observations of reaching the goal:



Infer an objective:



(then run RL or planning)



meta-training
tasks

Given a few observations of reaching the goal:



Infer an objective:



(then run RL or planning)

MAML with supervised loss:

$$\min_{\theta} \sum_{\text{task } i} \mathcal{L}_{\text{test}}^i(\theta - \alpha \nabla_{\theta} \mathcal{L}_{\text{train}}^i(\theta))$$

both positive & negatives

only positive examples
another form of weak supervision!

Test time: only need positive examples

$$\theta \leftarrow \theta - \alpha \nabla_{\theta} \mathcal{L}_{\text{train}}(\theta)$$

Novel Object Positioning via Visual Planning

Given 5 examples of success

Visual MPC with learned objective



infer goal classifier

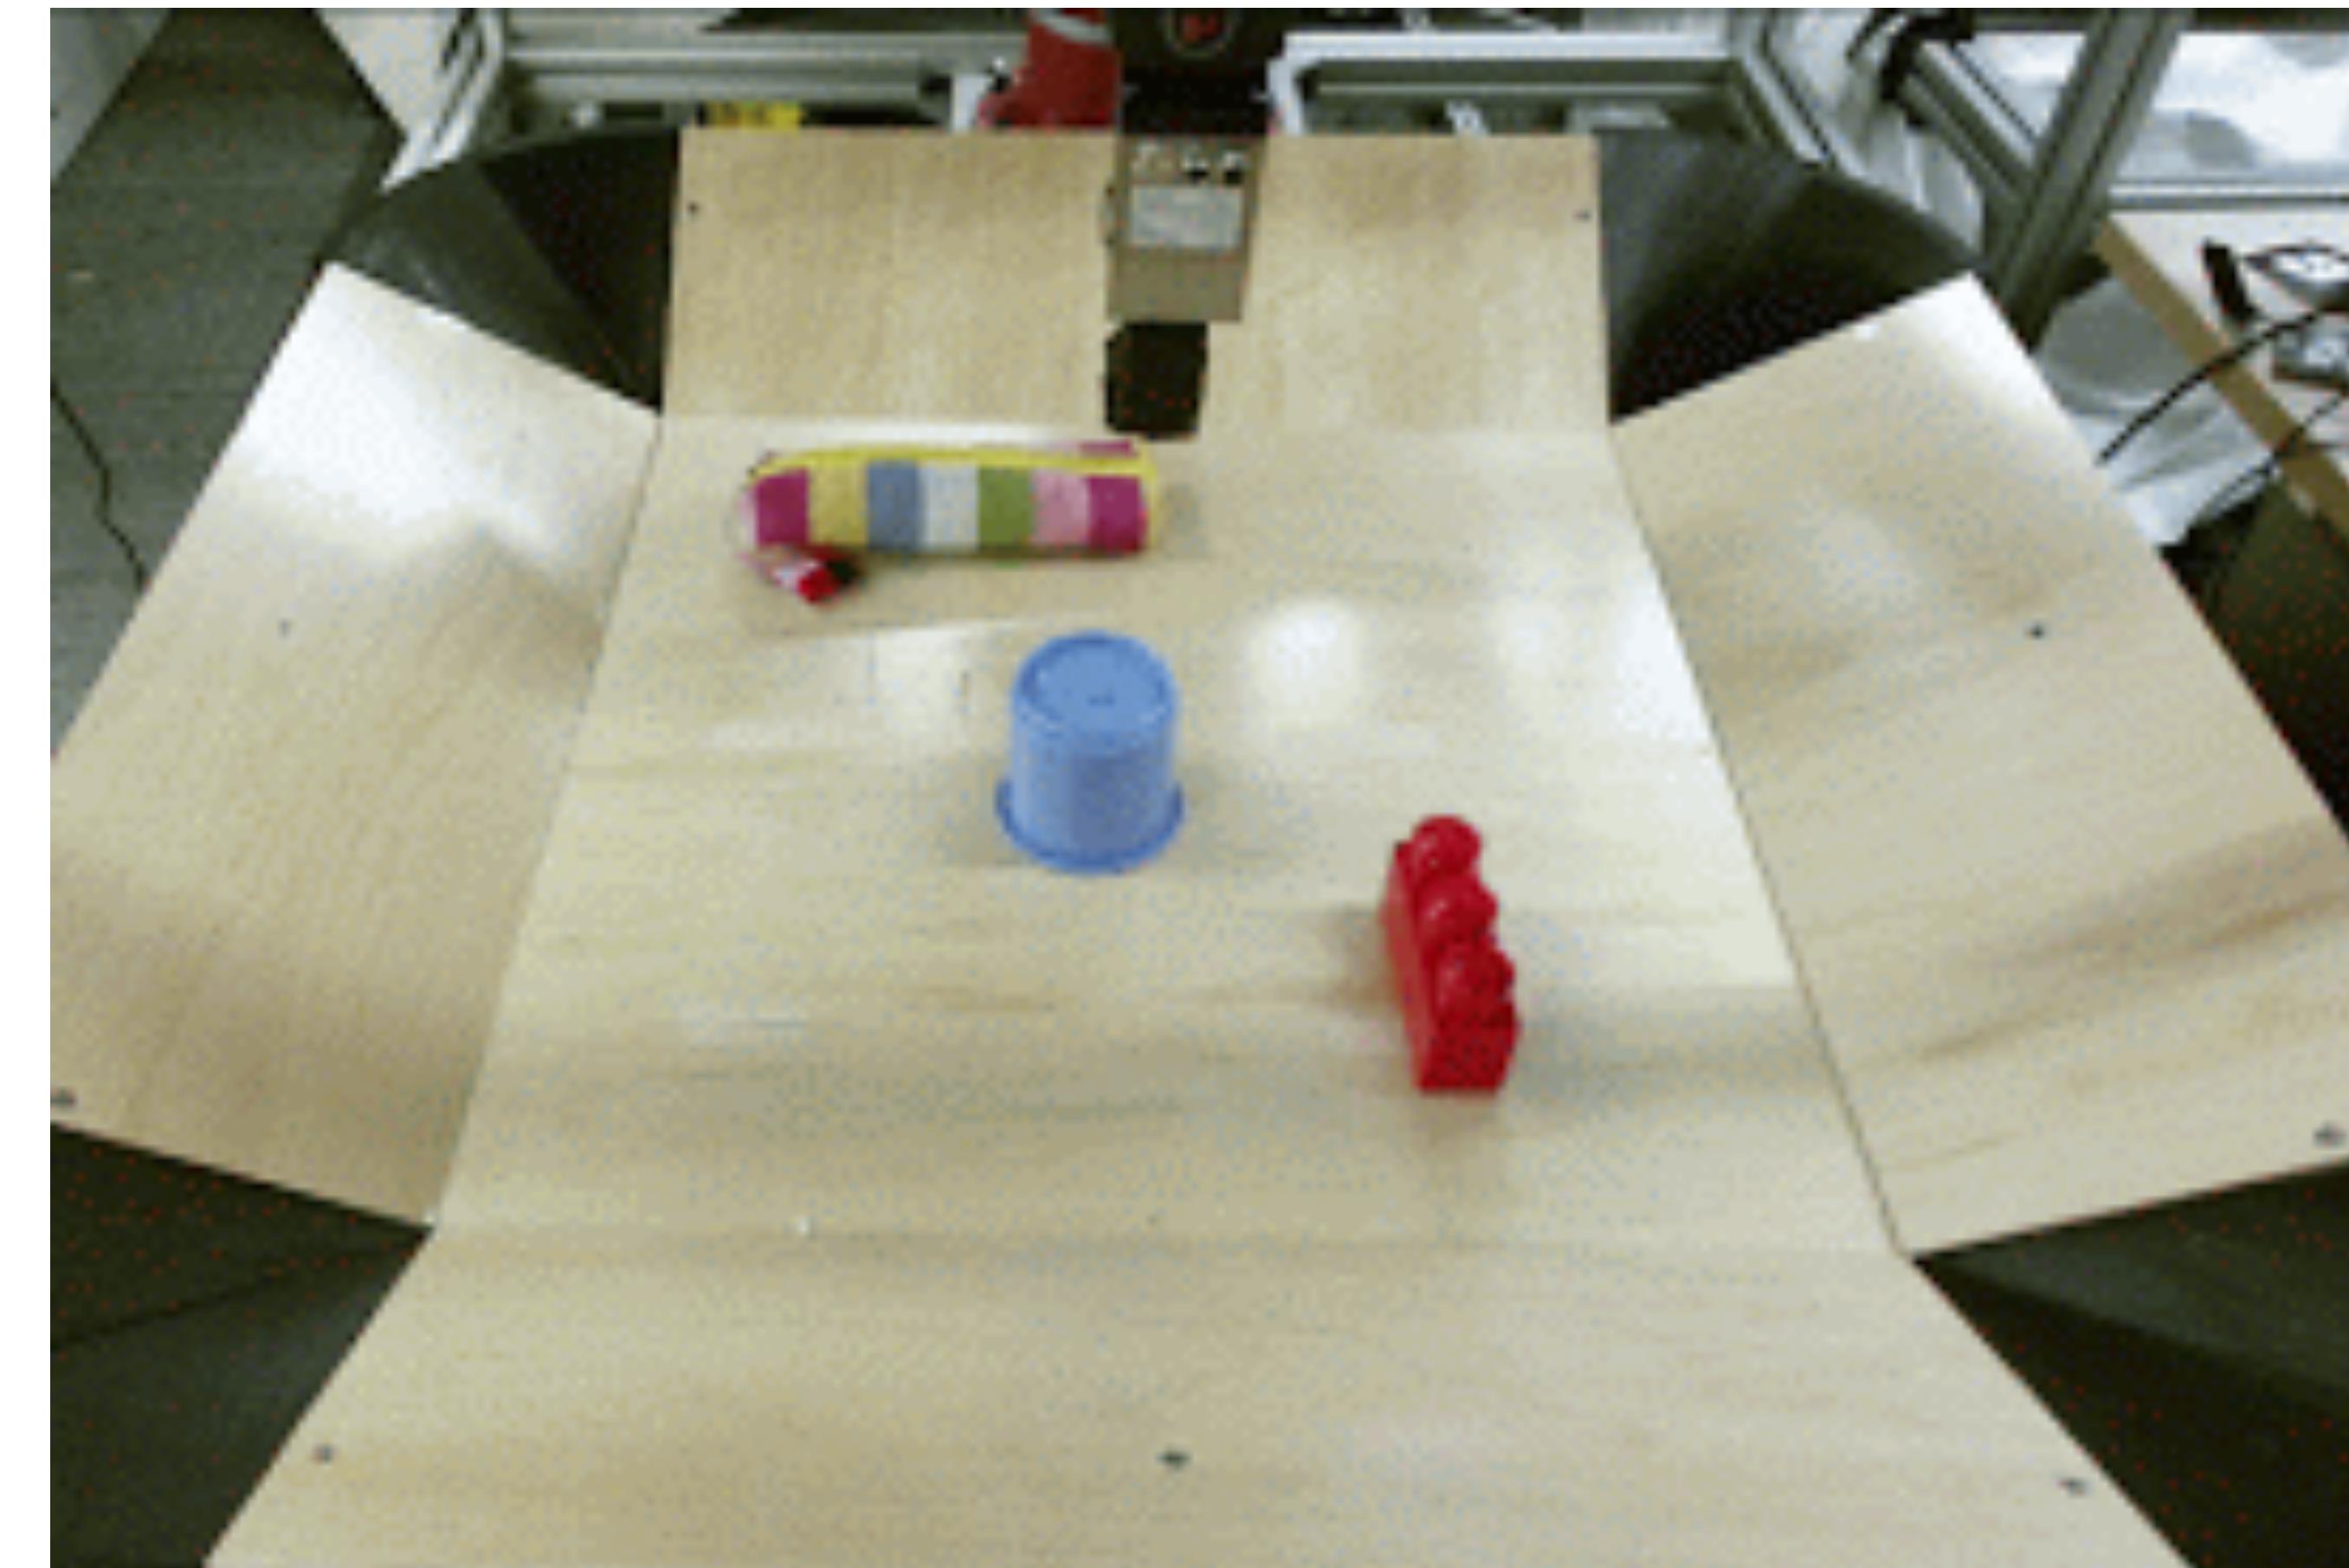
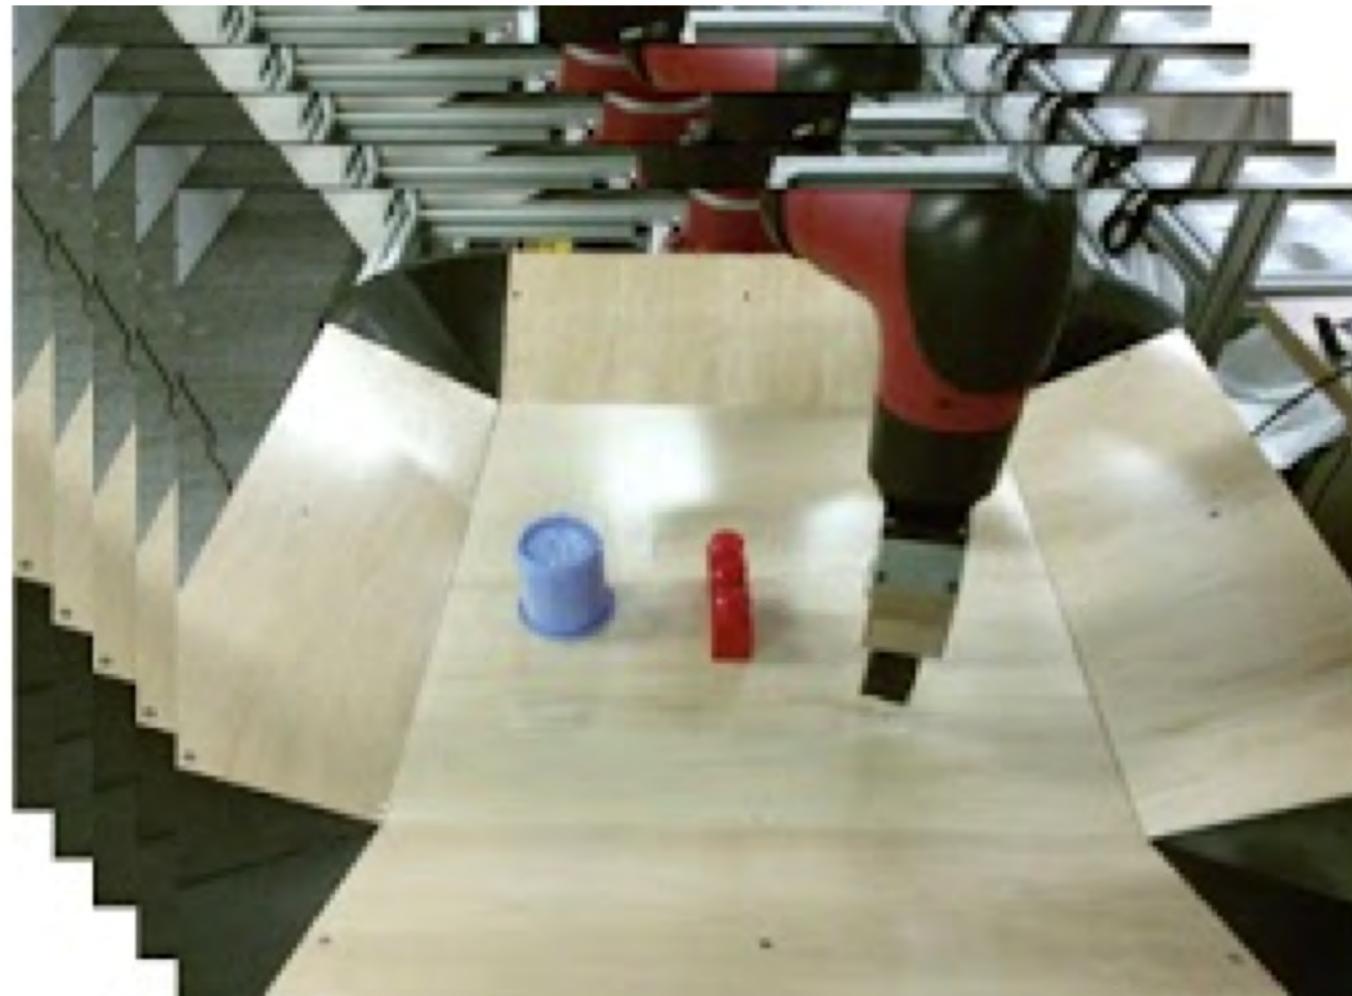
visual MPC w.r.t.
goal classifier



Novel Object Positioning via Visual Planning

Visual MPC with learned objective

Given 5 examples of success



Rope Manipulation via Reinforcement Learning

Given 5 examples of success



infer goal classifier



visual RL w.r.t. goal classifier

RL policy with learned objective



Rope Manipulation via Reinforcement Learning

RL policy with learned objective

Given 5 examples of success



Questions?

Meta-Learning Foundations

Finn, Abbeel, Levine, *Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks*. ICML '17

Finn & Levine, *Meta-Learning and Universality: Deep Representations and Gradient Descent can Approximate any Learning Algorithm*. ICLR '18

Grant, Finn, Levine, Darrell, Griffiths, *Recasting Gradient-Based Meta-Learning as Hierarchical Bayes*. ICLR '18

Finn*, Xu*, Levine, *Probabilistic Model-Agnostic Meta-Learning*. under review '18

Meta-Learning for Control

Finn*, Yu*, Zhang, Abbeel, Levine, *One-Shot Visual Imitation Learning via Meta-Learning*. CoRL '17

Yu*, Finn*, Xie, Dasari, Zhang, Abbeel, Levine, *One-Shot Imitation from Observing Humans via Domain-Adaptive Meta-Learning*. RSS '18

Clavera*, Nagabandi*, Fearing, Abbeel, Levine, Finn, *Learning to Adapt: Meta-Learning for Model-Based Control*. under review '18

Xu*, Ratner, Dragan, Levine, Finn, *Learning a Prior over Intent via Meta-Inverse Reinforcement Learning*. under review '18

Gupta, Eysenbach, Finn, Levine, *Unsupervised Meta-Learning for Reinforcement Learning*. under review '18

Xie, Singh, Levine, Finn, *Few-Shot Goal Inference for Visuomotor Learning and Planning*. under review '18